

CS 581

Advanced Artificial Intelligence

April 15, 2024

Announcements / Reminders

- Please follow the Week 13 To Do List instructions (if you haven't already)
- Programming Assignment #03: OPTIONAL/NOT FOR CREDIT
- **FINAL EXAM is on Monday (04/22/2024) in RE 104!**
 - different room!!!
 - IGNORE Registrar's FINAL EXAM date
 - Section 02: contact Mr. Charles Scott (scott@iit.edu) to make arrangements

Plan for Today

- Reinforcement Learning: Introduction
- Q-Learning

Refresher

- **MDPs**



- Value function V , action-value function Q , policy π
- Bellman equations
- Value iteration
- Policy iteration

- **Multi-armed bandits**

- Exploration vs exploitation trade-off
- ϵ -greedy approach

Markov Decision Process

- An MDP is defined by:
 - A set of states $s \in S$
 - A set of actions $a \in A$
 - A transition function $T(s, a, s')$
 - Prob that a from s leads to s'
 - i.e., $P(s' | s, a)$
 - Also called the model
 - A reward function $R(s, a, s')$
 - Sometimes just $R(s)$ or $R(s')$
 - A start state (or distribution)
 - Maybe a terminal state





Markov Decision Process

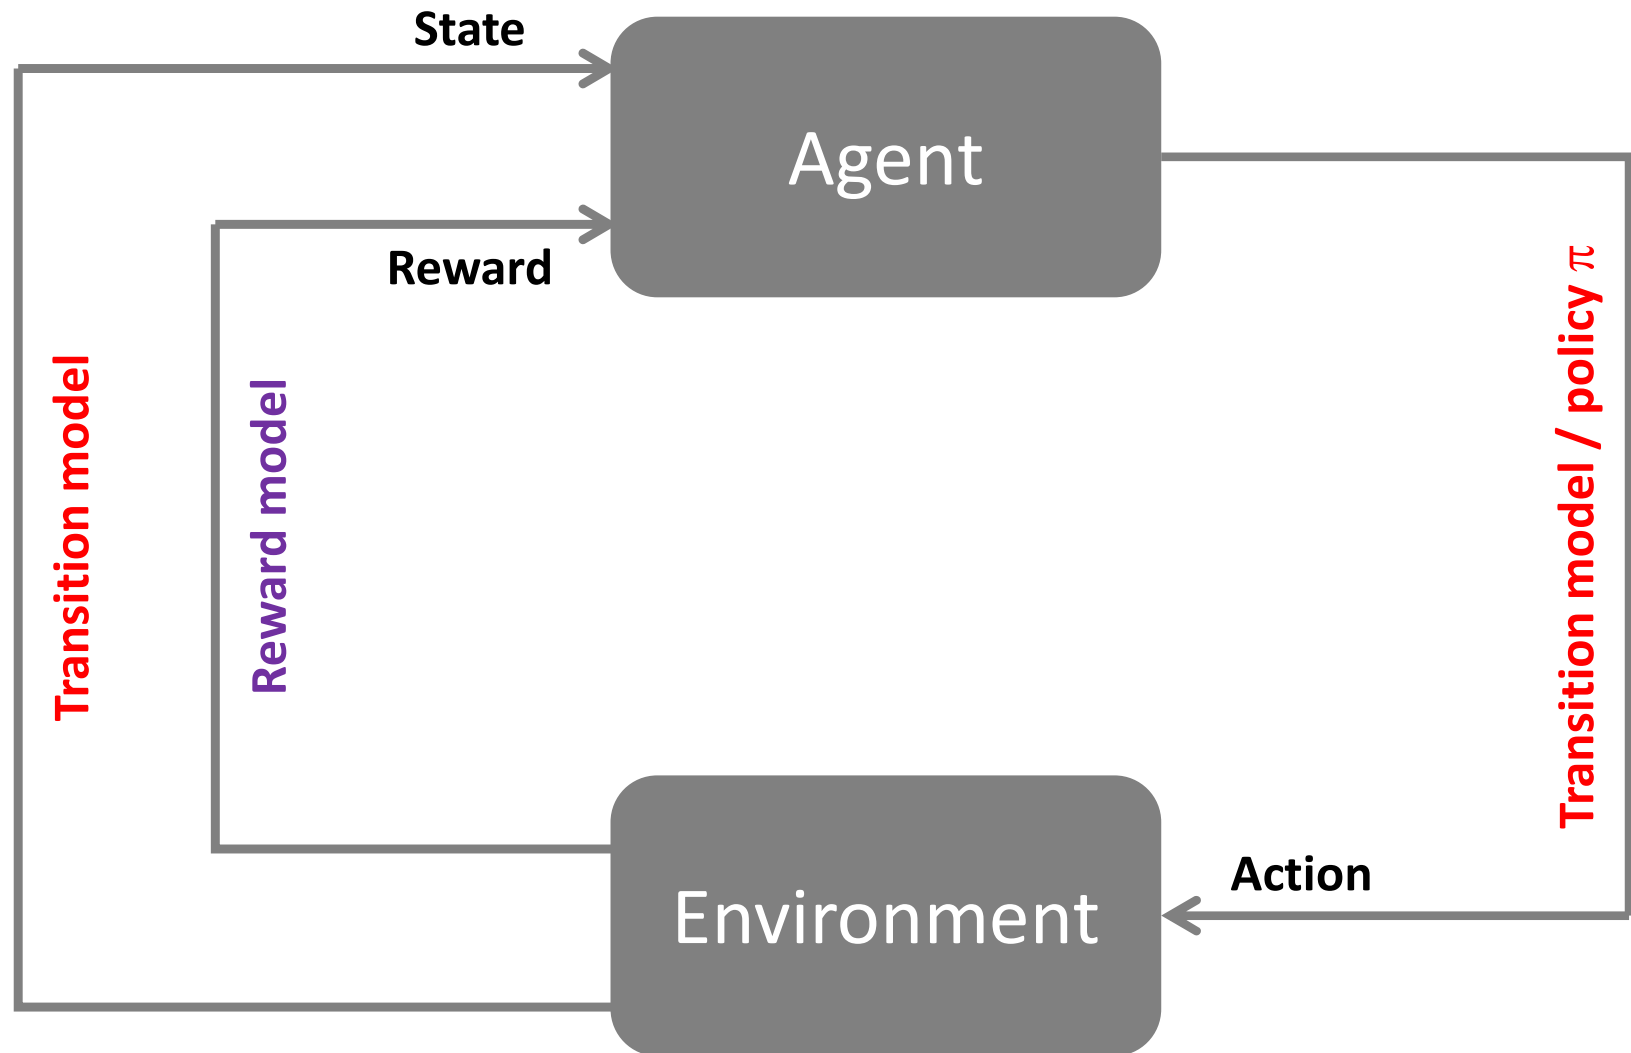
- An MDP is defined by:

- A set of states $s \in S$
- A set of actions $a \in A$
- A transition function $T(s,a,s')$
 - Prob that a from s leads to s'
 - i.e., $P(s' | s,a)$
 - Also called the model
- A reward function $R(s, a, s')$
 - Sometimes just $R(s)$ or $R(s')$
- A start state (or distribution)
- Maybe a terminal state

Model



Markov Decision Process



Solving MDPs

- **Offline algorithms:**
 - Value iteration
 - Policy iteration
 - Linear programming
- **Online algorithms:**
 - Approximation algorithms such as Monte Carlo planning

Bellman Equations

Bellman Equation

The **utility** of a **state** is the expected reward for the next transition plus the discounted utility of the **next state**, assuming that the agent chooses the optimal action:

$$U(s) = \max_{a \in A(s)} \sum_{s'} P(s' | s, a) * [R(s, a, s') + \gamma * U(s')]$$

Bellman Equation

The **utility** of a **state** is the expected reward for the next transition plus the discounted utility of the **next state**, assuming that the agent chooses the optimal action:

$$U(s) = \max_{a \in A(s)} \sum_{s'} P(s' | s, a) * [R(s, a, s') + \gamma * U(s')]$$

Utility / value of **current state s**

Bellman Equation

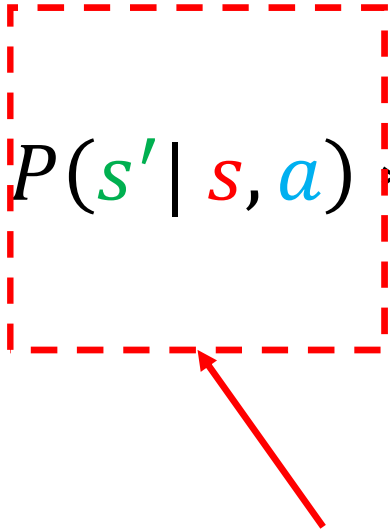
The **utility** of a **state** is the expected reward for the next transition plus the discounted utility of the **next state**, assuming that the agent chooses the optimal action:

$$U(s) = \max_{a \in A(s)} \sum_{s'} P(s' | s, a) * [R(s, a, s') + \gamma * U(s')]$$

Expected “long-term” **utility** / value after applying ONE specific **action** a [Need Environment Model!]

Bellman Equation

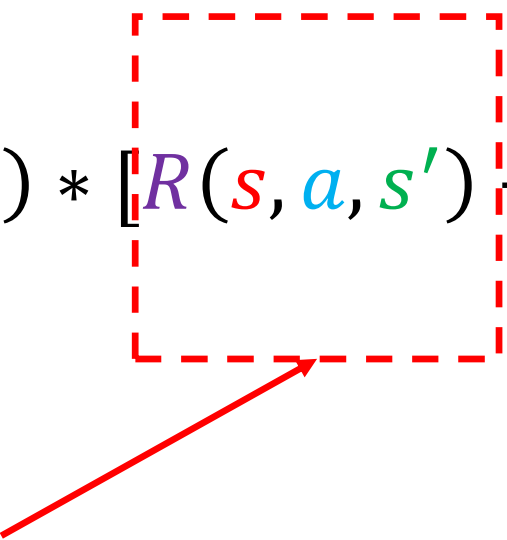
The **utility** of a **state** is the expected reward for the next transition plus the discounted utility of the **next state**, assuming that the agent chooses the optimal action:

$$U(s) = \max_{a \in A(s)} \sum_{s'} P(s' | s, a) * [R(s, a, s') + \gamma * U(s')]$$


Probability of transitioning FROM **current** state **s** TO **future** state **s'** after applying **action a** [Need Environment Model!]

Bellman Equation

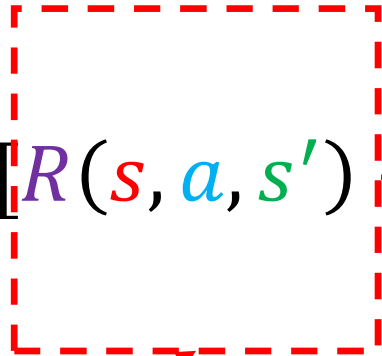
The **utility** of a **state** is the expected reward for the next transition plus the discounted utility of the **next state**, assuming that the agent chooses the optimal action:

$$U(s) = \max_{a \in A(s)} \sum_{s'} P(s' | s, a) * [R(s, a, s') + \gamma * U(s')]$$


Reward after transitioning FROM **current** state **s** TO **future** state **s'** after applying **action a** [Need Environment Model!]

Bellman Equation

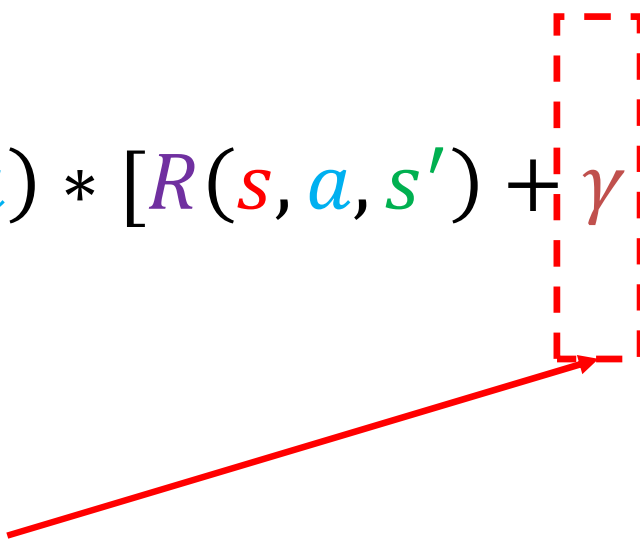
The **utility** of a **state** is the expected reward for the next transition plus the discounted utility of the **next state**, assuming that the agent chooses the optimal action:

$$U(s) = \max_{a \in A(s)} \sum_{s'} P(s' | s, a) * [R(s, a, s') + \gamma * U(s')]$$


CURRENT / SINGLE transition reward
[Need Environment Model!]

Bellman Equation

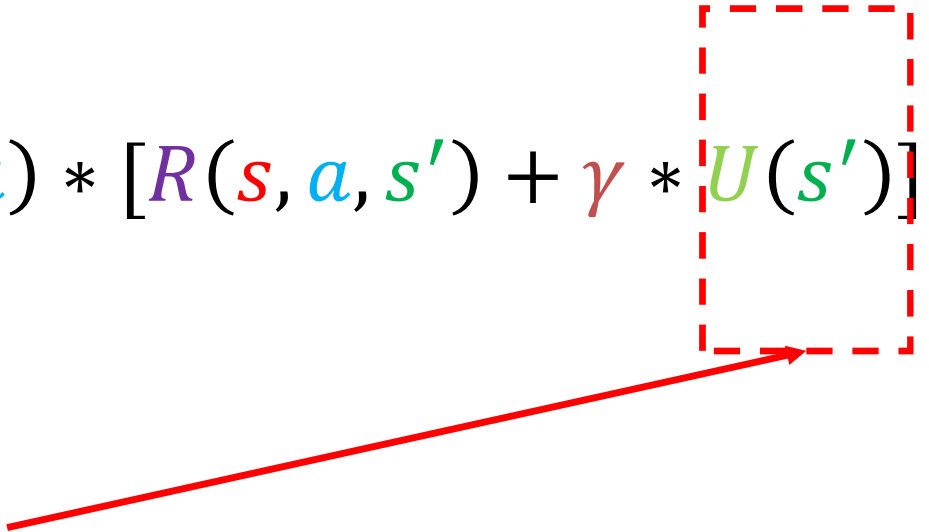
The **utility** of a **state** is the expected reward for the next transition plus the discounted utility of the **next state**, assuming that the agent chooses the optimal action:

$$U(s) = \max_{a \in A(s)} \sum_{s'} P(s' | s, a) * [R(s, a, s') + \gamma * U(s')]$$


Discount factor

Bellman Equation

The **utility** of a **state** is the expected reward for the next transition plus the discounted utility of the **next state**, assuming that the agent chooses the optimal action:

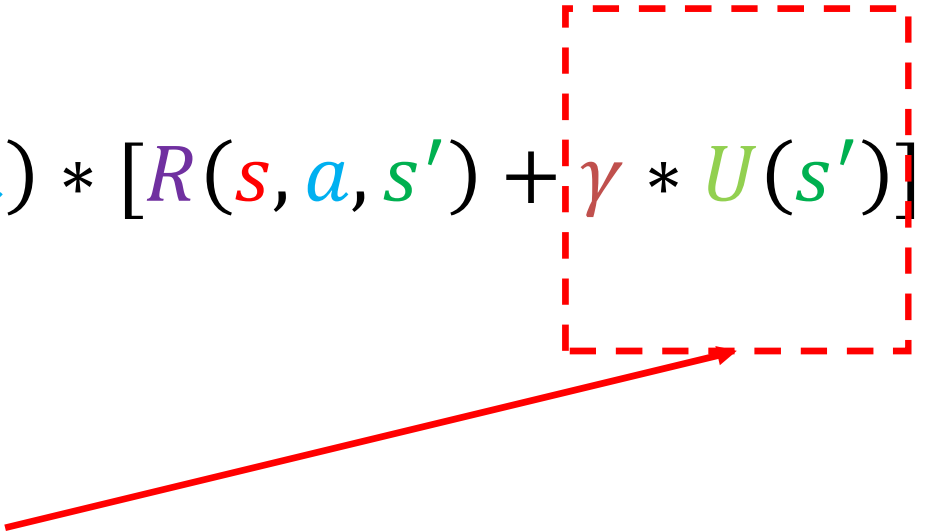
$$U(s) = \max_{a \in A(s)} \sum_{s'} P(s' | s, a) * [R(s, a, s') + \gamma * U(s')]$$


Future state s' utility

[can be a rough estimate at the beginning]

Bellman Equation

The **utility** of a **state** is the expected reward for the next transition plus the discounted utility of the **next state**, assuming that the agent chooses the optimal action:

$$U(s) = \max_{a \in A(s)} \sum_{s'} P(s' | s, a) * [R(s, a, s') + \gamma * U(s')]$$


Discounted future state s' utility

[can be a rough estimate at the beginning]

Bellman Equation

The **utility** of a **state** is the expected reward for the next transition plus the discounted utility of the **next state**, assuming that the agent chooses the optimal action:

$$U(s) = \max_{a \in A(s)} \sum_{s'} P(s' | s, a) * [R(s, a, s') + \gamma * U(s')]$$

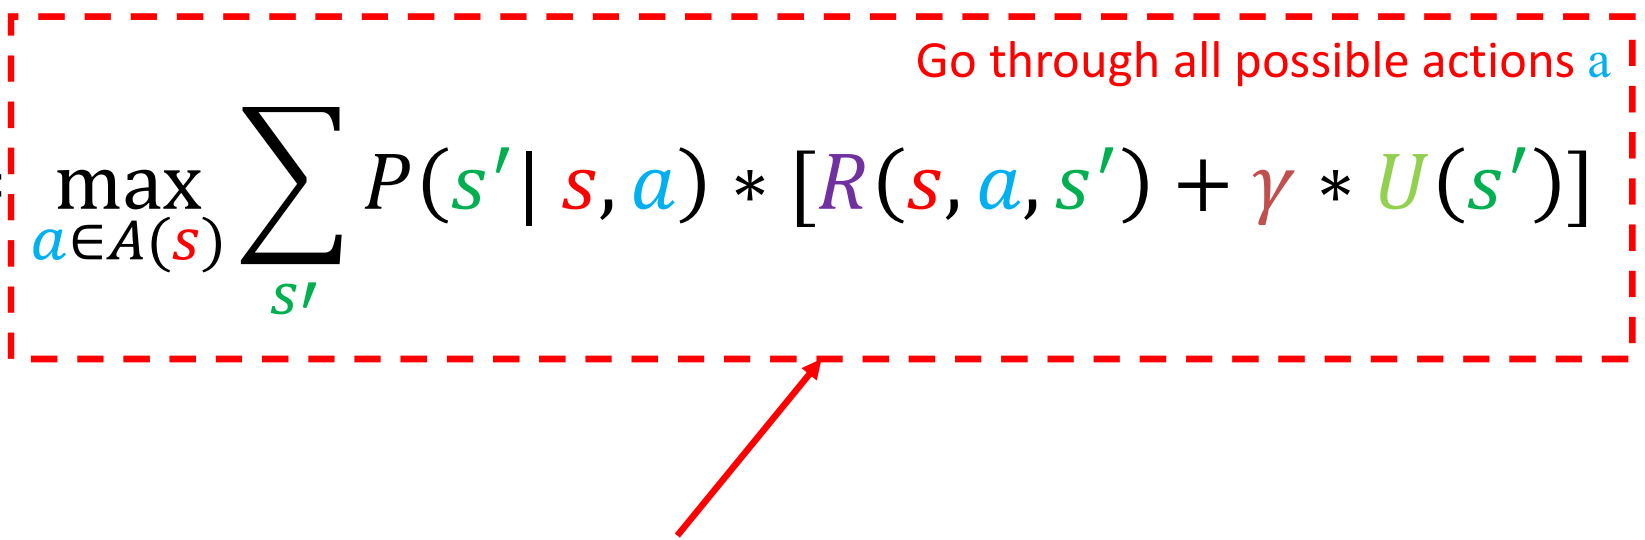
BEST Expected “long-term” **utility** / value after applying ONE specific BEST **action** a [Need Environment Model!]

Bellman Equation

The **utility** of a **state** is the expected **reward** for the next transition plus the discounted **utility** of the **next state**, assuming that the agent chooses the optimal action:

$$U(s) = \max_{a \in A(s)} \sum_{s'} P(s' | s, a) * [R(s, a, s') + \gamma * U(s')]$$

Go through all possible actions a



BEST Expected “long-term” **utility** / value after applying ONE specific BEST **action** a [Need Environment Model!]

Bellman Equation

The **utility** of a **state** is the expected reward for the next transition plus the discounted utility of the **next state**, assuming that the agent chooses the optimal action:

$$U(s) = \max_{a \in A(s)} \sum_{s'} P(s' | s, a) * [R(s, a, s') + \gamma * U(s')]$$

Go through all possible future states s'

Expected “long-term” **utility** / value after applying ONE specific **action** a [Need Environment Model!]

Expected Utility Given Policy π

The **utility** of a **state** is the expected reward for the next transition plus the discounted utility of the **next state**, assuming that the agent uses policy π :

$$U^\pi(s) = \sum_{s'} P(s' | s, \pi(s)) * [R(s, \pi(s), s') + \gamma * U^\pi(s')]$$

Bellman Optimality

The **utility** of a **state** is the expected **reward** for the next transition plus the discounted **utility** of the **next state**, assuming that the agent uses policy π :

$$U^\pi(s) = \sum_{s'} P(s' | s, \pi(s)) * [R(s, \pi(s), s') + \gamma * U^\pi(s')]$$

Bellman Update

Iterative utility update at $i+1$ iteration can be calculated with:

$$U_{i+1}(s) = \max_{a \in A(s)} \sum_{s'} P(s' | s, a) * [R(s, a, s') + \gamma * U_i(s')]$$

Bellman Equation: Example

3	0.812	0.868	0.918	+ 1
2	0.762		0.660	- 1
1	0.705	0.655	0.611	0.388
	1	2	3	4

Note: ALL non-terminal transitions have a reward $r = -0.04$

$$\begin{aligned}
 U(1, 1) = -0.04 + \gamma \max[& 0.8U(1, 2) + 0.1U(2, 1) + 0.1U(1, 1), & (Up) \\
 & 0.9U(1, 1) + 0.1U(1, 2), & (Left) \\
 & 0.9U(1, 1) + 0.1U(2, 1), & (Down) \\
 & 0.8U(2, 1) + 0.1U(1, 2) + 0.1U(1, 1)]. & (Right)
 \end{aligned}$$

Value Iteration Algorithm

function VALUE-ITERATION(mdp, ϵ) **returns** a utility function

inputs: mdp , an MDP with states S , actions $A(s)$, transition model $P(s' | s, a)$,
rewards $R(s)$, discount γ

ϵ , the maximum error allowed in the utility of any state

local variables: U, U' , vectors of utilities for states in S , initially zero

δ , the maximum change in the utility of any state in an iteration

repeat

$U \leftarrow U'; \delta \leftarrow 0$

for each state s **in** S **do**

$U'[s] \leftarrow R(s) + \gamma \max_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s']$

if $|U'[s] - U[s]| > \delta$ **then** $\delta \leftarrow |U'[s] - U[s]|$

until $\delta < \epsilon(1 - \gamma)/\gamma$

return U

Value Iteration Algorithm

function VALUE-ITERATION(mdp, ϵ) **returns** a utility function

inputs: mdp , an MDP with states S , actions $A(s)$, transition model $P(s' | s, a)$,
rewards $R(s)$, discount γ

ϵ , the maximum error allowed in the utility of any state

local variables: U, U' , vectors of utilities for states in S , initially zero

δ , the maximum change in the utility of any state in an iteration

repeat

N states: N Bellman Equations to iteratively “solve”

$U \leftarrow U'; \delta \leftarrow 0$

for each state s **in** S **do**

$U'[s] \leftarrow R(s) + \gamma \max_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s']$

if $|U'[s] - U[s]| > \delta$ **then** $\delta \leftarrow |U'[s] - U[s]|$

until $\delta < \epsilon(1 - \gamma)/\gamma$

return U

Value Iteration Algorithm

function VALUE-ITERATION(mdp, ϵ) **returns** a utility function

inputs: mdp , an MDP with states S , actions $A(s)$, transition model $P(s' | s, a)$,
rewards $R(s)$, discount γ

ϵ , the maximum error allowed in the utility of any state

local variables: U, U' , vectors of utilities for states in S , initially zero

δ , the maximum change in the utility of any state in an iteration

repeat

Repeat infinite times: guaranteed convergence

$U \leftarrow U'; \delta \leftarrow 0$

for each state s **in** S **do**



$U'[s] \leftarrow R(s) + \gamma \max_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s']$

if $|U'[s] - U[s]| > \delta$ **then** $\delta \leftarrow |U'[s] - U[s]|$

until $\delta < \epsilon(1 - \gamma)/\gamma$

return U

Value Iteration: Convergence



Policy Iteration:

- Start with initial policy π_0
- Policy iteration algorithm involves (alternates between) two steps
 - Policy evaluation: given a policy π_i , calculate $U_i = U^{\pi_i}$, the utility of each state if π_i were to be executed
 - Policy improvement: calculate a new MEU policy π_{i+1} , using a one step look-ahead based on U_i

$$\pi^*(s) = \operatorname{argmax}_{a \in A(s)} \sum_{s'} P(s' | s, a) * [R(s, a, s') + \gamma * U(s')]$$

Bellman Update

Iterative utility update at $i+1$ iteration can be calculated with:

$$U_{i+1}(s) = \max_{a \in A(s)} \sum_{s'} P(s' | s, a) * [R(s, a, s') + \gamma * U_i(s')]$$

Policy Iteration Algorithm

function POLICY-ITERATION(mdp) **returns** a policy

inputs: mdp , an MDP with states S , actions $A(s)$, transition model $P(s' | s, a)$

local variables: U , a vector of utilities for states in S , initially zero

π , a policy vector indexed by state, initially random

repeat

$U \leftarrow \text{POLICY-EVALUATION}(\pi, U, mdp)$

$unchanged? \leftarrow \text{true}$

for each state s **in** S **do**

if $\max_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s'] > \sum_{s'} P(s' | s, \pi[s]) U[s']$ **then do**

$\pi[s] \leftarrow \operatorname{argmax}_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s']$

$unchanged? \leftarrow \text{false}$

until $unchanged?$

return π

Policy Iteration Algorithm

function POLICY-ITERATION(mdp) **returns** a policy

inputs: mdp , an MDP with states S , actions $A(s)$, transition model $P(s' | s, a)$

local variables: U , a vector of utilities for states in S , initially zero

π , a policy vector indexed by state, initially random

repeat

$U \leftarrow \text{POLICY-EVALUATION}(\pi, U, mdp)$

$unchanged? \leftarrow \text{true}$

Policy evaluation

for each state s **in** S **do**

if $\max_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s'] > \sum_{s'} P(s' | s, \pi[s]) U[s']$ **then do**

$\pi[s] \leftarrow \operatorname{argmax}_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s']$

$unchanged? \leftarrow \text{false}$

until $unchanged?$

return π

Policy Iteration Algorithm

function POLICY-ITERATION(mdp) **returns** a policy

inputs: mdp , an MDP with states S , actions $A(s)$, transition model $P(s' | s, a)$

local variables: U , a vector of utilities for states in S , initially zero

π , a policy vector indexed by state, initially random

repeat

$U \leftarrow \text{POLICY-EVALUATION}(\pi, U, mdp)$

$unchanged? \leftarrow \text{true}$

for each state s **in** S **do**

if $\max_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s'] > \sum_{s'} P(s' | s, \pi[s]) U[s']$ **then do**

$\pi[s] \leftarrow \operatorname{argmax}_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s']$

$unchanged? \leftarrow \text{false}$

Policy improvement

until $unchanged?$

return π

Policy Iteration Algorithm

function POLICY-ITERATION(mdp) **returns** a policy

inputs: mdp , an MDP with states S , actions $A(s)$, transition model $P(s' | s, a)$

local variables: U , a vector of utilities for states in S , initially zero

π , a policy vector indexed by state, initially random

repeat

$U \leftarrow \text{POLICY-EVALUATION}(\pi, U, mdp)$

$unchanged? \leftarrow \text{true}$

for each state s **in** S **do**

if $\max_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s'] > \sum_{s'} P(s' | s, \pi[s]) U[s']$ **then do**

$\pi[s] \leftarrow \operatorname{argmax}_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s']$

$unchanged? \leftarrow \text{false}$

Recalculate policy (find new MEU policy) for all s

until $unchanged?$

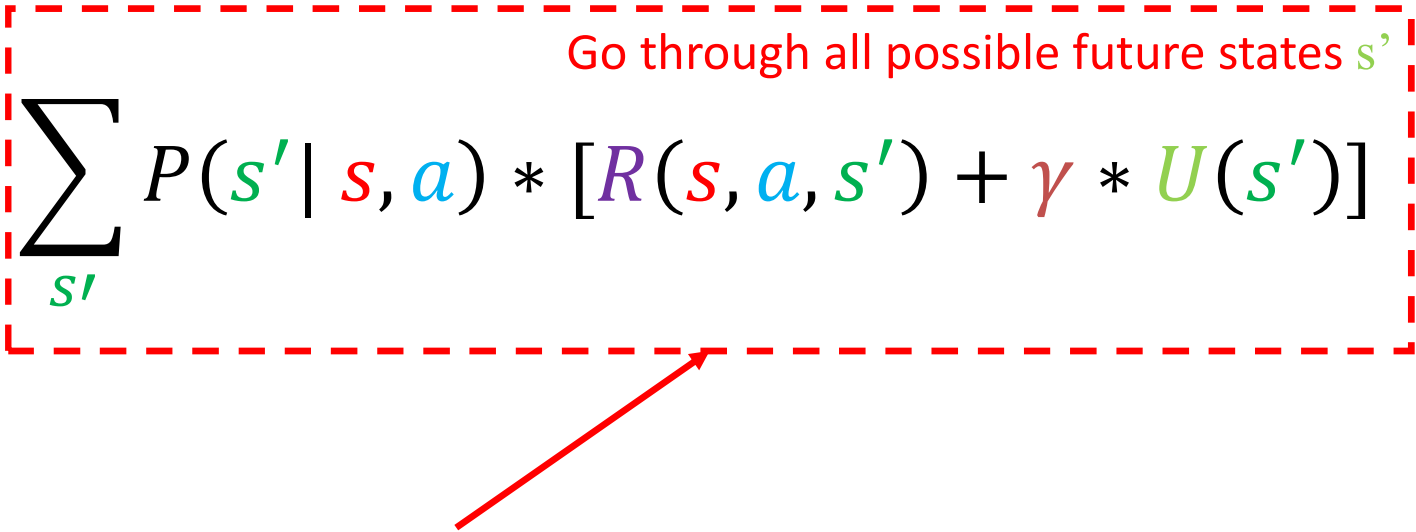
return π

Bellman Equation

The **utility** of a **state** is the expected reward for the next transition plus the discounted utility of the **next state**, assuming that the agent chooses the optimal action:

$$U(s) = \max_{a \in A(s)} \sum_{s'} P(s' | s, a) * [R(s, a, s') + \gamma * U(s')]$$

Go through all possible future states s'



Expected “long-term” **utility** / value after applying ONE specific **action** a [Need Environment Model!]

Policy Iteration Algorithm

function POLICY-ITERATION(mdp) **returns** a policy

inputs: mdp , an MDP with states S , actions $A(s)$, transition model $P(s' | s, a)$

local variables: U , a vector of utilities for states in S , initially zero

π , a policy vector indexed by state, initially random

repeat

$U \leftarrow \text{POLICY-EVALUATION}(\pi, U, mdp)$

$unchanged? \leftarrow \text{true}$

for each state s **in** S **do**

if $\max_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s'] > \sum_{s'} P(s' | s, \pi[s]) U[s']$ **then do**

Better action found

$\pi[s] \leftarrow \operatorname{argmax}_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s']$

$unchanged? \leftarrow \text{false}$

until $unchanged?$

return π

Policy Iteration Algorithm

function POLICY-ITERATION(mdp) **returns** a policy

inputs: mdp , an MDP with states S , actions $A(s)$, transition model $P(s' | s, a)$

local variables: U , a vector of utilities for states in S , initially zero

π , a policy vector indexed by state, initially random

repeat

$U \leftarrow \text{POLICY-EVALUATION}(\pi, U, mdp)$

$unchanged? \leftarrow \text{true}$

for each state s **in** S **do**

if $\max_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s'] > \sum_{s'} P(s' | s, \pi[s]) U[s']$ **then do**

$\pi[s] \leftarrow \operatorname{argmax}_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s']$

Update policy for state s

$unchanged? \leftarrow \text{false}$

until $unchanged?$

return π

Policy Iteration Algorithm

function POLICY-ITERATION(mdp) **returns** a policy

inputs: mdp , an MDP with states S , actions $A(s)$, transition model $P(s' | s, a)$

local variables: U , a vector of utilities for states in S , initially zero

π , a policy vector indexed by state, initially random

repeat

$U \leftarrow \text{POLICY-EVALUATION}(\pi, U, mdp)$

$unchanged? \leftarrow \text{true}$

for each state s **in** S **do**

if $\max_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s'] > \sum_{s'} P(s' | s, \pi[s]) U[s']$ **then do**

$\pi[s] \leftarrow \operatorname{argmax}_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s']$

$unchanged? \leftarrow \text{false}$

Policy updated -> remember!

until $unchanged?$

return π

Reinforcement Learning

Main Machine Learning Categories

Supervised learning

Supervised learning is one of the most common techniques in machine learning. It is based on **known relationship(s) and patterns within data** (for example: relationship between inputs and outputs).

Frequently used types: **regression**, and **classification**.

Unsupervised learning

Unsupervised learning involves finding underlying patterns within data. Typically used in **clustering** data points (similar customers, etc.)

Reinforcement learning

Reinforcement learning is inspired by behavioral psychology. It is **based on a rewarding / punishing an algorithm**.



Rewards and punishments are based on algorithm's action within its environment.

Markov Decision Process

- An MDP is defined by:

- A set of states $s \in S$
- A set of actions $a \in A$
- A transition function $T(s, a, s')$
 - Prob that a from s leads to s'
 - i.e., $P(s' | s, a)$
 - Also called the model
- A reward function $R(s, a, s')$
 - Sometimes just $R(s)$ or $R(s')$
- A start state (or distribution)
- Maybe a terminal state

Model





Reinforcement Learning


- An MDP is defined by:

- A set of states $s \in S$
- A set of actions $a \in A$
- A transition function $T(s, a, s')$
 - Prob that a from s leads to s'
 - i.e., $P(s' | s, a)$
 - Also called the model
- A reward function $R(s, a, s')$
 - Sometimes just $R(s)$ or $R(s')$
- A start state (or distribution)
- Maybe a terminal state

UNKNOWN



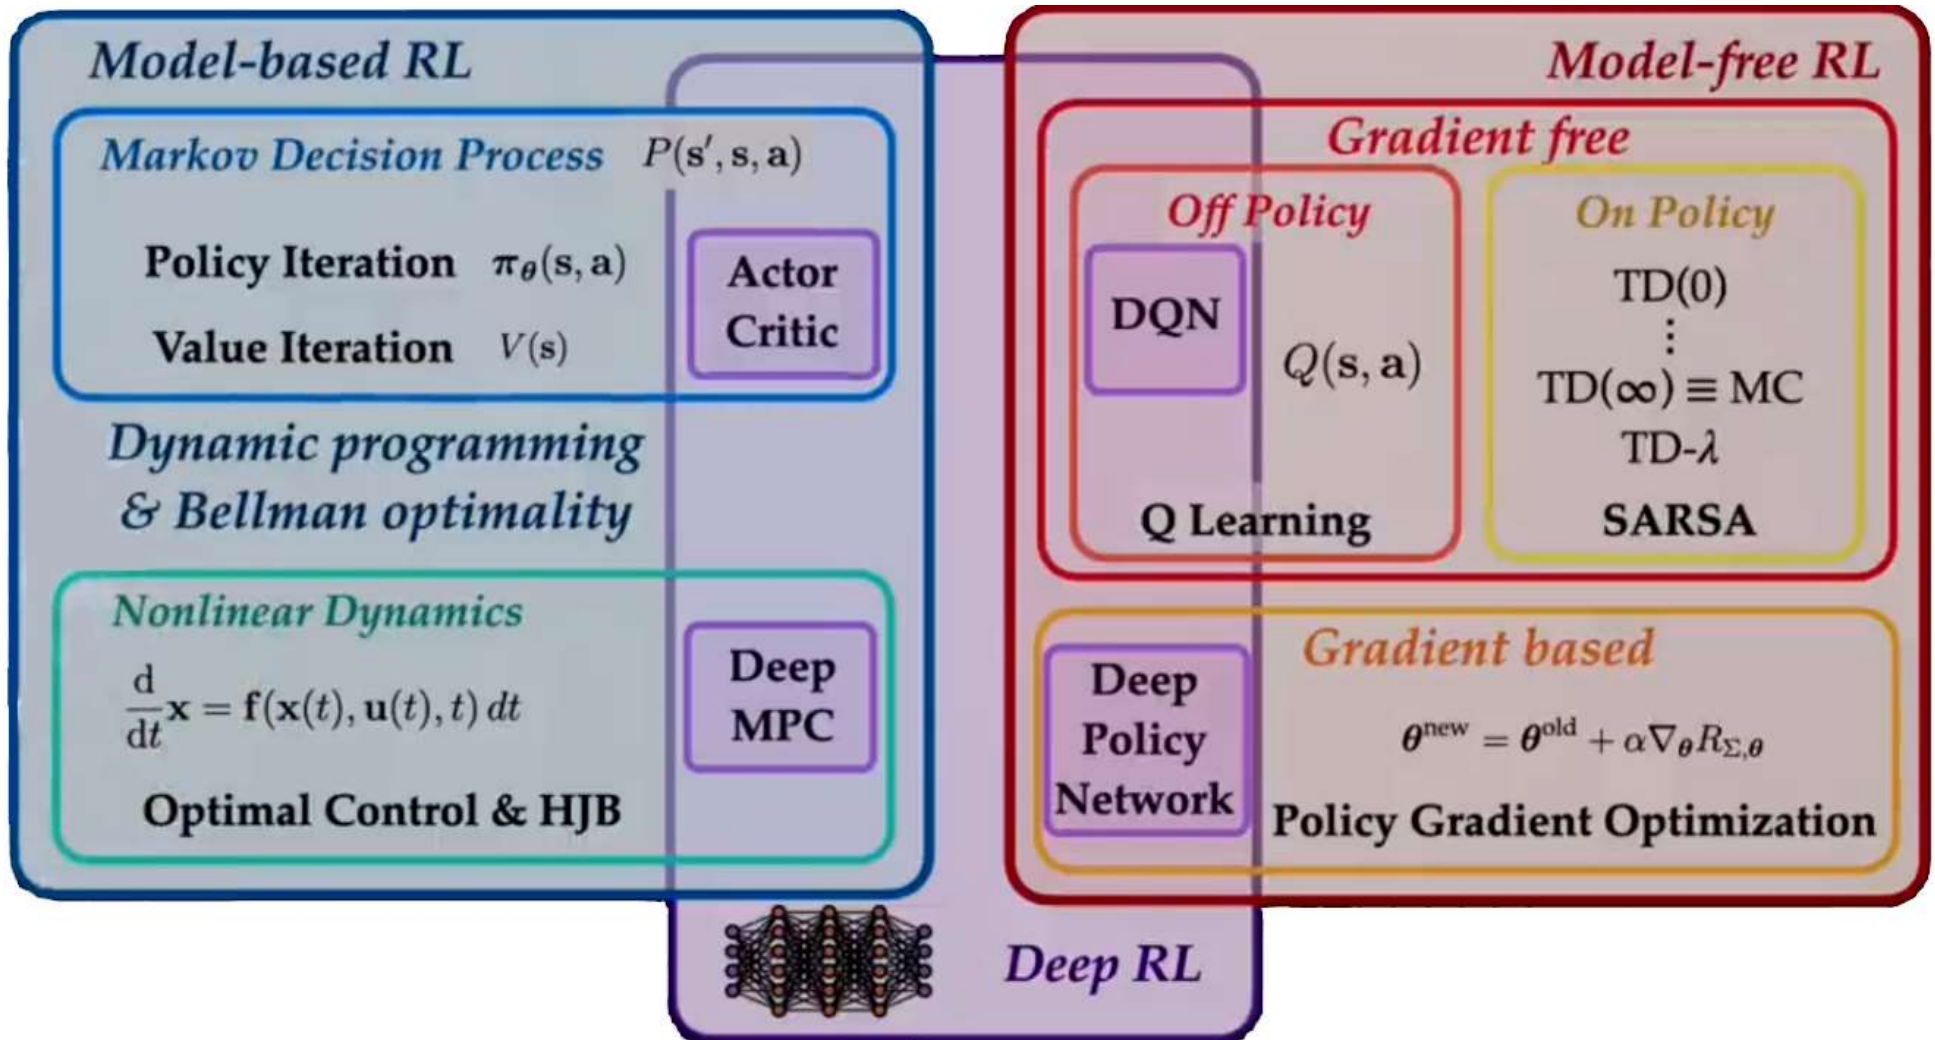
RL: Agents and Environments



MDPs vs. Reinforcement Learning

- MDPs are building blocks for RL
- RL has the additional complexity that the agent does not have access to the full specification of the MDP. For example:
 - Transition probabilities are often unknown
 - Reward function is often unknown

Reinforcement Learning Approaches



HJB: Hamilton-Jacobi-Bellman

RL: Prediction and Control

- **Prediction**
 - Given a policy, estimate the utility/value function
- **Control**
 - Learn the optimal policy

Model-free vs. Model-based

- **Model-free:**
 - The agent does not have and does not learn a model of the how the environment works
- **Model-based:**
 - The agent learns/improves a model of the environment
- **Note: we are not talking about an approximate “model” of a state representation**
 - rather, we mean model of the environment, such as transition probabilities

RL: On-Policy and Off-Policy

- **On-Policy RL:** the agent consistently follows its current policy while exploring the environment (even if suboptimal)
 - SARSA
- **Off-Policy RL:** on the other hand, allows the agent to deviate from its current policy and try different actions (even if suboptimal)
 - Q-Learning

RL: On-Policy and Off-Policy

- **On-Policy RL:** The behavior/experience is generated by the same policy π that we are trying to improve
-
- **Off-Policy RL:** The behavior/experience is generated by a behavior policy b and we are trying to learn/improve policy π

Passive vs. Active RL

- **Passive Reinforcement Learning**

- agent policy π is known and fixed
 - = agent knows which action to pick NOW
- learning state utilities $U(s)$
 - possibly environment model (transition function, reward function, etc.) as well

- **Active Reinforcement Learning**

- agent has learn what to do as well

Approaches

- **Prediction**
 - Monte Carlo methods
 - Temporal-difference learning, specifically TD(0)
 - Unified view: TD(λ)
- **Control**
 - Monte Carlo methods
 - Temporal-difference learning: SARSA, N-step TD, TD(λ)
 - Q-learning
- **Approximate methods**
 - MC prediction
 - TD prediction
 - Semi-gradient SARSA control

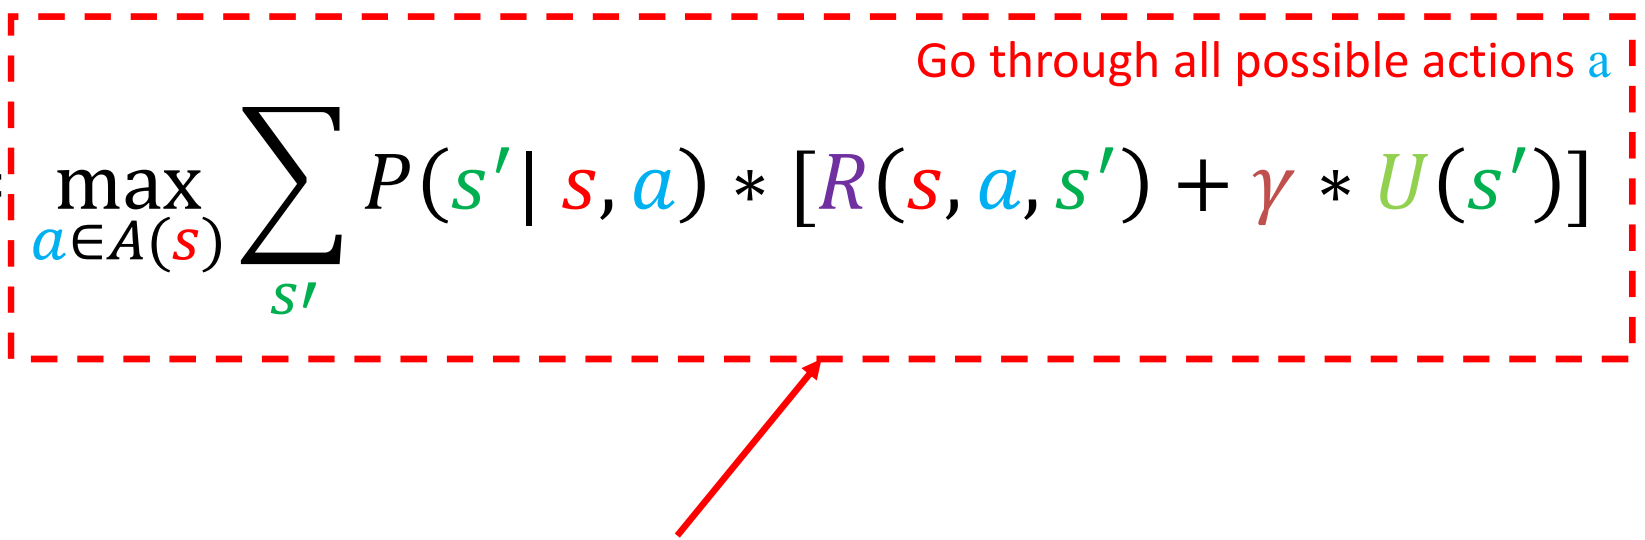
Q - Learning

Bellman Equation

The **utility** of a **state** is the expected **reward** for the next transition plus the discounted **utility** of the **next state**, assuming that the agent chooses the optimal action:

$$U(s) = \max_{a \in A(s)} \sum_{s'} P(s' | s, a) * [R(s, a, s') + \gamma * U(s')]$$

Go through all possible actions a



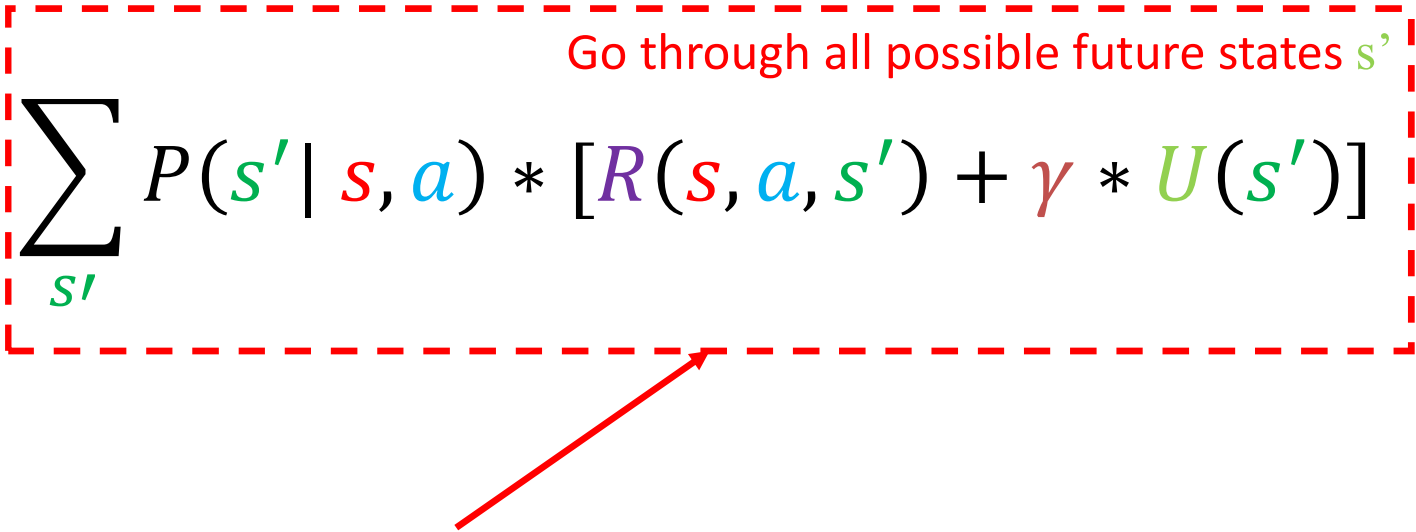
BEST Expected “long-term” **utility** / value after applying ONE specific BEST **action** a [Need Environment Model!]

Bellman Equation

The **utility** of a **state** is the expected reward for the next transition plus the discounted utility of the **next state**, assuming that the agent chooses the optimal action:

$$U(s) = \max_{a \in A(s)} \sum_{s'} P(s' | s, a) * [R(s, a, s') + \gamma * U(s')]$$

Go through all possible future states s'



Expected “long-term” **utility** / value after applying ONE specific **action** a [Need Environment Model!]

Notation

- $P(s' | s, a)$ Probability of arriving at state s' given we are at state s and take action a
- $R(s, a, s')$ The reward the agent receives when it transitions from state s to state s' via action a
- $\pi(s)$ The action recommended by policy π at state s
- π^* Optimal policy
- $U^\pi(s)$ The expected utility obtained via executing policy π starting at state s
- $U^{\pi^*}(s)$ is often abbreviated as $U(s)$
- $Q(s, a)$ expected utility of taking action a at state s
- γ Discount factor $[0, 1]$

Bellman Equation

The **utility** of a **state** is the expected reward for the next transition plus the discounted utility of the **next state**, assuming that the agent chooses the optimal action:

$$U(s) = \max_{a \in A(s)} \sum_{s'} P(s' | s, a) * [R(s, a, s') + \gamma * U(s')]$$


Go through all possible future states s'

$Q(s, a)$

Quality-function: expected utility of taking action a at state s

Utility of State and Q-Function


$$U(s) = \max_{a \in A(s)} \sum_{s'} P(s' | s, a) * [R(s, a, s') + \gamma * U(s')]$$




$$U(s) = \max_{a \in A(s)} Q(s, a)$$

Q-Function/Utility of State/Policy

$$U(s) = \max_{a \in A(s)} \sum_{s'} P(s' | s, a) * [R(s, a, s') + \gamma * U(s')]$$



$$U(s) = \max_{a \in A(s)} Q(s, a)$$



$$\pi_s^* = \operatorname{argmax}_{a \in A(s)} Q(s, a)$$

I DON'T NEED
to know the
TRANSITION
function.

NO need to know
what the
next state
s' is

Q-Learning Agent

function Q-LEARNING-AGENT(*percept*) **returns** an action

inputs: *percept*, a percept indicating the current state s' and reward signal r'

persistent: Q , a table of action values indexed by state and action, initially zero

N_{sa} , a table of frequencies for state–action pairs, initially zero

s, a, r , the previous state, action, and reward, initially null

if TERMINAL?(s) **then** $Q[s, \text{None}] \leftarrow r'$

if s is not null **then**

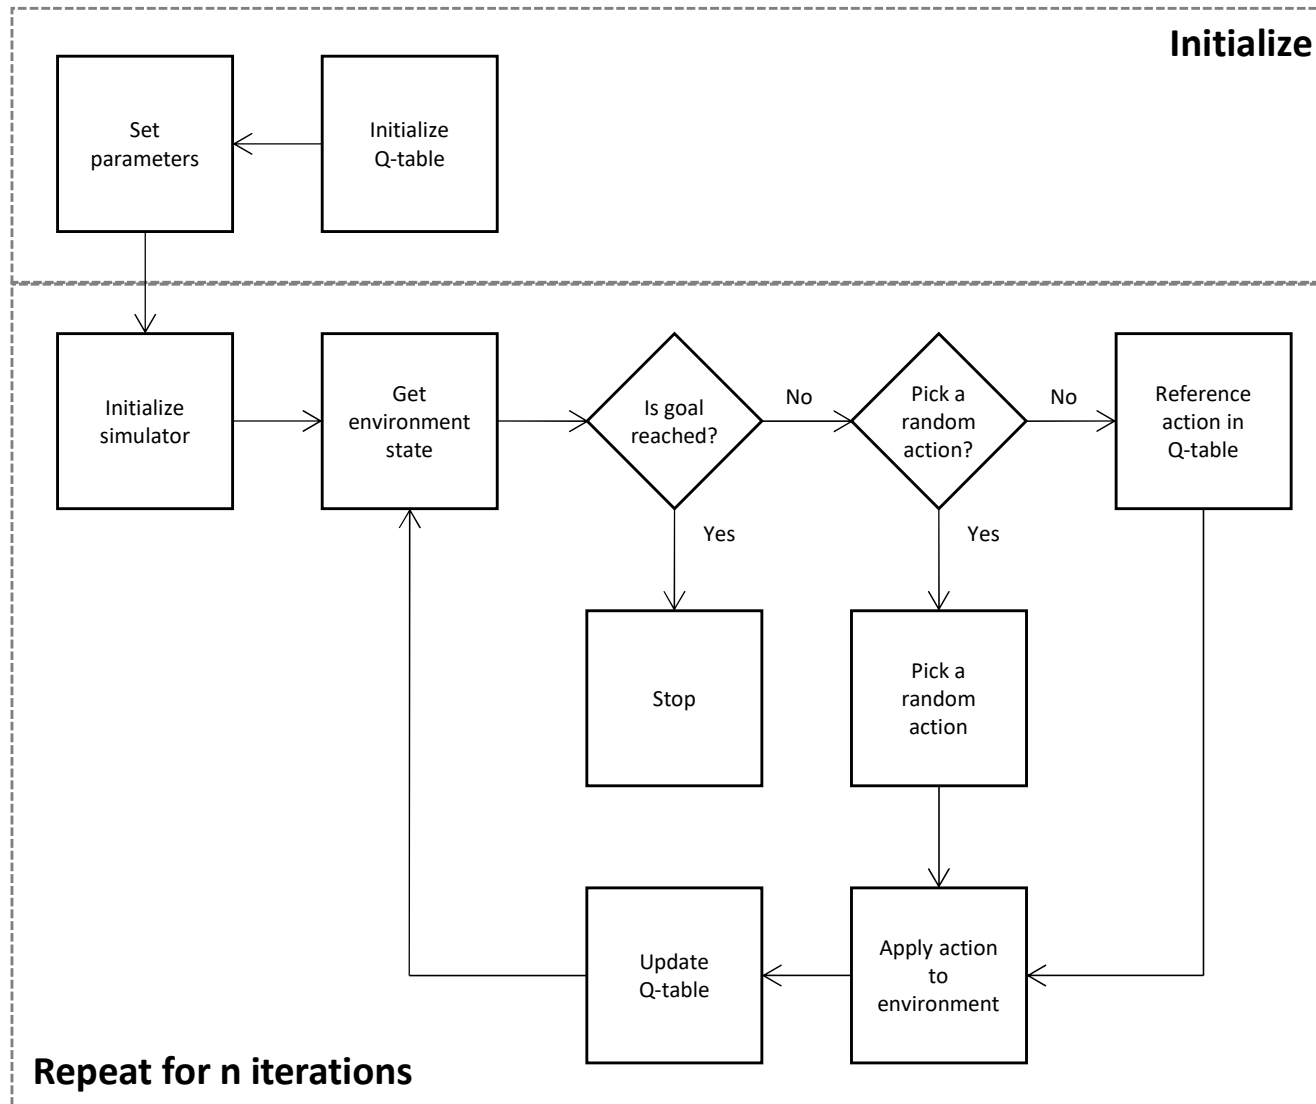
 increment $N_{sa}[s, a]$

$Q[s, a] \leftarrow Q[s, a] + \alpha(N_{sa}[s, a])(r + \gamma \max_{a'} Q[s', a'] - Q[s, a])$

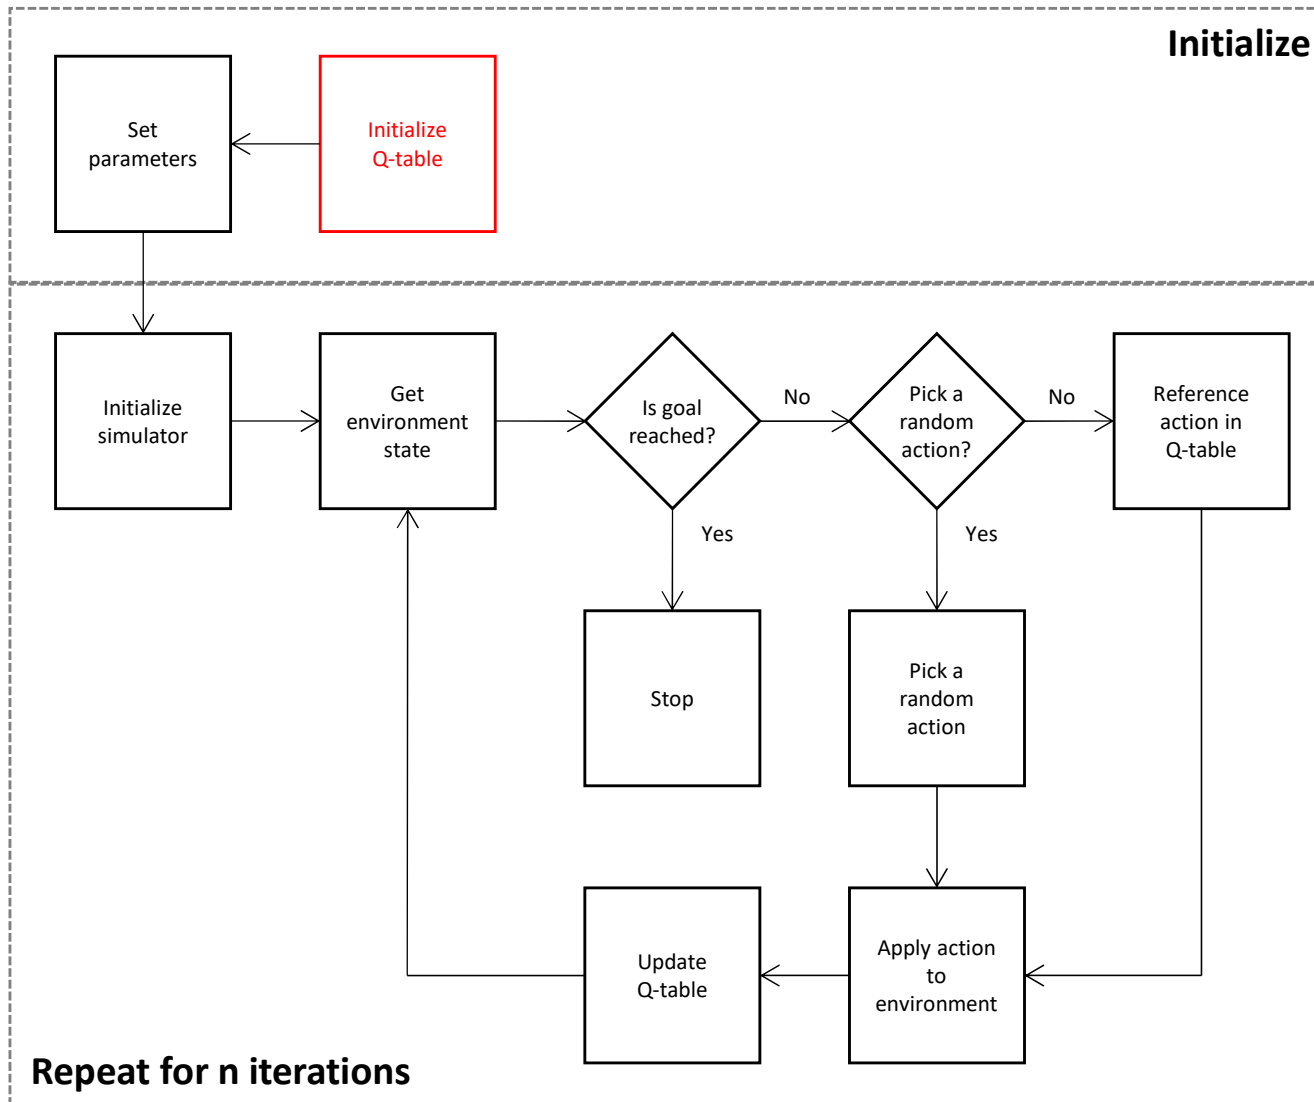
$s, a, r \leftarrow s', \operatorname{argmax}_{a'} f(Q[s', a'], N_{sa}[s', a']), r'$

return a

Q-Learning Algorithm



Q-Learning Algorithm



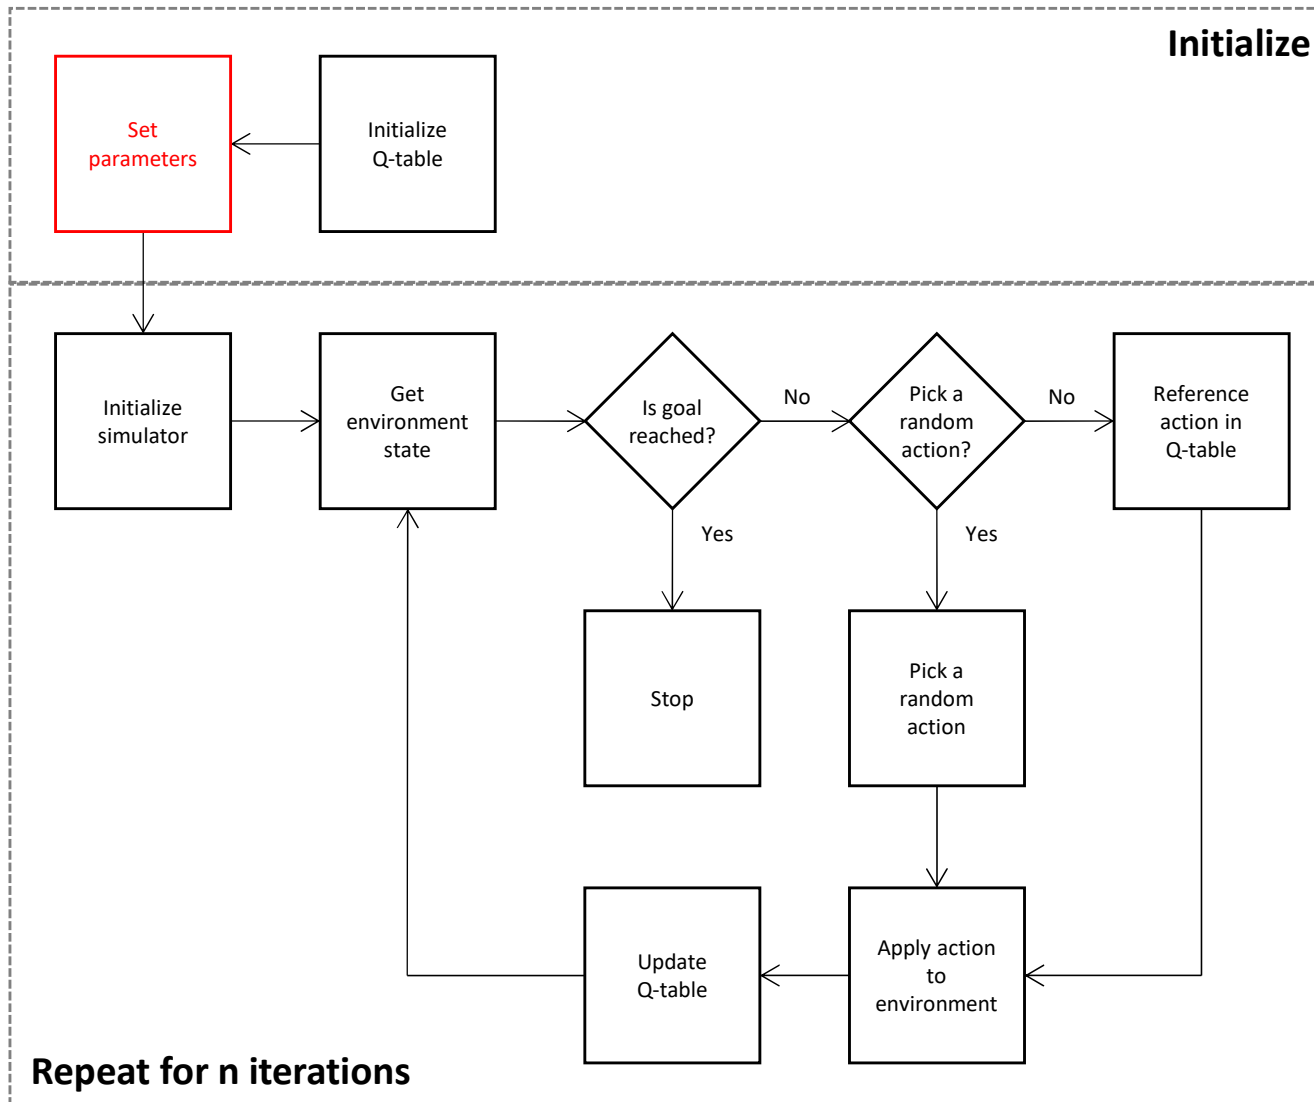
Initialize Q-table:

Set up and initialize (all values set to 0) a table where:

- rows represent **possible states**
- columns represent **actions**

Note that additional states can be added to the table when encountered.

Q-Learning Algorithm



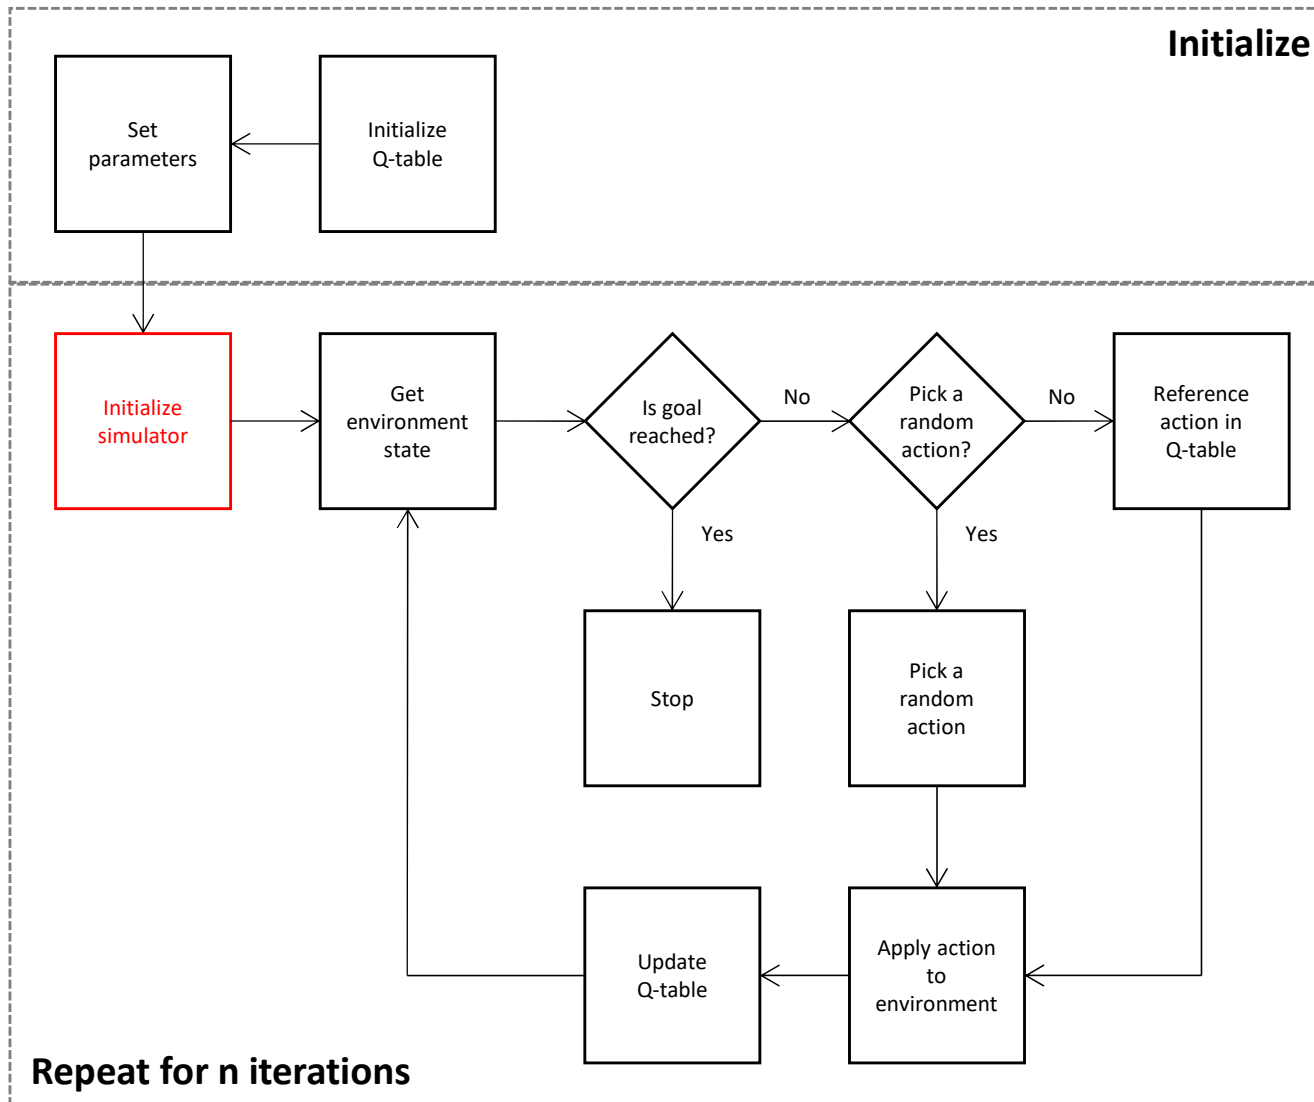
Set parameters:

Set and initialize **hyperparameters** for the Q-learning process.

Hyperparameters include:

- **chance of choosing a random action:** a threshold for **choosing a random action over an action from the Q-table**
- **learning rate:** a parameter that describes **how quickly the algorithm should learn from rewards** in different states
 - high: faster learning with erratic Q-table changes
 - low: gradual learning with possibly more iterations
- **discount factor:** a parameter that **describes how valuable are future rewards**. It tells the algorithm whether it should seek “immediate gratification” (small) or “long-term reward” (large)

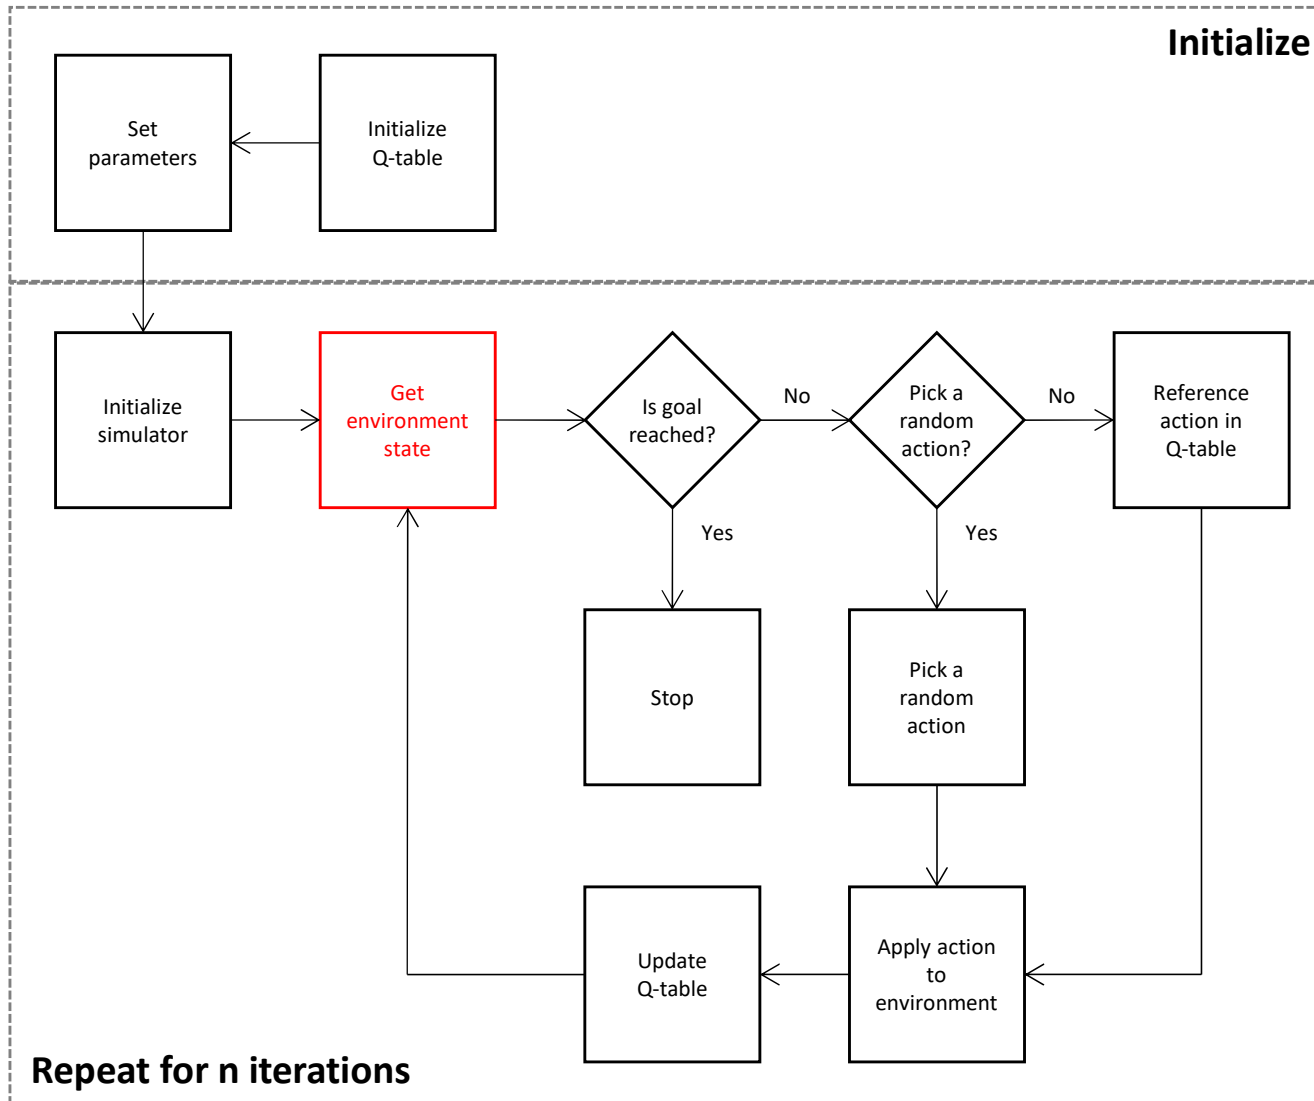
Q-Learning Algorithm



Initialize simulator:

Reset the simulated environment to its initial state and place the agent in a neutral state.

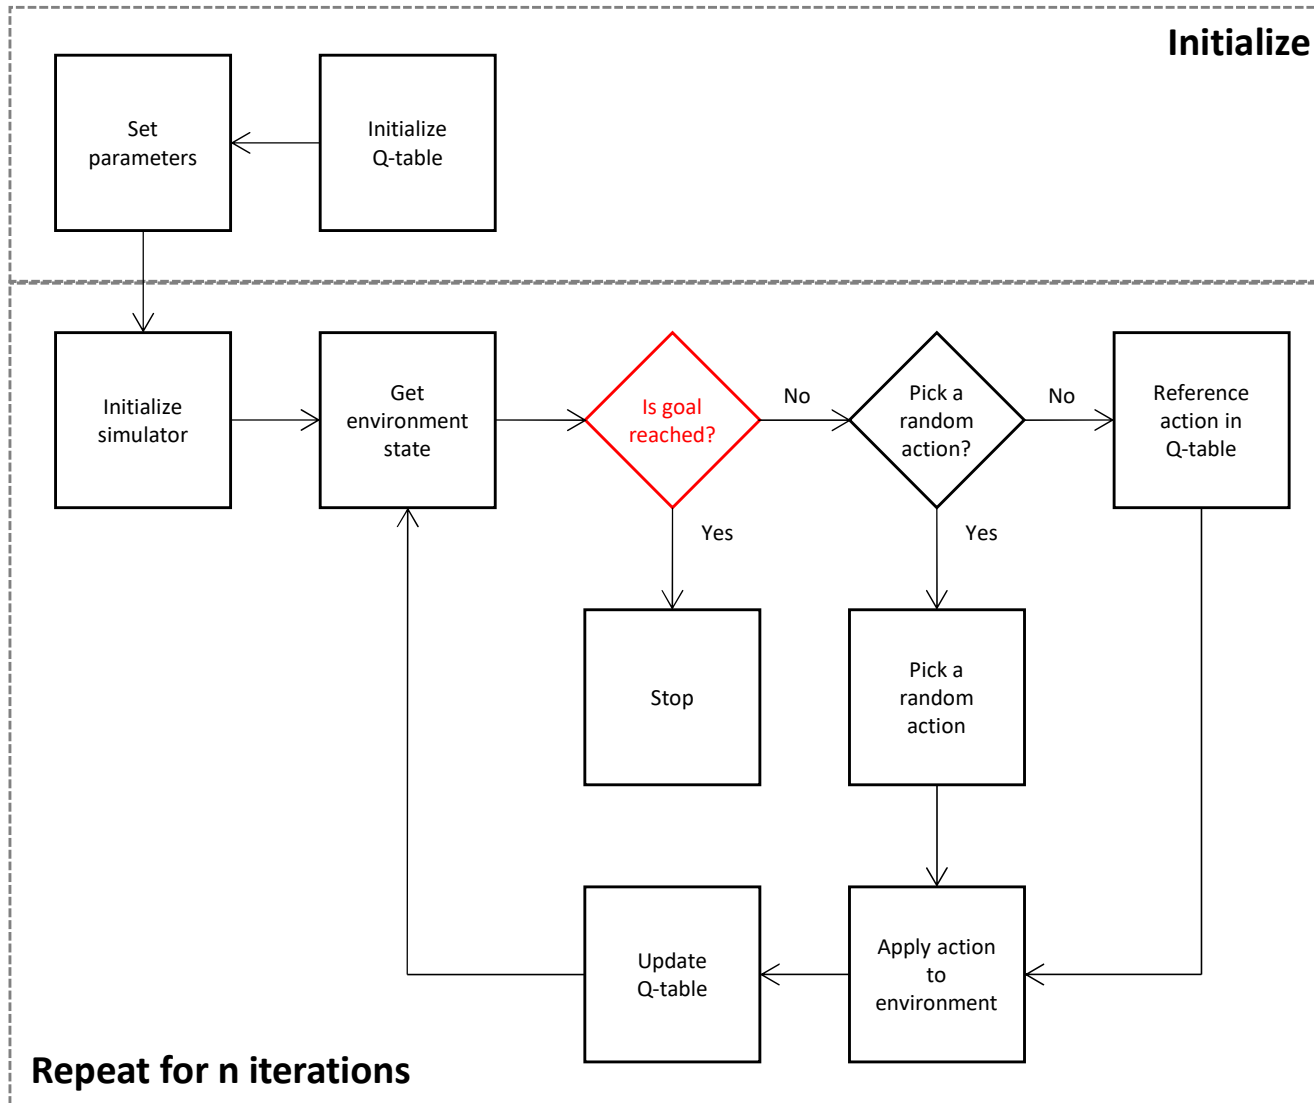
Q-Learning Algorithm



Get environment state:

Report the current state of the environment. Typically a vector of values representing all relevant variables.

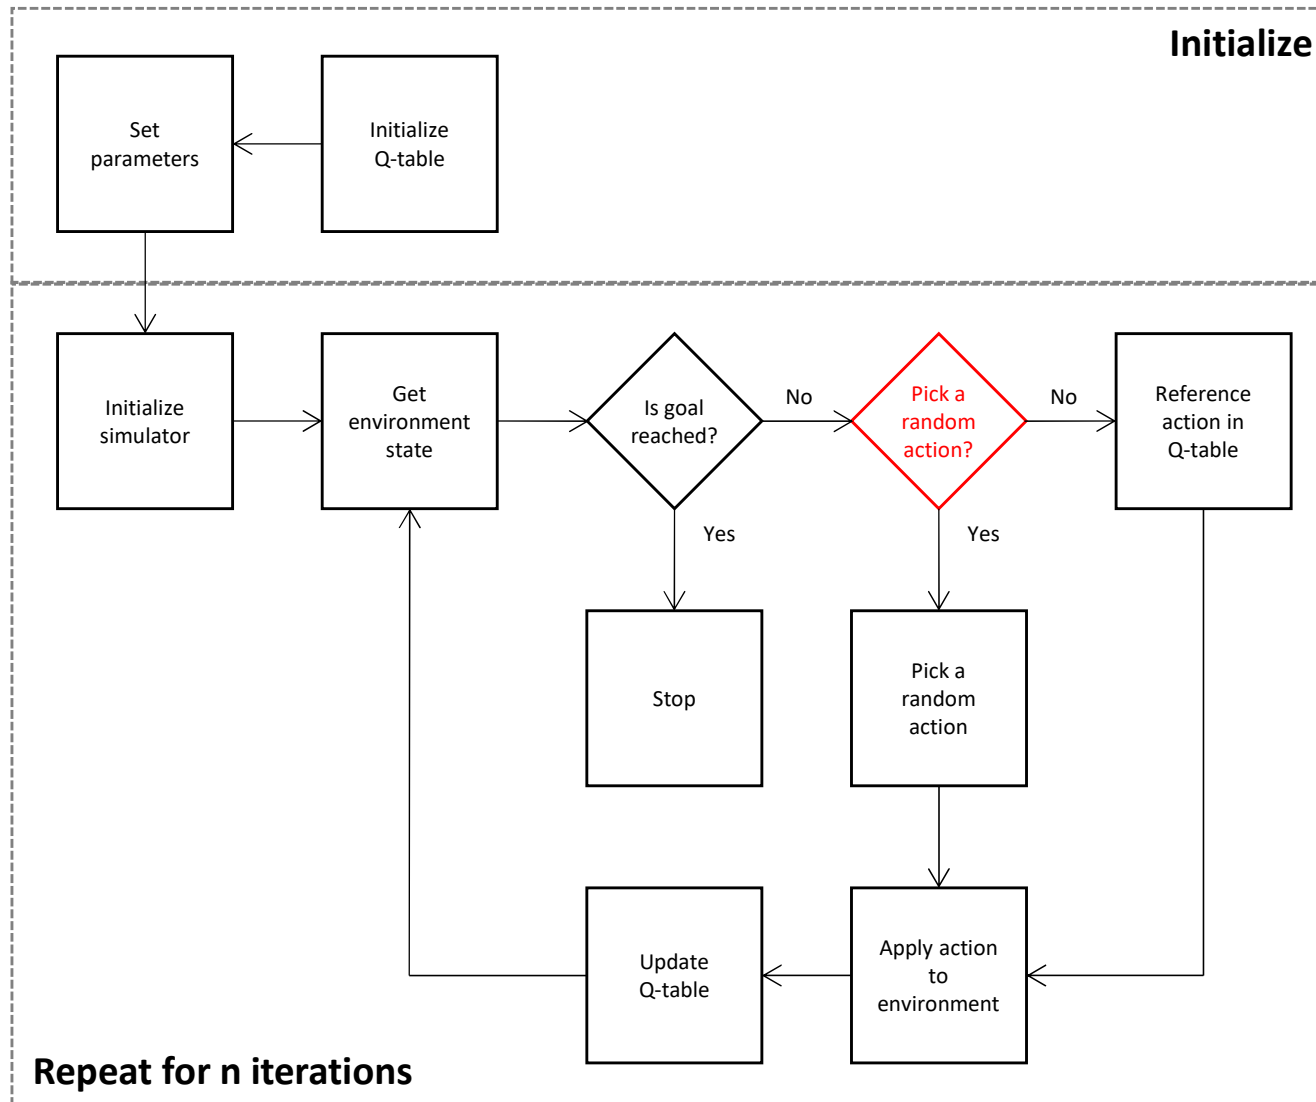
Q-Learning Algorithm



Is goal reached?:

Verify if the goal of the simulation has been achieved. It could be decided with the agent arriving in expected final state or by some simulation parameter.

Q-Learning Algorithm

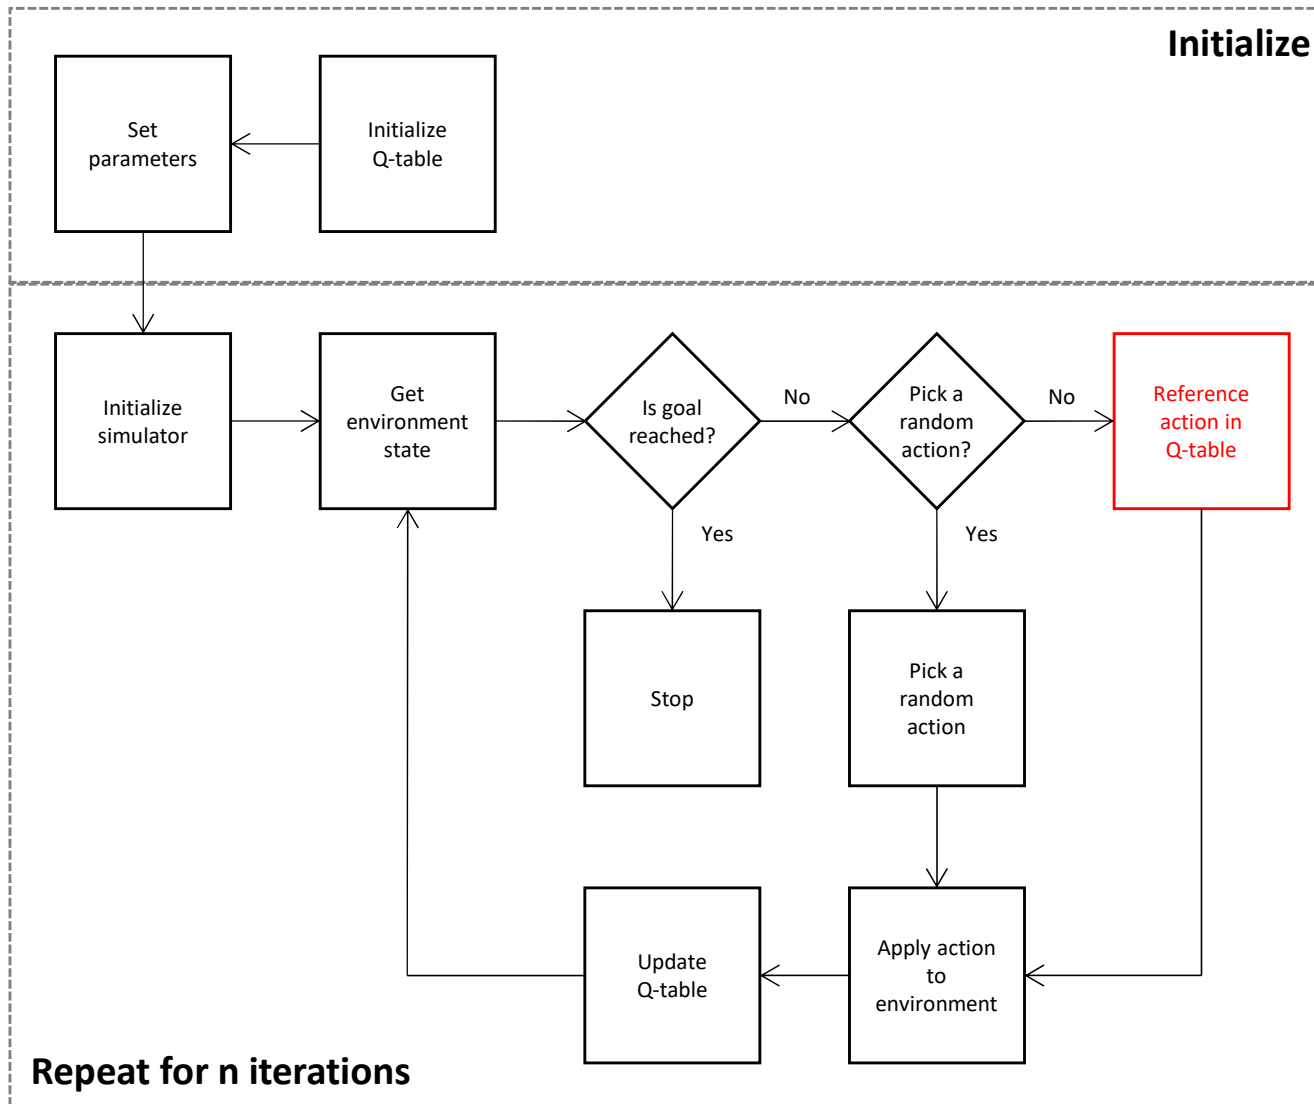


Pick a random action?:

Decide whether next action should be picked at random or not (it will be selected based on Q-table data then).

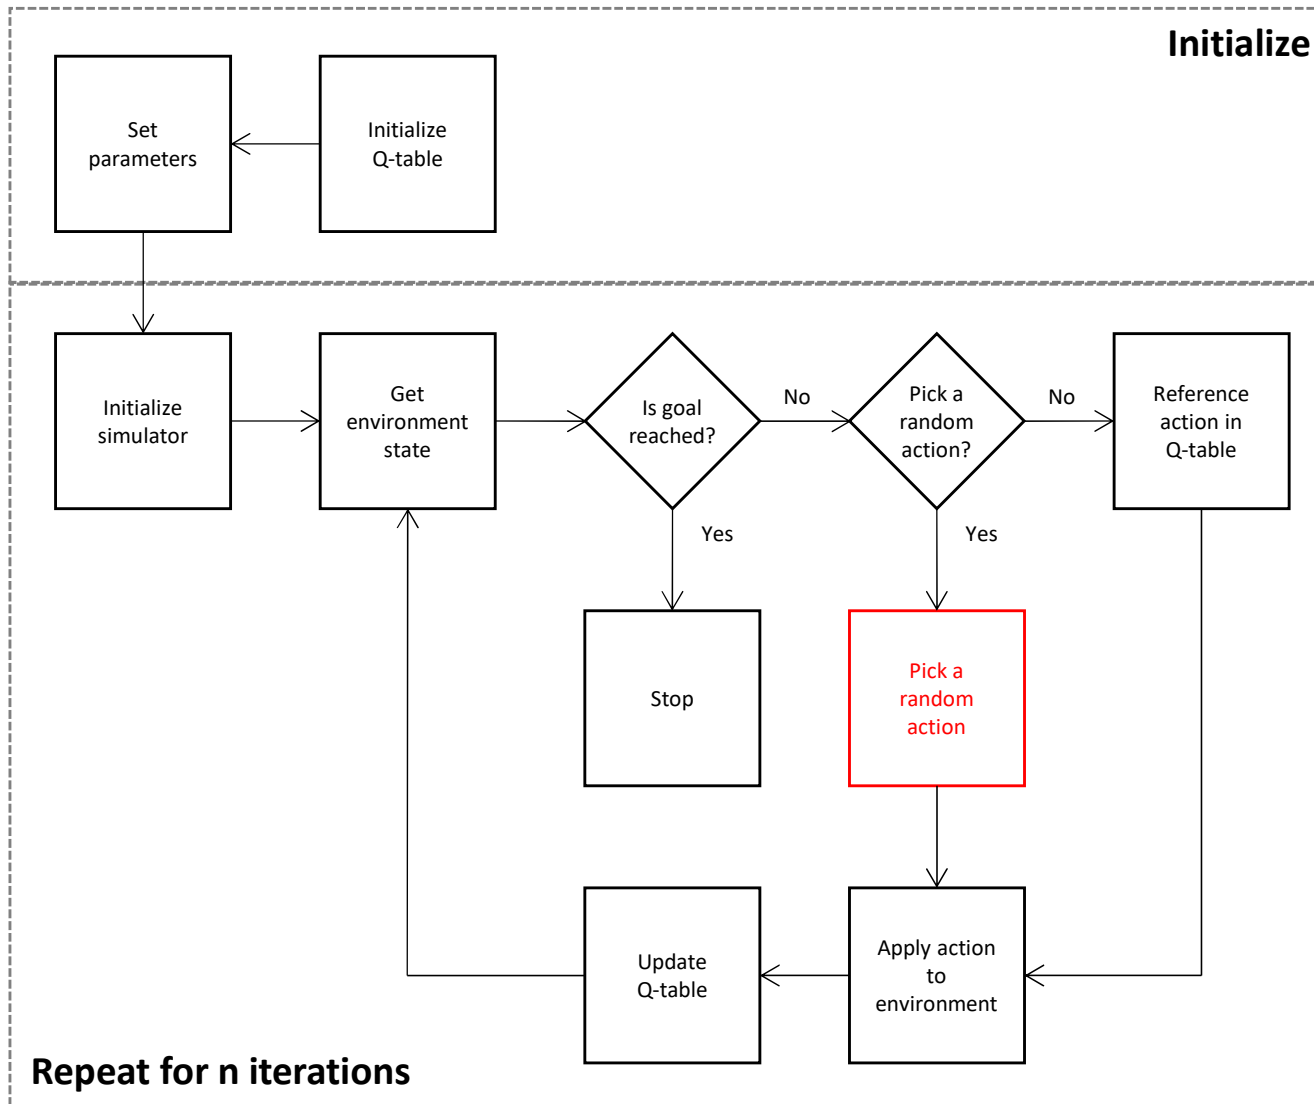
Use the **chance of choosing a random action hyperparameter** to decide.

Q-Learning Algorithm



Reference action in Q-table:
Next action decision will be based on data from the Q-table **given the current state of the environment.**

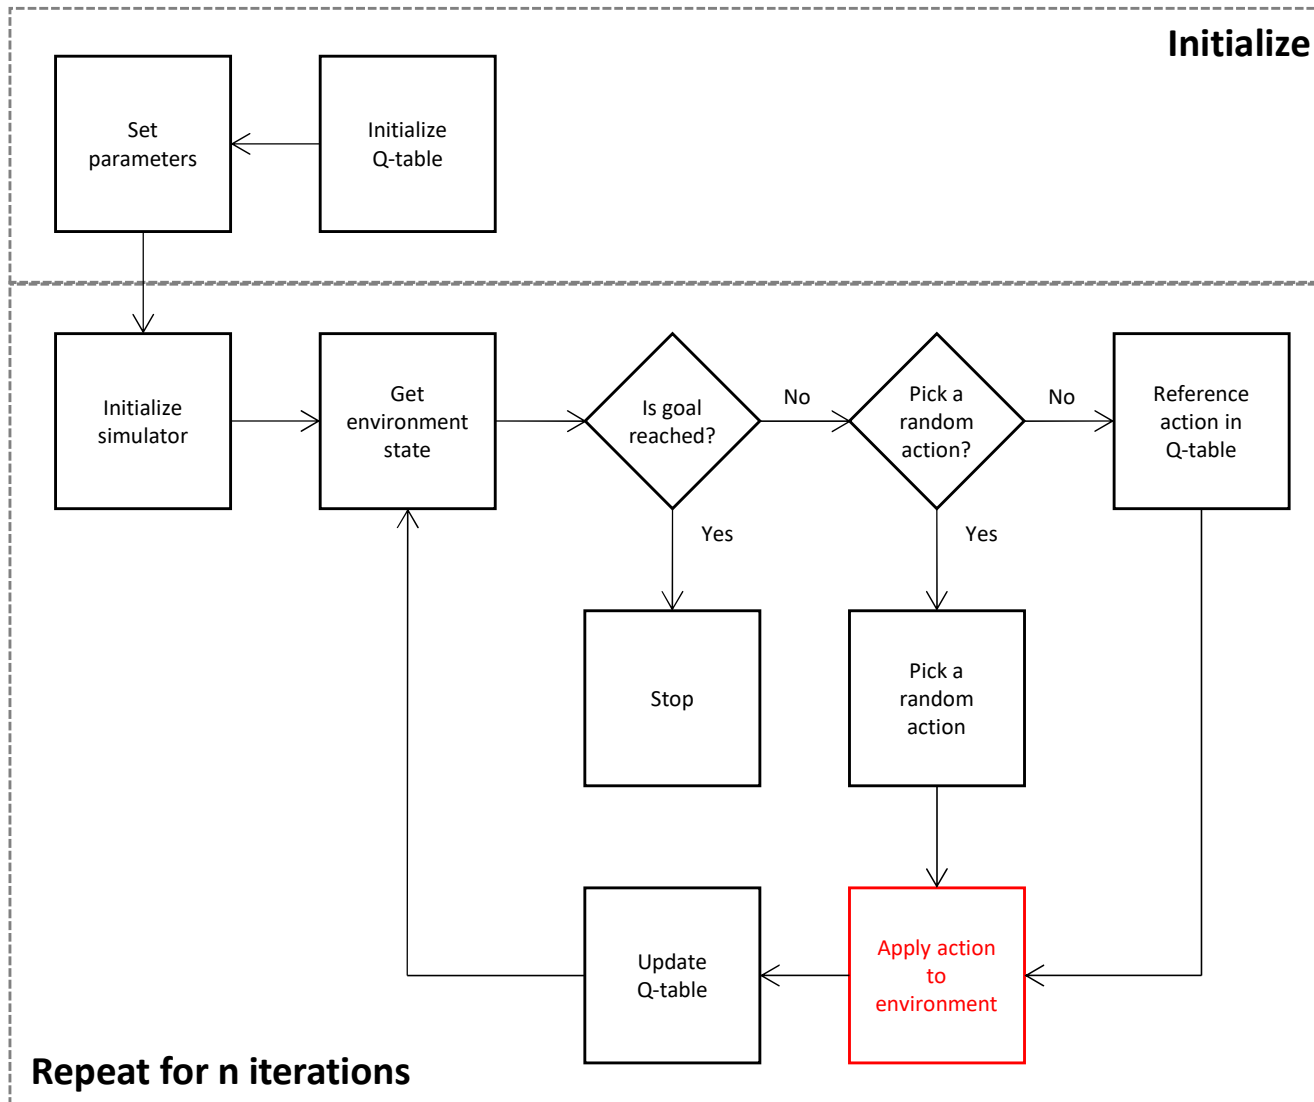
Q-Learning Algorithm



Pick a random action:

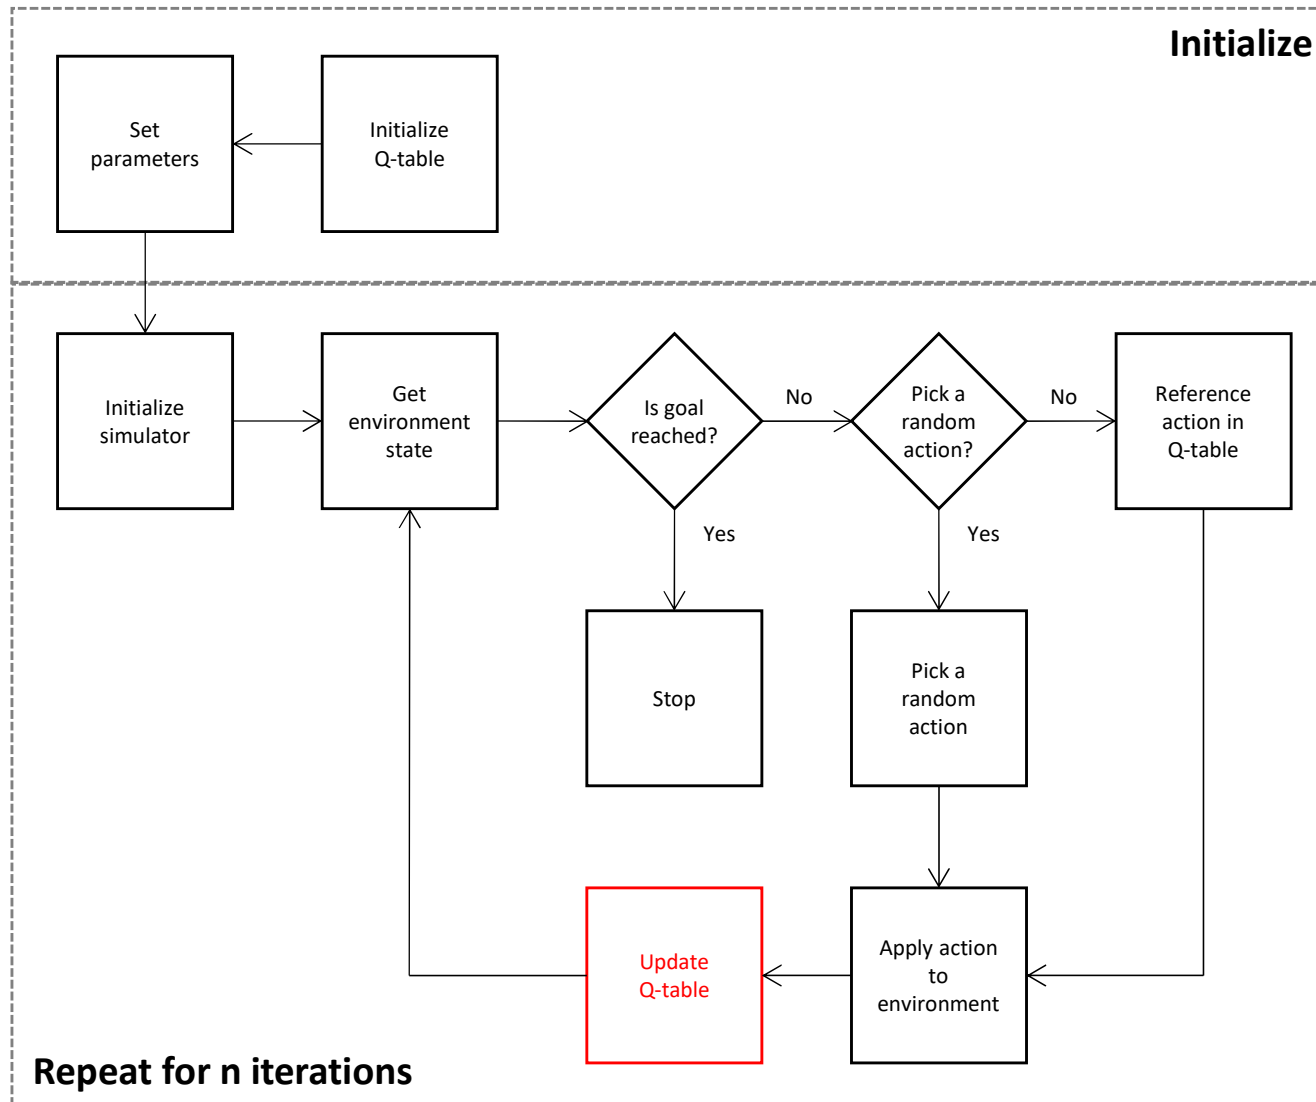
Pick any of the available actions at random. Helpful with exploration of the environment.

Q-Learning Algorithm



Apply action to environment:
Apply the action to the environment to change it. Each action will have its own reward.

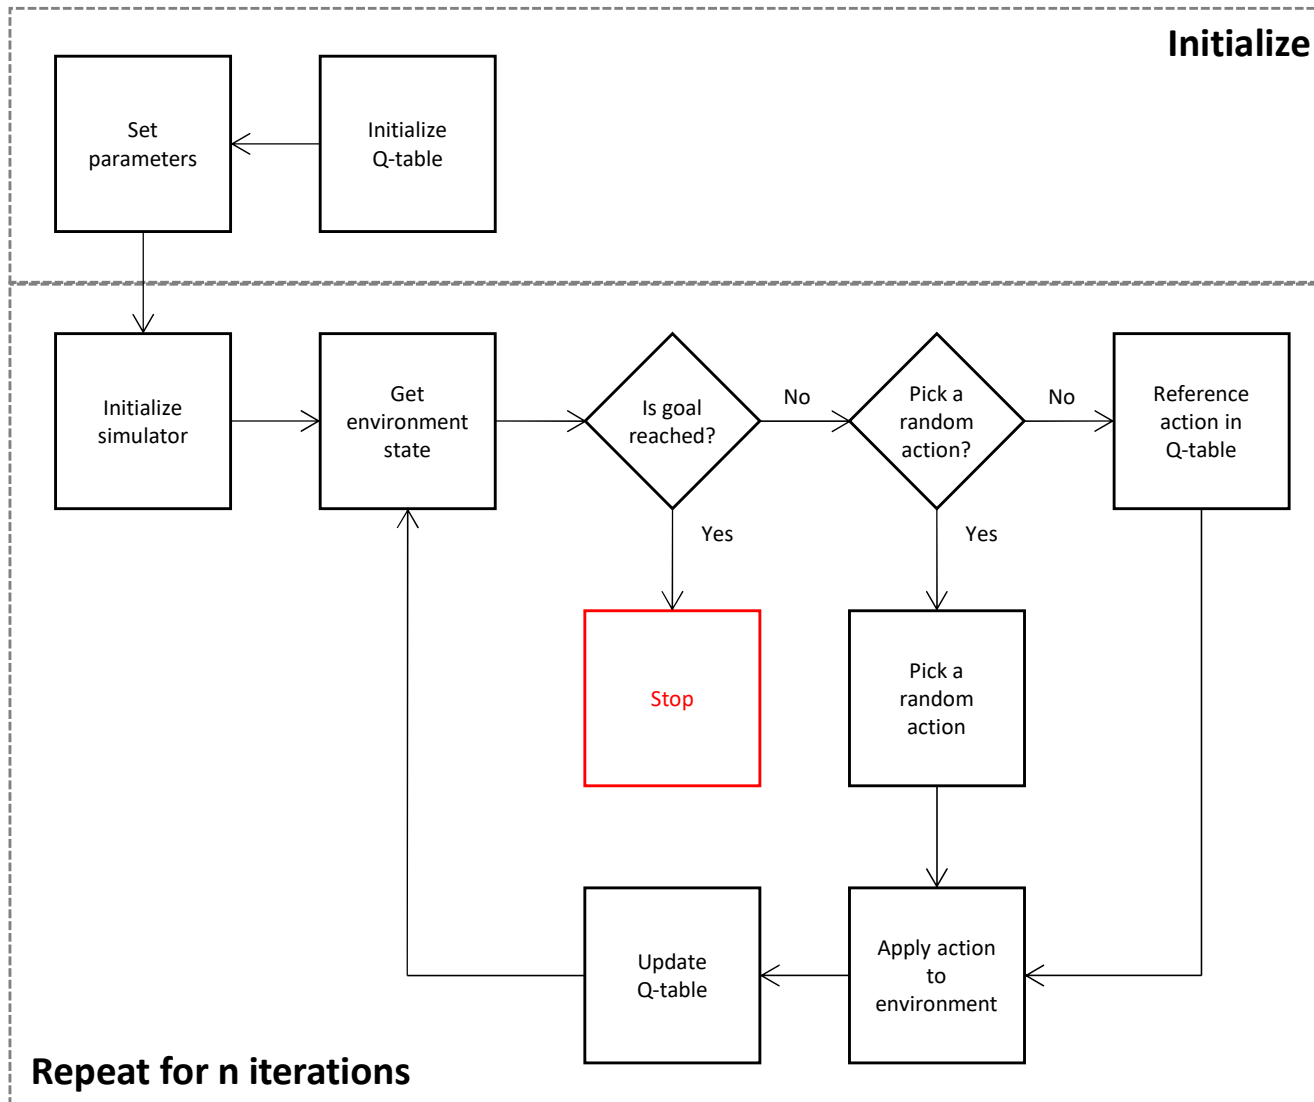
Q-Learning Algorithm



Update Q-table:

Update the Q-table given the reward resulting from recently applied action (feedback from the environment).

Q-Learning Algorithm



Stop:
Stop the learning process

Q – Learning: Example

Q-Learning Agent

function Q-LEARNING-AGENT(*percept*) **returns** an action

inputs: *percept*, a percept indicating the current state s' and reward signal r'

persistent: Q , a table of action values indexed by state and action, initially zero

N_{sa} , a table of frequencies for state–action pairs, initially zero

s, a, r , the previous state, action, and reward, initially null

if TERMINAL?(s) **then** $Q[s, \text{None}] \leftarrow r'$

if s is not null **then**

increment $N_{sa}[s, a]$


ASSUME: 1

$Q[s, a] \leftarrow Q[s, a] + \alpha(N_{sa}[s, a])(r + \gamma \max_{a'} Q[s', a'] - Q[s, a])$

$s, a, r \leftarrow s', \operatorname{argmax}_{a'} f(Q[s', a'], N_{sa}[s', a']), r'$

return a

Q-Learning Algorithm



Q-table		Actions			
		↑	↓	→	←
States	1	0	0	0	0
	2	0	0	0	0

	n	0	0	0	0

Rewards:

Move into car: -100

Move into pedestrian: -1000

Move into empty space: 100

Move into goal: 500

Action:

Reward:

Q-table value:


ASSUME: 1

$$Q[s, a] \leftarrow Q[s, a] + \alpha(N_{sa}[s, a])(r + \gamma \max_{a'} Q[s', a'] - Q[s, a])$$

Discount Factor γ

- The discount factor γ is a number between 0 and 1.
- The discount factor describes the preference of an agent for current rewards over future rewards.
- When γ is close to 0, rewards in the distant future are viewed as insignificant.
- When γ is 1, discounted rewards are exactly equivalent to additive rewards, so additive rewards are a special case of discounted rewards.
- Discounting appears to be a good model of both animal and human preferences over time.

Q-Learning Algorithm



Q-table		Actions			
		↑	↓	→	←
States	1	0	0	0	0
	2	0	0	0	0

	n	0	0	0	0

Rewards:

Move into car: -100

Move into pedestrian: -1000

Move into empty space: 100

Move into goal: 500

Action:

Reward:


Q-table value:

$$Q(\text{state}, \text{action}) = (1 - \text{alpha}) * Q(\text{state}, \text{action}) + \text{alpha} * (\text{reward} + \text{gamma} * Q(\text{next state}, \text{all actions}))$$

← Learning rate
Discount

Current value
Maximum value of all actions on next state →

Q-Learning Algorithm



Q-table		Actions			
		↑	↓	→	←
States	1	0	0	0	0
	2	0	0	0	0

	n	0	0	0	0

Rewards:

Move into car: -100

Move into pedestrian: -1000

Move into empty space: 100

Move into goal: 500

Action:

Reward:


Q-table value:

$$Q(\text{state}, \text{action}) = (1 - \text{alpha}) * Q(\text{state}, \text{action}) + \text{alpha} * (\text{reward} + \text{gamma} * Q(\text{next state}, \text{all actions}))$$

Learning rate Discount

Current value Maximum value of all actions on next state

Q-Learning Algorithm



Q-table		Actions			
		↑	↓	→	←
States	1	0	0	0	0
	2	0	0	0	0

	n	0	0	0	0

Rewards:

Move into car: -100

Move into pedestrian: -1000

Move into empty space: 100

Move into goal: 500

Action:

Reward:


Q-table value:

$$Q(\text{state}, \text{action}) = (1 - \text{alpha}) * Q(\text{state}, \text{action}) + \text{alpha} * (\text{reward} + \text{gamma} * Q(\text{next state}, \text{all actions}))$$

← Learning rate
Discount

Current value
Maximum value of all actions on next state →

Q-Learning Algorithm



Action: →

Reward:   -100

Q-table value:

$$Q(1, \text{east}) = (1 - 0.1) * 0 + 0.1 * (-100 + 0.6 * \max \text{ of } Q(2, \text{all actions}))$$

Q-table		Actions			
		↑	↓	→	←
States	1	0	0	0	0
	2	0	0	0	0

	n	0	0	0	0

Rewards:


Move into car: -100

Move into pedestrian: -1000

Move into empty space: 100

Move into goal: 500

Q-Learning Algorithm



Action: →

Reward: -100

Q-table value:

$$Q(1, \text{east}) = (1 - 0.1) * 0 + 0.1 * (-100 + 0.6 * 0) = -10$$

Q-table		Actions			
		↑	↓	→	←
States	1	0	0	-10	0
	2	0	0	0	0

	n	0	0	0	0

Rewards:














































Move into car: -100

Move into pedestrian: -1000

Move into empty space: 100

Move into goal: 500

Q-Learning Algorithm

Action: →

Reward:   -1000

Q-table value:

$$Q(2, \text{south}) = (1 - 0.1) * 0 + 0.1 * (-1000 + 0.6 * \max \text{ of } Q(3, \text{all actions}))$$

Q-table		Actions			
		↑	↓	→	←
States	1	0	0	-10	0
	2	0	0	0	0

	n	0	0	0	0

Rewards:


Move into car: -100

Move into pedestrian: -1000

Move into empty space: 100

Move into goal: 500

Q-Learning Algorithm



Action: →

Reward: -1000

Q-table value:

$$Q(2, \text{south}) = (1 - 0.1) * 0 + 0.1 * (-1000 + 0.6 * 0) = -100$$

Q-table		Actions			
		↑	↓	→	←
States	1	0	0	-10	0
	2	0	-100	0	0

	n	0	0	0	0

Rewards:

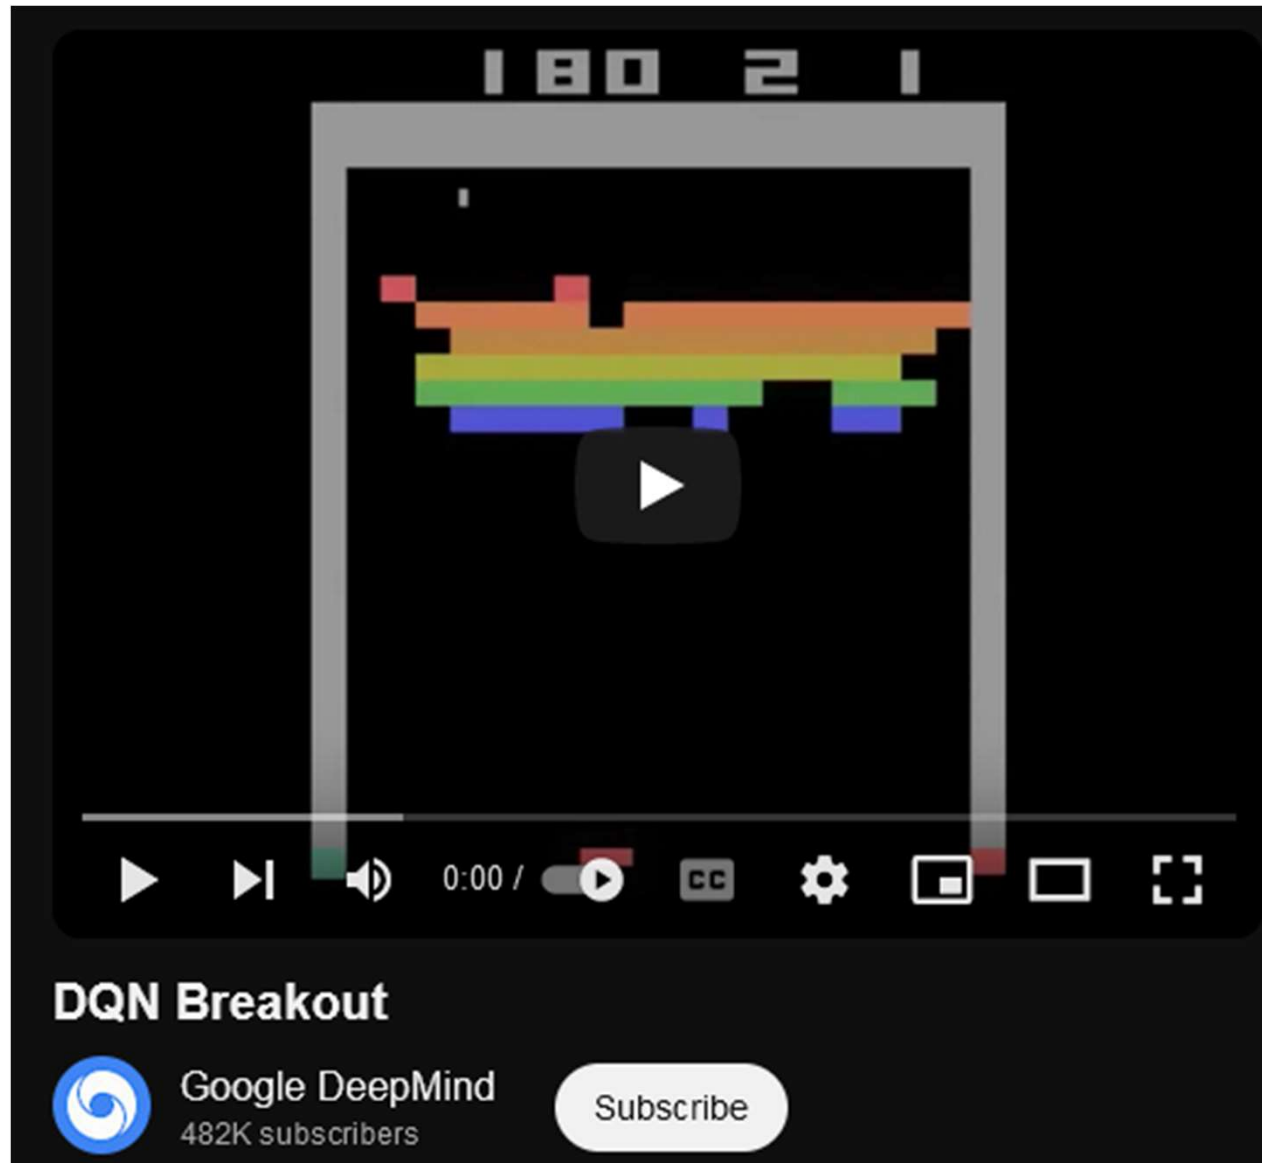
Move into car: -100

Move into pedestrian: -1000

Move into empty space: 100

Move into goal: 500

Deep Q-Learning



Source: <https://www.youtube.com/watch?v=TmPfTpjtdgg>