

CS 581

Advanced Artificial Intelligence

January 22, 2024

Announcements / Reminders

- Please follow the Week 02 To Do List instructions (if you haven't already)
- Written Assignment #01: to be posted soon
- Programming Assignment #01: to be posted soon
- **UPDATED Exam Dates:**
 - **Midterm:**
 - WAS: ~~02/28/2024 during Wednesday lecture time~~
 - IS: 02/21/2024 during Wednesday lecture time
 - **Final:**
 - WAS: ~~04/24/2024 during Wednesday lecture time~~
 - IS: 04/22/2024 during Monday lecture time

Teaching Assistants

Name	e-mail	Office hours
Gawade, Vishal	vgawade@hawk.iit.edu	TBD
Zhou, Xiaoting	xzhou70@hawk.iit.edu	TBD

TAs will:

- assist you with your assignments,
- hold office hours to answer your questions,
- grade your assignments (**a specific TA will be assigned to you**).

Take advantage of their time and knowledge!

DO NOT email them with questions unrelated to lab grading.

Make time to meet them during their office hours.

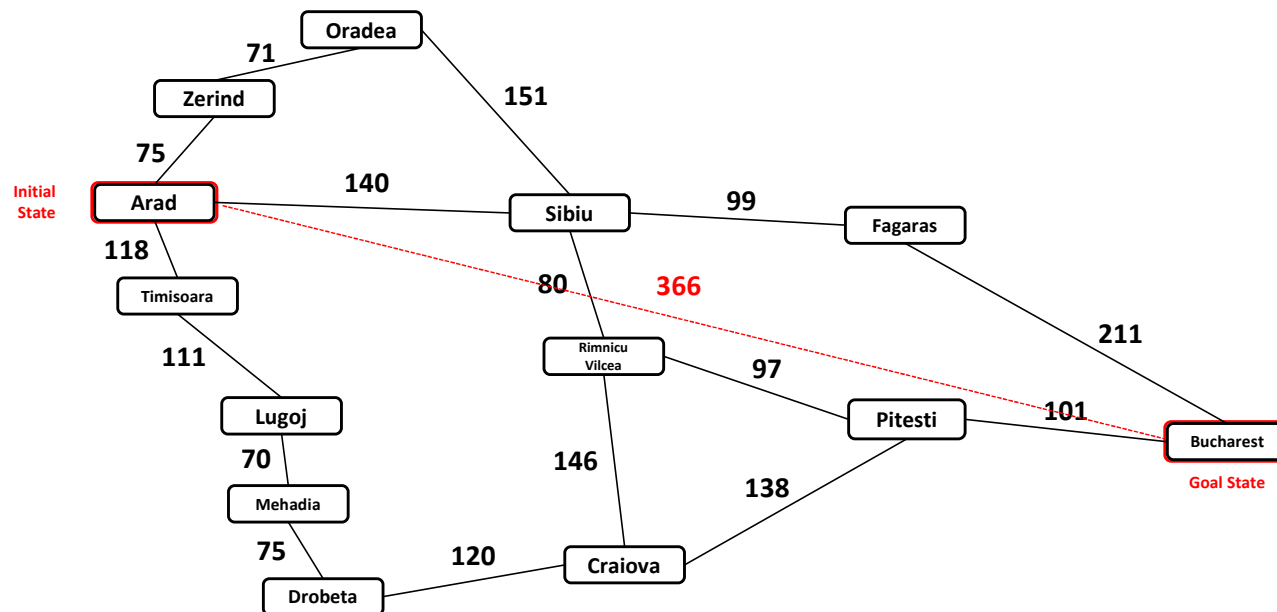
Add a [CS581 Spring 2024] prefix to your email subject when contacting TAs, please.

Plan for Today

- Solving problems by Searching
 - A* Algorithm: continued
 - Local Search Algorithms

What Made A* Work Well?

- Straight-line heuristics is **admissible**: it never overestimates the cost.



- An **admissible heuristics** is guaranteed to give you the optimal solution

A* Evaluation Function

A* Evaluation Function:

$$f(n) = g(\text{State}_n) + h(\text{State}_n)$$

where:

- $g(n)$ - initial node to node n path cost
- $h(n)$ - **estimated cost** of the best path that continues from node n to a goal node

So:

$$f(n) < C^*$$

where:

- C^* - **optimal** path cost

Admissible Heuristic: Proof

An **admissible heuristic** $h()$ is guaranteed to give you the optimal solution. Why? Proof by contradiction:

- Say: the algorithm returned a suboptimal path ($C > C^*$)
- So: there exist a node n on C^* not expanded on C :

If so: $f(n) > C^*$

$$f(n) = g(n) + h(n) \quad (\text{by definition})$$

$$f(n) = g^*(n) + h(n) \quad (\text{because } n \text{ is on } C^*)$$


$$f(n) \leq g^*(n) + h^*(n) \quad (\text{if } h(n) \text{ admissible: } h(n) \leq h^*(n))$$

But that would mean that:

$$f(n) \leq C^* \quad (\text{contradiction!})$$

What Made A* Work Well?

- Straight-line heuristics is **consistent**: its estimate is getting better and better as we get closer to the goal




- Every **consistent** heuristics is **admissible heuristics**, but not the other way around

A*: Search Contours

How does A* “direct” the search progress?


$$f(n) = g(\text{State}_n) + h(\text{State}_n)$$



A*: SURELY Expanded Nodes

Nodes **INSIDE** contours (ellipsoids) have:


$$f(n) = g(\text{State}_n) + h(\text{State}_n) < \text{contour value}$$



A*: POSSIBLY Expanded Nodes

Nodes **ON** contours (ellipsoids) have:


$$f(n) = g(\text{State}_n) + h(\text{State}_n) = \text{contour value}$$



A*: NOT Expanded Nodes


Nodes **OUTSIDE** contours (ellipsoids) have:

$$f(n) = g(\text{State}_n) + h(\text{State}_n) > \text{contour value}$$



Dominating Heuristics

We can have more than one available heuristics.
For example $h_1(n)$ and $h_2(n)$.




$h_2(n)$ dominates $h_1(n)$ if and only if

$$h_2(n) > h_1(n) \text{ for every } n$$

Dominating Heuristics

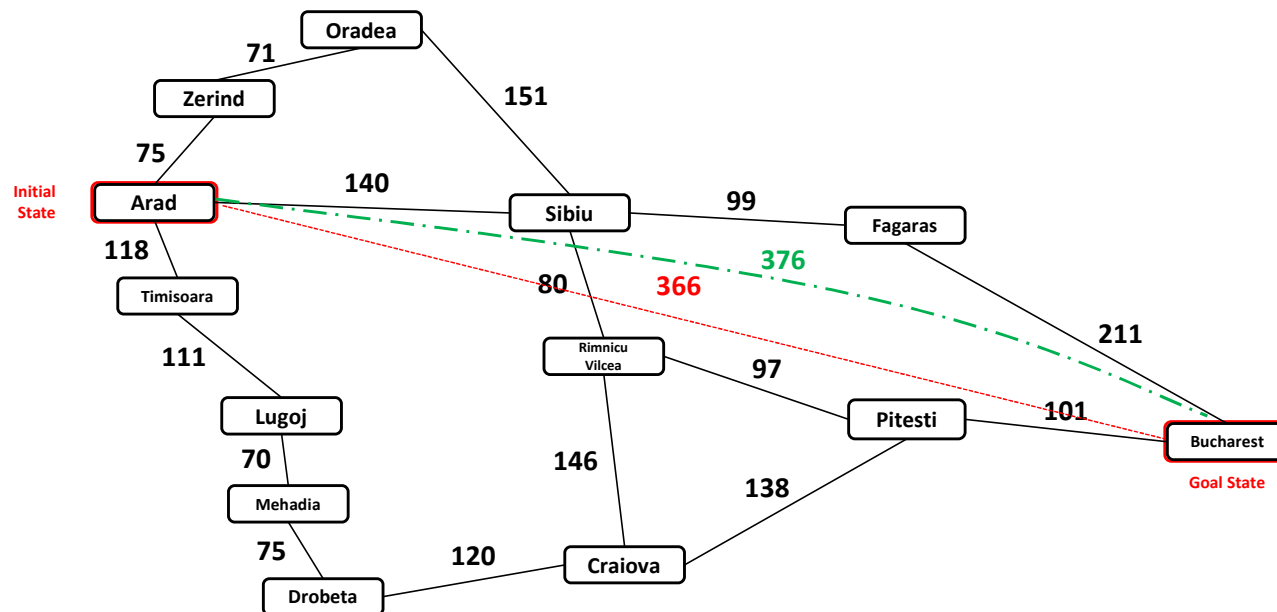
We can have more than one available heuristics.
For example $h_1(n)$ and $h_2(n)$.



Heuristics $h_2(n)$ estimate is closer to actual cost than $h_1(n)$. $h_2(n)$ dominates $h_1(n)$. Use $h_2(n)$.

Dominating Heuristics

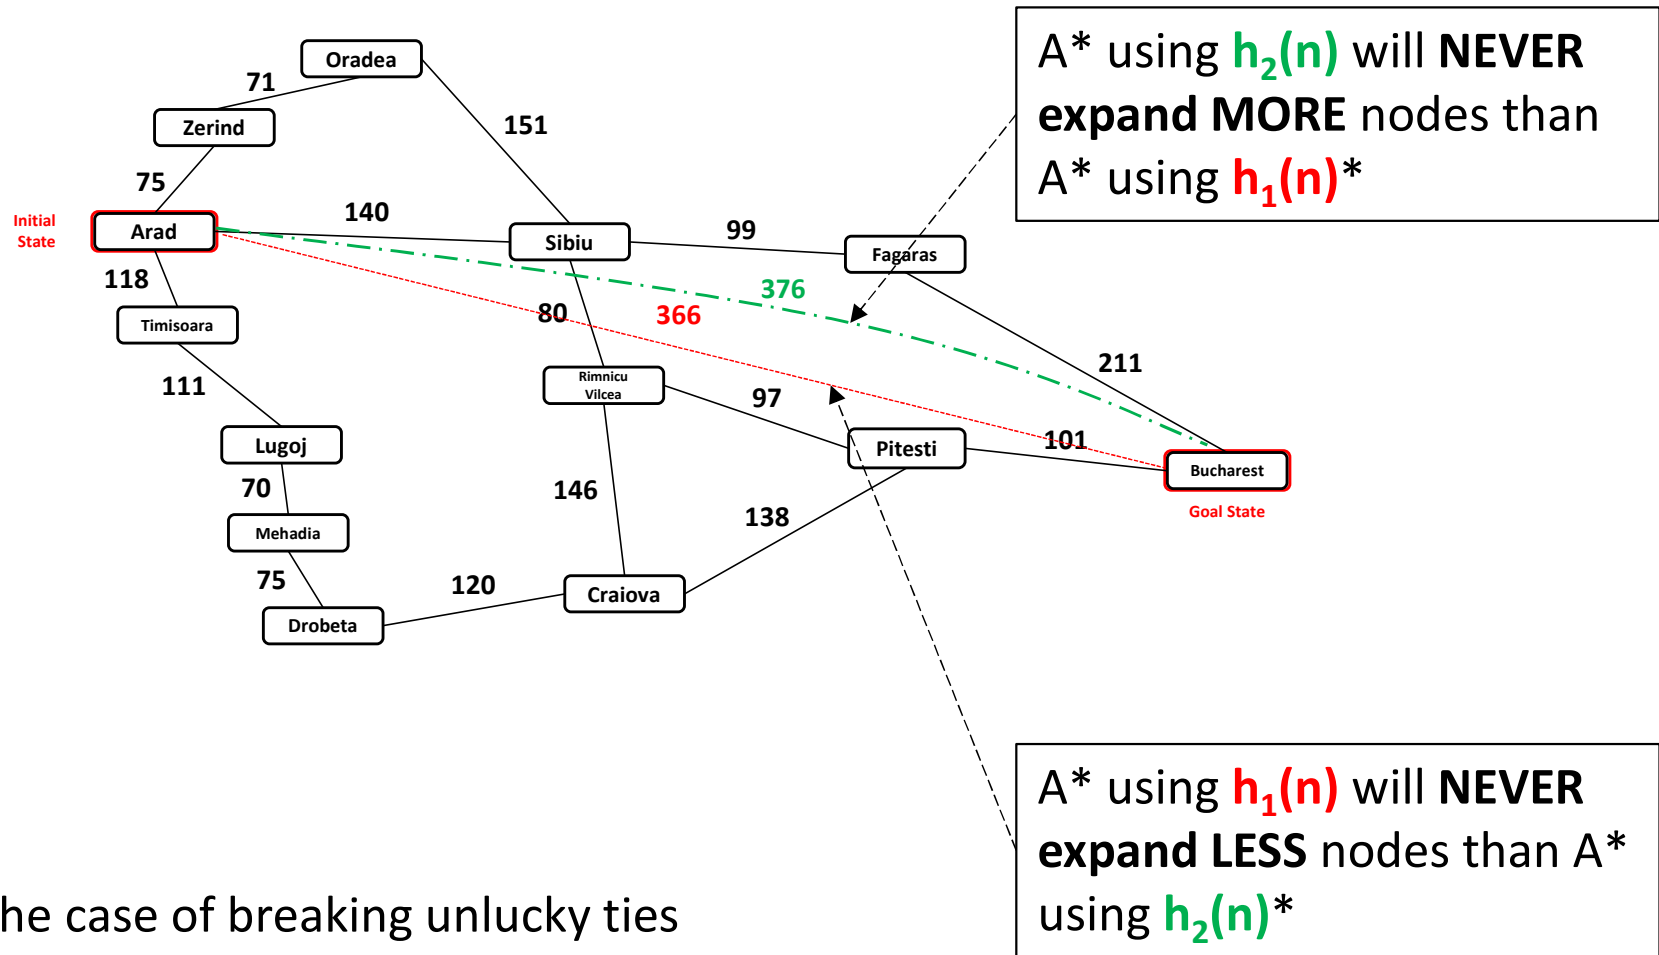
We can have more than one available heuristics.
For example $h_1(n)$ and $h_2(n)$.



If you have multiple admissible heuristics where none dominates the other: Let $h(n) = \max(h_1(n), h_2(n), \dots, h_m(n))$

Domination \rightarrow Efficiency

Heuristics $h_2(n)$ estimate is closer to actual cost than $h_1(n)$. $h_2(n)$ dominates $h_1(n)$. Use $h_2(n)$.



* Except for in the case of breaking unlucky ties

Domination → Efficiency: Why?

A* Evaluation Function:

$$f(n) = g(\text{State}_n) + h(\text{State}_n)$$

where:

- $g(n)$ - initial node to node n path cost
- $h(n)$ - **estimated cost** of the best path that continues from node n to a goal node

So:

$$f(n) < C^*$$

where:

- C^* - **optimal** path cost

Domination \rightarrow Efficiency: Why?

With

$$f(n) < C^*$$

and

$$f(n) = g(\text{State}_n) + h(\text{State}_n)$$

We get:

$$g(\text{State}_n) + h(\text{State}_n) < C^*$$

$$h(\text{State}_n) < C^* - g(\text{State}_n)$$

Domination \rightarrow Efficiency: Why?

If

$$h_1(\text{State}_n) < h_2(\text{State}_n)$$

then:

$$h_1(\text{State}_n) < h_2(\text{State}_n) < C^* - g(\text{State}_n)$$

Domination \rightarrow Efficiency: Why?

If

$$h_1(\text{State}_n) < h_2(\text{State}_n)$$

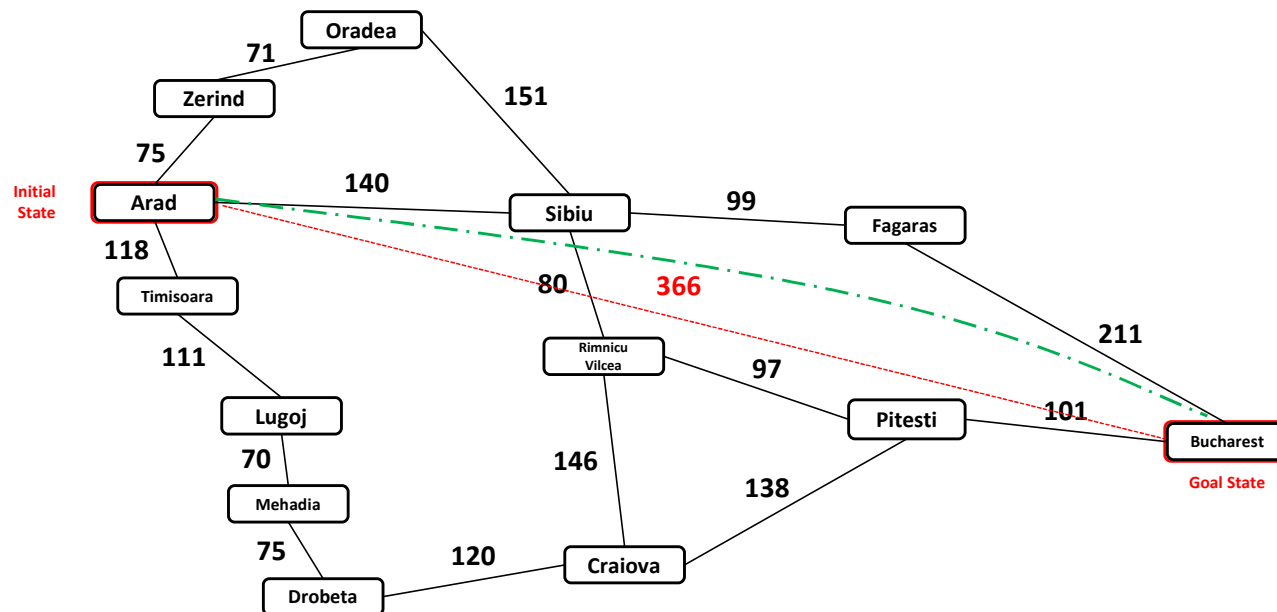
then:

$$h_1(\text{State}_n) < h_2(\text{State}_n) < C^* - g(\text{State}_n)$$

All A* frontier candidates selected using $h_1(n)$ will include **AT LEAST** all A* frontier candidates selected using $h_2(n)$. Possibly more.

Domination \rightarrow Efficiency: But?

$h_2(n)$ dominates $h_1(n)$. Always use $h_2(n)$?



Generally yes, but $h_2(n)$ vs. $h_1(n)$ heuristic computation time may be a deciding factor here.

Heuristic and Search Performance

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

Consider an 8-puzzle game and two admissible heuristics:

- $h_1()$ – number of misplaced tiles (not counting blank)
- $h_2()$ – Manhattan distance

Heuristic and Search Performance

7	2	4
5		6
8	3	1


Start State

	1	2
3	4	5
6	7	8


Goal State

- $h_1(\text{Start}) = 8$ All 8 blocks are out of place ○

Heuristic and Search Performance



Start State



Goal State

- $h_1(\text{Start}) = 8$
- $h_2(\text{Start}) = 3 + 1 + 2 + 2 + 2 + 3 + 3 + 2 = 18$

Heuristic and Search Performance

7	2	4
5		6
8	3	1

Start State

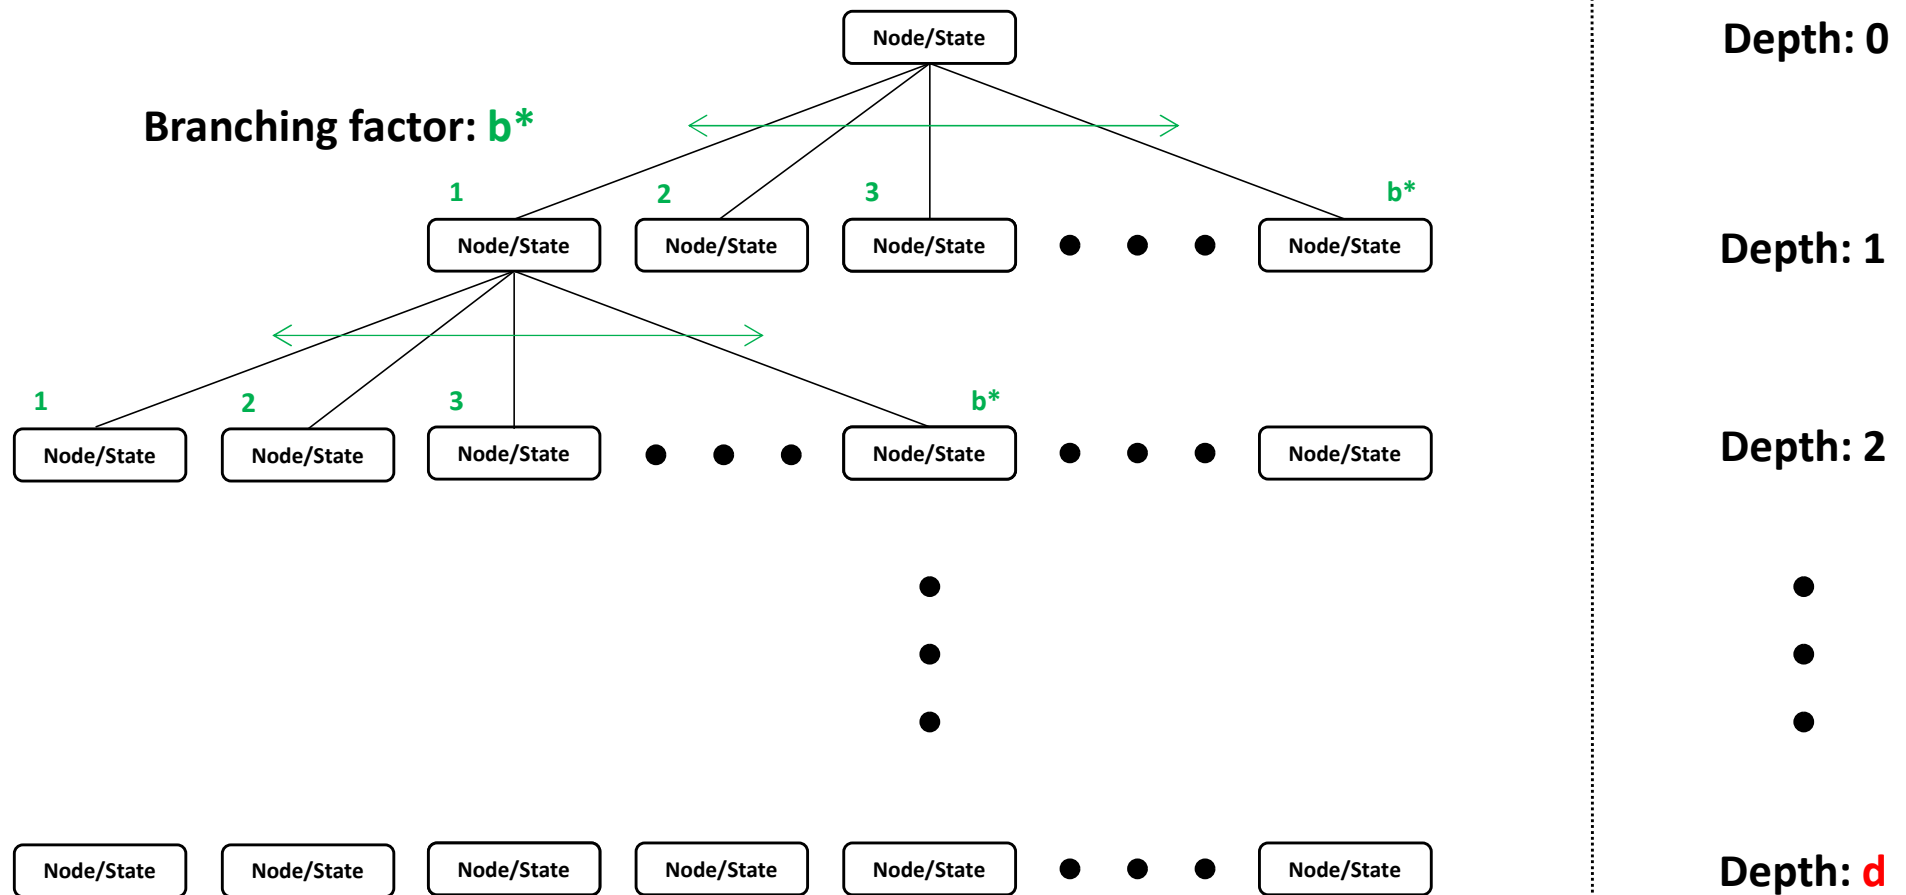
	1	2
3	4	5
6	7	8

Goal State

$$h_1(\text{Start}) = 8 < h_2(\text{Start}) = 18 < C^* = 26$$

Neither heuristic overestimates true solution cost

h() Quality: Effective Branching Factor



Total number of nodes / states: $1 + b^* + b^{*2} + b^{*3} + \dots + b^{*d} \rightarrow N + 1$

h() Quality: Effective Branching Factor

- If the total number of nodes generated is N and the solution depth is d , then
 - b^* is the branching factor that a **uniform tree of depth d** would need to have in order to contain $N+1$ nodes

$$N+1 = 1 + b^* + (b^*)^2 + \dots + (b^*)^d$$

- If A* finds a solution at depth 4 using 40 nodes, what is b^* ?

$$\approx 2.182$$

- A good heuristic function achieves $b^* \approx 1$

8-puzzle: Heuristics Comparison

d	Search Cost (nodes generated)			Effective Branching Factor		
	BFS	$A^*(h_1)$	$A^*(h_2)$	BFS	$A^*(h_1)$	$A^*(h_2)$
6	128	24	19	2.01	1.42	1.34
8	368	48	31	1.91	1.40	1.30
10	1033	116	48	1.85	1.43	1.27
12	2672	279	84	1.80	1.45	1.28
14	6783	678	174	1.77	1.47	1.31
16	17270	1683	364	1.74	1.48	1.32
18	41558	4102	751	1.72	1.49	1.34
20	91493	9905	1318	1.69	1.50	1.34
22	175921	22955	2548	1.66	1.50	1.34
24	290082	53039	5733	1.62	1.50	1.36
26	395355	110372	10080	1.58	1.50	1.35
28	463234	202565	22055	1.53	1.49	1.36

Where (data are averaged over 100 puzzles for each solution length $d = 6$ to 28):

- BFS – Breadth First Search
- $A^*(h_1)$ – A^* using $h_1()$ – number of misplaced tiles
- $A^*(h_2)$ – A^* using $h_2()$ – Manhattan distance

Can We Make A^* Even Faster?
(Sometimes at a cost!)

Weighted A* Evaluation Function

Weighted A* Evaluation Function:


$$f(n) = g(\text{State}_n) + W * h(\text{State}_n)$$

where:


- $g(n)$ - initial node to node n path cost
- $h(n)$ - **estimated cost** of the best path that continues from node n to a goal node
- $W > 1$

Here, weight W makes $h(\text{State}_n)$ (perhaps only “sometimes”) inadmissible. It **becomes potentially more accurate** = less expansions!

Weighted A*: “Good Enough” Search



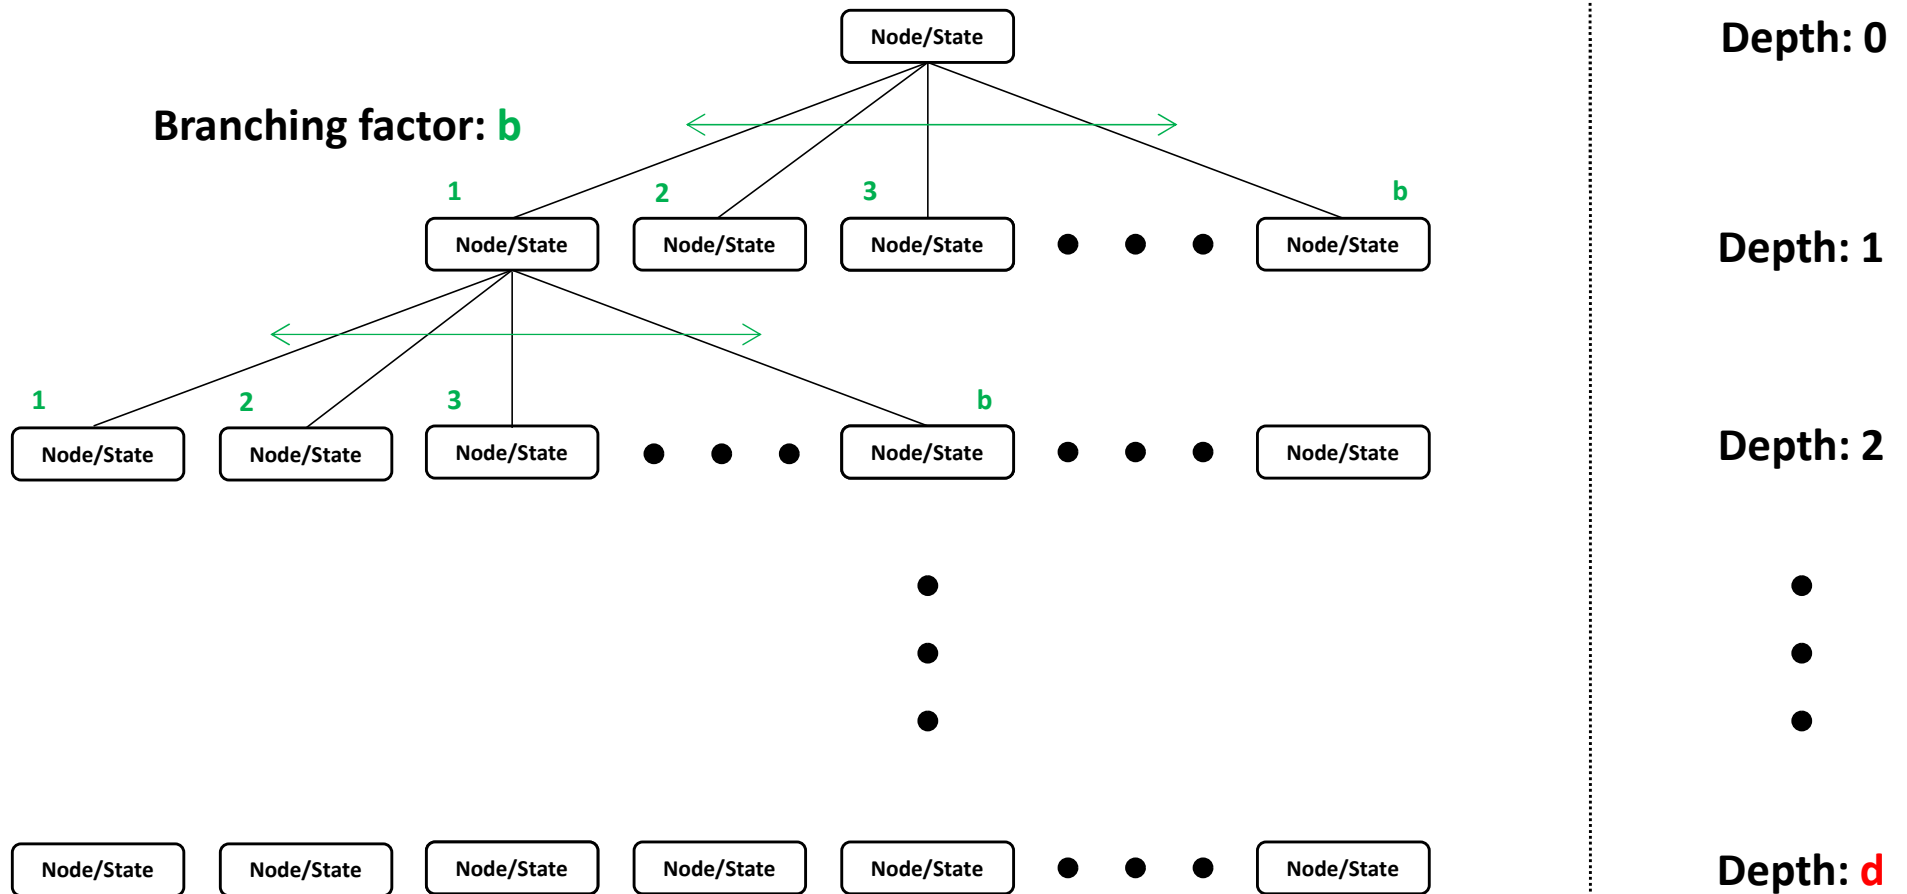
(a)



(b)


Figure 3.21 Two searches on the same grid: (a) an A* search and (b) a weighted A* search with weight $W = 2$. The gray bars are obstacles, the purple line is the path from the green start to red goal, and the small dots are states that were reached by each search. On this particular problem, weighted A* explores 7 times fewer states and finds a path that is 5% more costly.

Search Tree Challenges: Size



What are the possible **A* search problems** here? PLURAL!

A* Search Challenges: Complexity



Straight-line distance to Goal state					
State	I	A	B	C	G
h(State)	7	2	3	4	0

INITIAL STATE: I
GOAL STATE: G

Frontier	Parent	C	I	I				
	Node	B	B	A				
	f(Node)	6	7	7				
Reached	Parent	----	I	C	I			
	Key/State	I	A	B	C			
	Path cost	0	5	3	2			


State Space **Graph**

Frontier / Reached

Search **Tree**

We need to STORE multiple structures in memory:

- **Frontier** queue
- **Reached** table
- search **tree**




This may be too much to handle...



Frontier node

A* Search Challenges: Solutions



Straight-line distance to Goal state					
State	I	A	B	C	G
h(State)	7	2	3	4	0

INITIAL STATE: I
GOAL STATE: G

Frontier	Parent	C	I	I				
	Node	B	B	A				
	f(Node)	6	7	7				
Reached	Parent	----						
	Key/State	I	A	C				
	Path cost	0		3				


State Space **Graph**

Frontier / Reached

Search **Tree**


A) “Merge” **Frontier** and **Reached** data into one structure :

- less memory used, but
- more complex algorithm
- potentially slower algorithm



Frontier node

A* Search Challenges: Solutions



Straight-line distance to Goal state					
State	I	A	B	C	G
h(State)	7	2	3	4	0

INITIAL STATE: I
GOAL STATE: G

Frontier	Parent	C	I	I				
	Node	B	B	A				
	f(Node)	6	7	7				

Reached	Parent	----	I	C	I			
	Key/State	I	A	B	C			
	Path cost	0	5	3	2			


State Space **Graph**

Frontier / Reached

Search **Tree**

B) Remove nodes from **Reached once we can prove** that they no longer be needed




- ensure that such nodes will not be “touched” again



Frontier node


Can We Remove It From Reached?

State Space graph provides a list of state neighbors.



It can be used to decide if the state will be needed or not.

A* Search Challenges: Solutions



Straight-line distance to Goal state					
State	I	A	B	C	G
h(State)	7	2	3	4	0

INITIAL STATE: I
GOAL STATE: G

Frontier	Parent	C	I	I				
	Node	B	B	A				
	f(Node)	6	7	7				

Reached	Parent	----	I	C	I			
	Key/State	I	A	B	C			
	Path cost	0	5	3	2			


State Space **Graph**

Frontier / Reached

Search **Tree**

C) Keep **reference counts** of the number of times a state can be reached.




- once a state has been visited **reference counts** times, it can be removed from **Reached**



Frontier node


State Space: Reference Counts

State Space graph provides the number of neighbors for each state.



It can be used as a **reference count**.

A* Search Challenges: Solutions



Straight-line distance to Goal state					
State	I	A	B	C	G
$h(\text{State})$	7	2	3	4	0

INITIAL STATE: I
GOAL STATE: G

Frontier	Parent	C	I	I				
	Node	B	B	A				
	$f(\text{Node})$	6	7	7				
Reached	Parent	----	I	C	I			
	Key/State	I	A	B	C			
	Path cost	0	5	3	2			

State Space **Graph**


Frontier / Reached

Search **Tree**

D) **Beam Search** → restrict the size of the **Frontier**:


- keep k best $f()$ -score elements
- keep only elements with $f()$ -score such that:

$f()$ -score within δ of $f()_{\text{best}}$



Frontier node

A* Search Challenges: Solutions



Straight-line distance to Goal state					
State	I	A	B	C	G
h(State)	7	2	3	4	0


INITIAL STATE: I
GOAL STATE: G

Frontier	Parent	C	I	I				
	Node	B	B	A				
	f(Node)	6	7	7				
Reached	Parent	----	I	C	I			
	Key/State	I	A	B	C			
	Path cost	0	5	3	2			

State Space **Graph**

Frontier / Reached

Search **Tree**



Reduce size



Frontier node

E) Use **Iterative-Deepening A*** search (IDA*)

Iterative Deepening **DFS**: Illustration


limit: 0



limit: 1




limit: 2




Iterative Deepening **DFS**: Illustration

limit: 3



A* Search Challenges: Solutions



Straight-line distance to Goal state					
State	I	A	B	C	G
h(State)	7	2	3	4	0


INITIAL STATE: I
GOAL STATE: G

Frontier	Parent	C	I	I				
	Node	B	B	A				
	f(Node)	6	7	7				
Reached	Parent	----	I	C	I			
	Key/State	I	A	B	C			
	Path cost	0	5	3	2			

State Space **Graph**

Frontier / Reached

Search **Tree**



Reduce size



Frontier node

F) Use **Recursive Best First Search**


- maintain $f()$ limit (best alternative ancestor path $f()$)
- if current node's $f()$ exceeds $f()$ limit, recursion unwinds

Recursive Best First Search

function RECURSIVE-BEST-FIRST-SEARCH(*problem*) **returns** a solution or *failure*
 solution, *fvalue* \leftarrow RBFS(*problem*, NODE(*problem*.INITIAL), ∞)
return *solution*

function RBFS(*problem*, *node*, *f_limit*) **returns** a solution or *failure*, and a new *f*-cost limit
 if *problem*.IS-GOAL(*node*.STATE) **then return** *node*
 successors \leftarrow LIST(EXPAND(*node*))
 if *successors* is empty **then return** *failure*, ∞
 for each *s* **in** *successors* **do** // update *f* with value from previous search
 s.f \leftarrow max(*s*.PATH-COST + *h*(*s*), *node.f*)
 while true do
 best \leftarrow the node in *successors* with lowest *f*-value
 if *best.f* > *f_limit* **then return** *failure*, *best.f*
 alternative \leftarrow the second-lowest *f*-value among *successors*
 result, *best.f* \leftarrow RBFS(*problem*, *best*, min(*f_limit*, *alternative*))
 if *result* \neq *failure* **then return** *result*, *best.f*

A* Search Challenges: Solutions



Straight-line distance to Goal state					
State	I	A	B	C	G
h(State)	7	2	3	4	0

INITIAL STATE: I
GOAL STATE: G

Frontier	Parent	C	I	I				
	Node	B	B	A				
	f(Node)	6	7	7				


Reached	Parent	----	I	C	I			
	Key/State	I	A	B	C			
	Path cost	0	5	3	2			

State Space **Graph**

Frontier / Reached

Search **Tree**

G) Use **Bidirectional A* Search**



Frontier node

Bidirectional A*

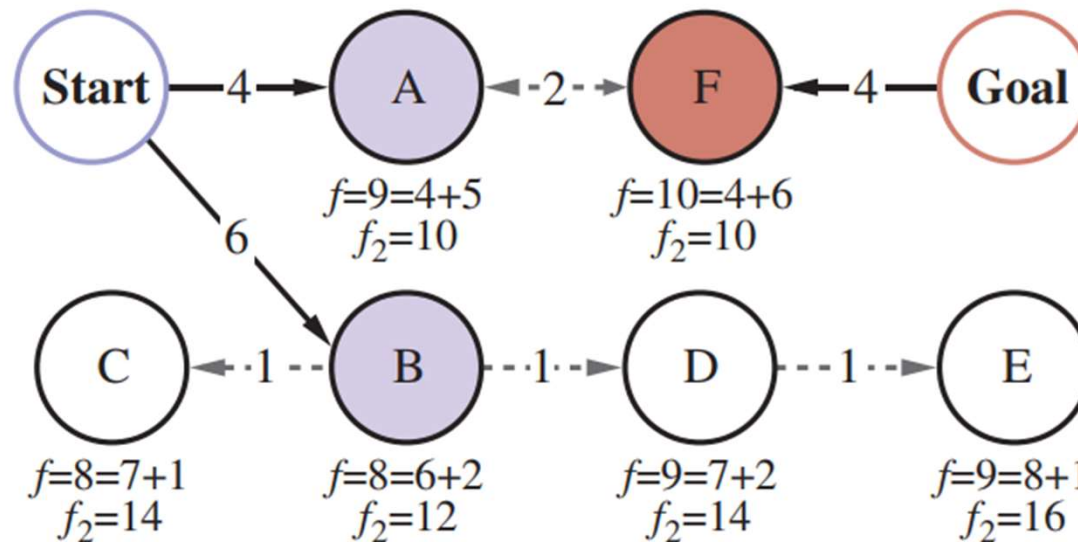


Figure 3.24 Bidirectional search maintains two frontiers: on the left, nodes A and B are successors of Start; on the right, node F is an inverse successor of Goal. Each node is labeled with $f = g + h$ values and the $f_2 = \max(2g, g + h)$ value. (The g values are the sum of the action costs as shown on each arrow; the h values are arbitrary and cannot be derived from anything in the figure.) The optimal solution, Start-A-F-Goal, has cost $C^* = 4 + 2 + 4 = 10$, so that means that a meet-in-the-middle bidirectional algorithm should not expand any node with $g > \frac{C^*}{2} = 5$; and indeed the next node to be expanded would be A or F (each with $g = 4$), leading us to an optimal solution. If we expanded the node with lowest f cost first, then B and C would come next, and D and E would be tied with A, but they all have $g > \frac{C^*}{2}$ and thus are never expanded when f_2 is the evaluation function.

Summary

Environment Assumptions

“Simple Environment”:

- Fully observable
- Single agent
- Deterministic
- Static
- Episodic or sequential
- Discrete
- Known to the agent

Defining Search Problem

Before agent searching, formulate a well-formed problem:

- Define a set of possible states: **State Space**
- Specify **Initial State**
- Specify **Goal State(s)** (there can be multiple)
- Define a FINITE set of possible **Actions** for EACH state in the State Space
- Come up with a **Transition Model** which describes what each action does
- Specify the **Action Cost Function**: a function that gives the cost of applying action a in state s

State Space: A Graph

$ACTIONS(A) = \{toB, toC\}$

$ACTIONS(B) = \{toF\}$

$ACTIONS(F) = \{toE\}$

Action: toB

$RESULT(A, toB) = B$

$C(A, toB, B) = 1$

INITIAL
STATE
A

STATE
B

STATE
F

Action: toE

$C(F, toE, E) = 1$

Action: toF

$C(B, toF, F) = 1$

Action: toE

$C(D, toE, E) = 1$

Action: toB

$C(C, toB, B) = 1$

$C(A, toC, C) = 1$

Action: toC

STATE
C

STATE
D

GOAL
STATE
E

$RESULT(A, toC) = C$

Action: toD

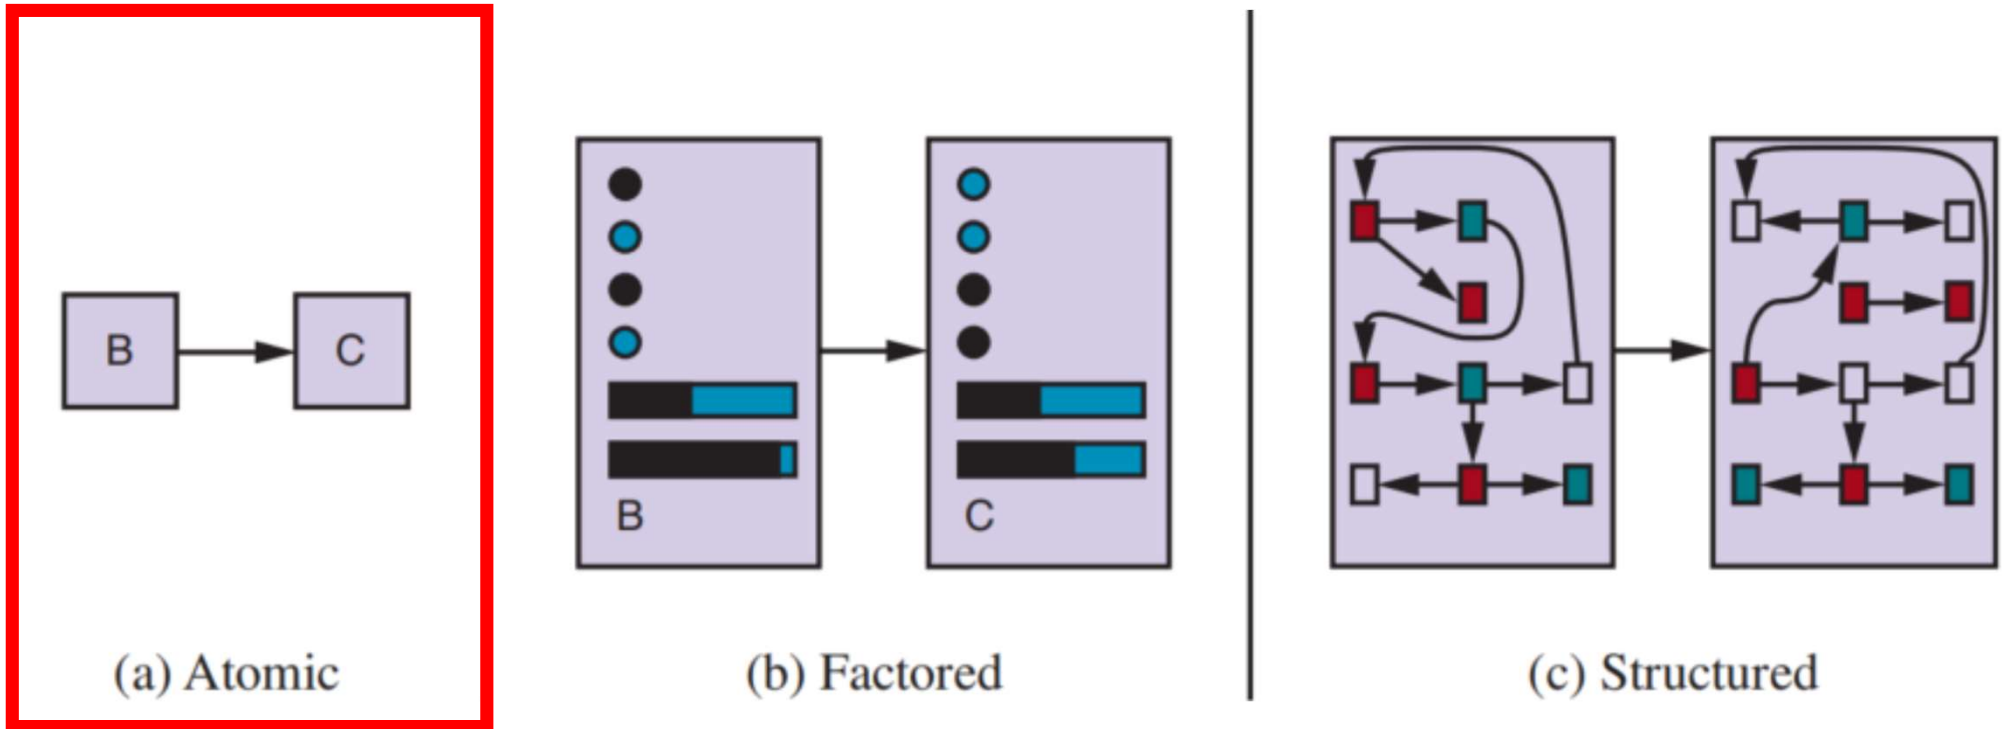
$C(C, toD, D) = 1$

$ACTIONS(C) = \{toB, toD\}$

$ACTIONS(D) = \{toE\}$

$ACTIONS(E) = \emptyset$

States: Usually Atomic Representation



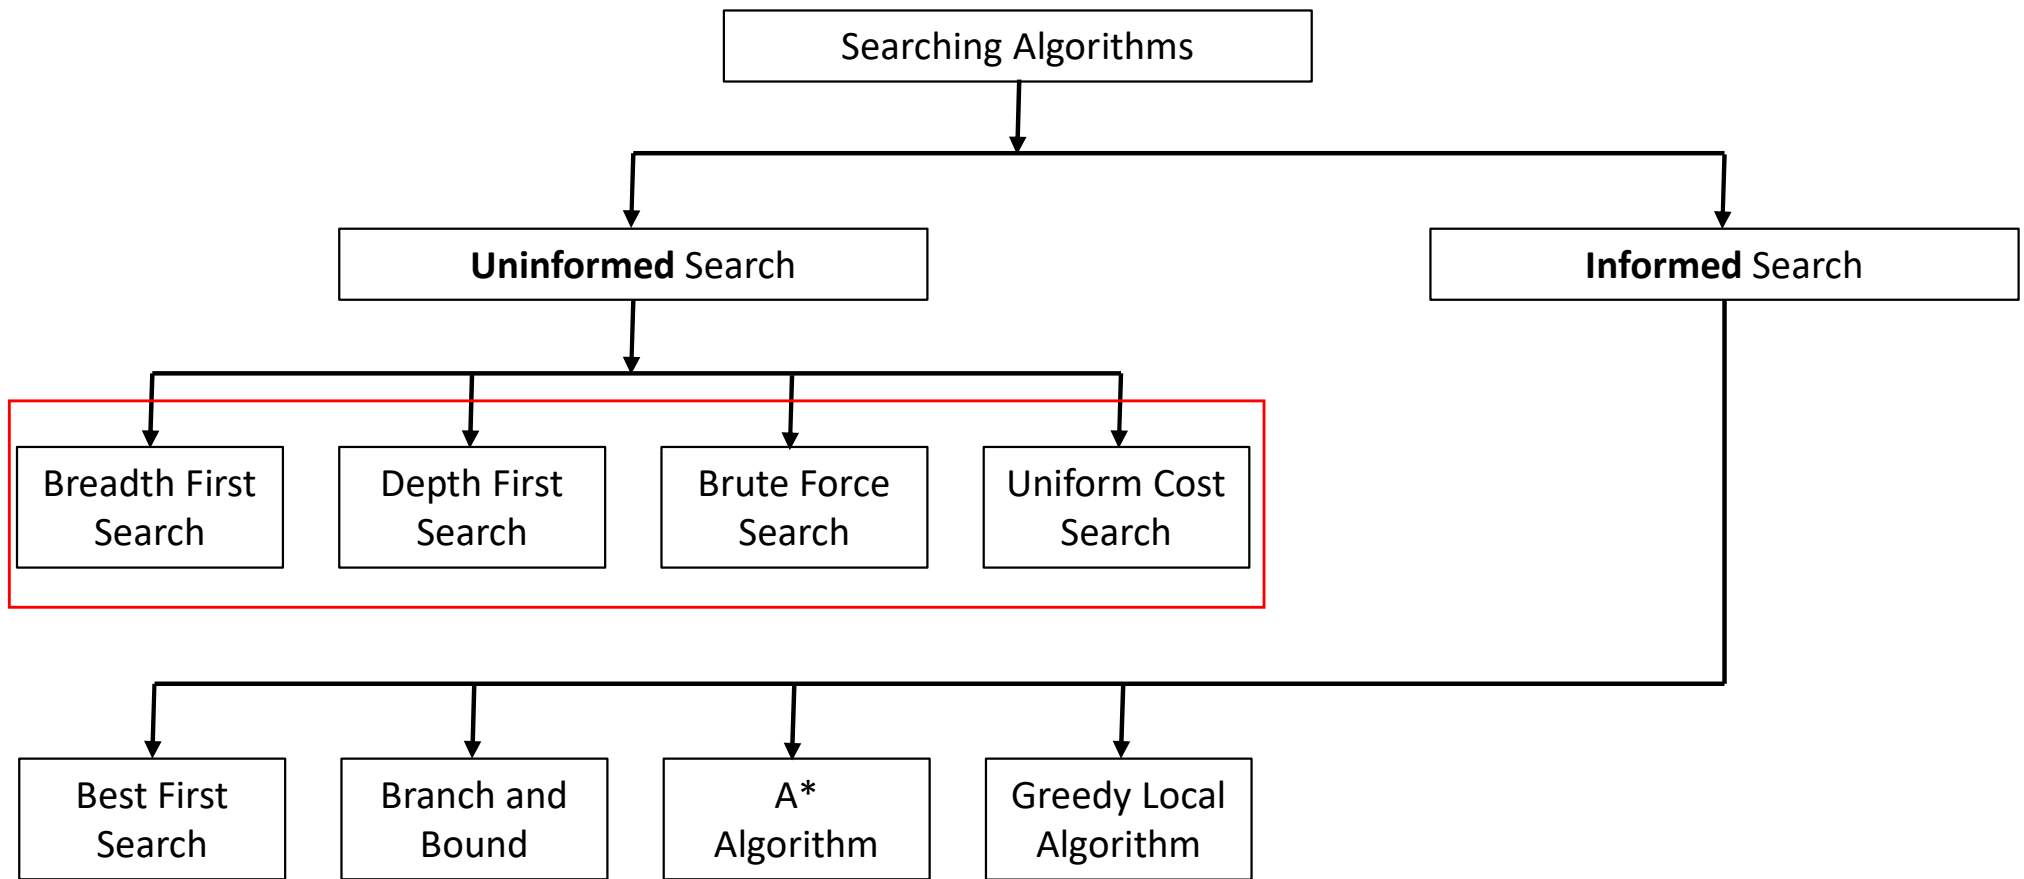
- **Atomic:** state representation has **NO** internal structure
- **Factored:** state representation includes fixed attributes (which can have values)
- **Structured:** state representation includes objects and their relationships

Judging Searching Performance

Search algorithms can be judged by:

- **Completeness**: Is the algorithm guaranteed to find a solution when there is one, and to correctly report failure when there is not?
- **Cost optimality**: Does it find a solution with the lowest path cost of all solutions?
- **Time complexity**: How long does it take to find a solution? (in seconds, actions, states, etc.)
- **Space complexity**: How much memory is needed to perform the search?

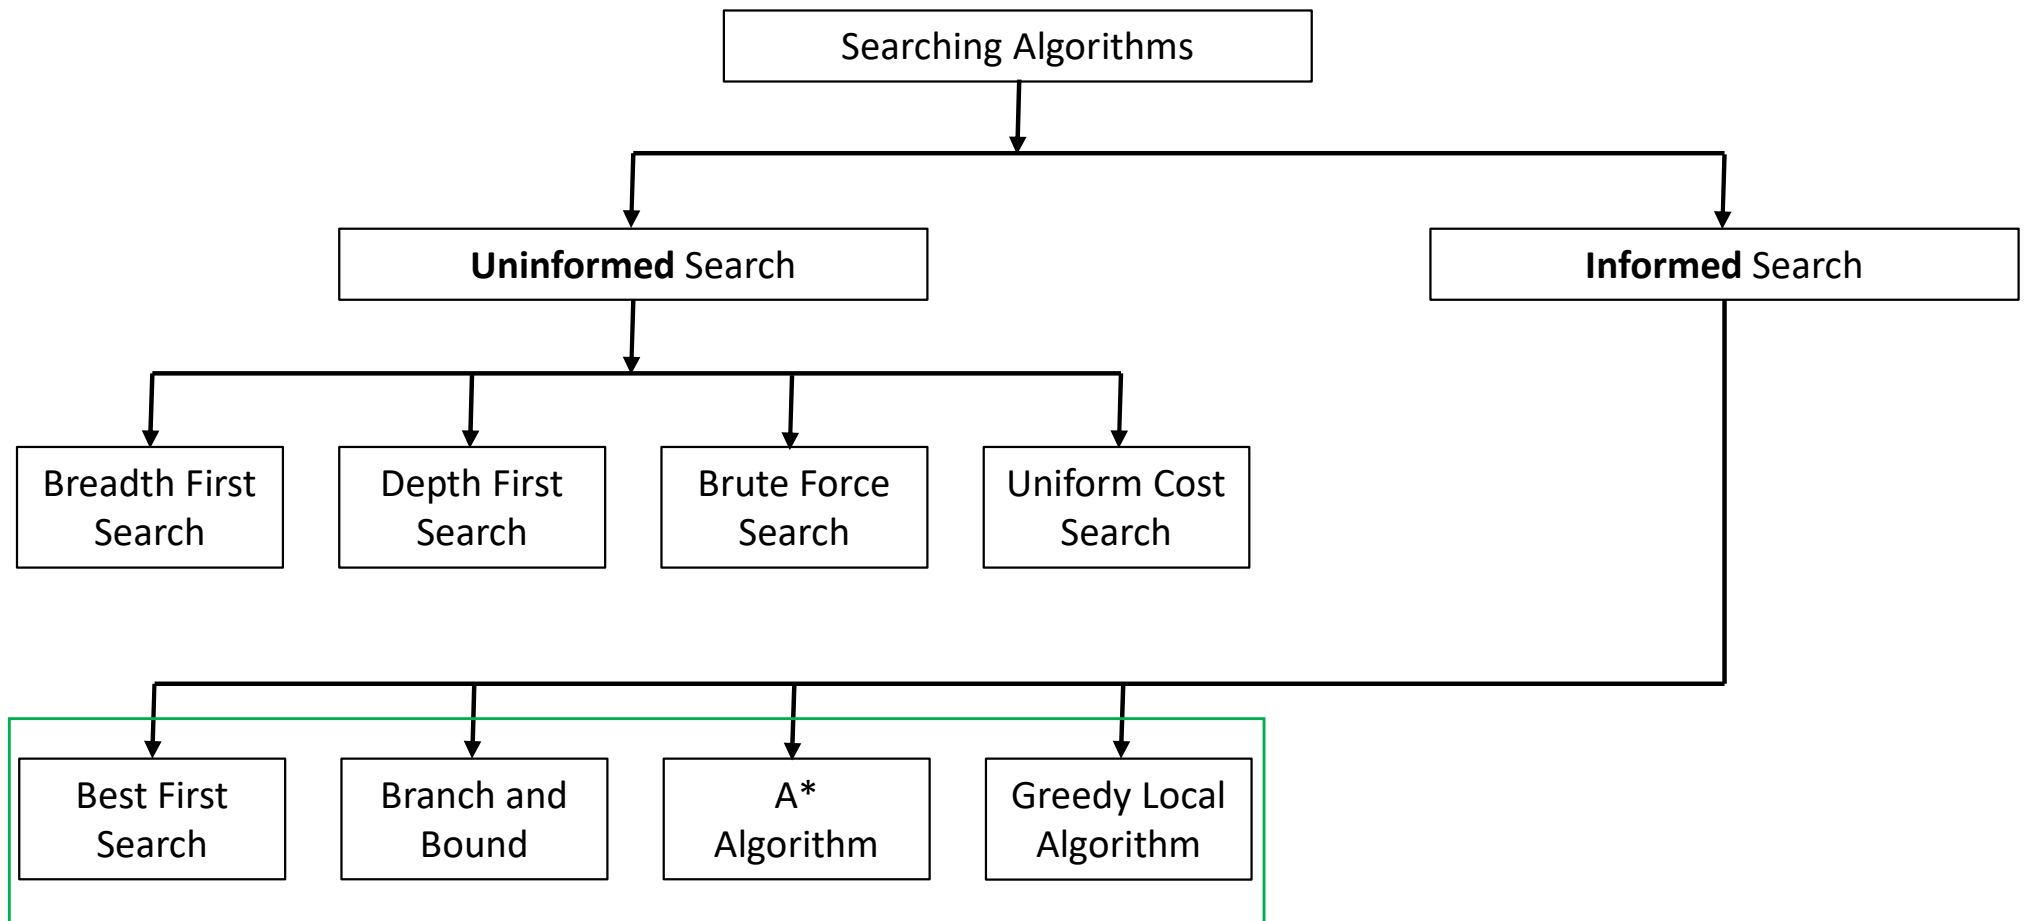
UNINFORMED Searching Algorithms



Uninformed search: an agent has no estimate how far it is from the goal

Informed search: an agent can estimate how far it is from the goal

INFORMED Searching Algorithms



Uninformed search: an agent has no estimate how far it is from the goal

Informed search: an agent can estimate how far it is from the goal

INFORMED Searching Algorithms

- **Greedy Best First Search** expands nodes with minimal $f(n) = h(n)$
 - complete, **not optimal**, but often efficient
- **A* Search** expands nodes with minimal $f(n) = g(n) + h(n)$
 - **complete and optimal IF $h(n)$ is admissible**
 - space complexity an issue
- **Bidirectional A* Search** is sometimes more efficient than pure A* Search
- **Iterative Deepening A* (IDA*)** is an iterative deepening version of A*, and thus addresses the space complexity issue

INFORMED Searching Algorithms

- RBFS (Recursive Best First Search) and SMA* (Simplified Memory-Bounded A*) are robust, optimal search algorithms that **use limited amounts of memory**. Given enough time, **they can solve problems for which A* runs out of memory**
- **Weighted A* Search** **focuses the search** towards the goal and **expands fewer nodes at the expense of search optimality**
- **Beam Search** puts a **limit on the size of Frontier**. That makes it **incomplete and suboptimal**. Often finds “good enough” solution and is faster than complete search

Finding / Selecting Heuristics

The performance of heuristic-based search depends on the quality of the heuristic function $h()$.

- Dominating heuristic is better but can be costlier to compute
- Ideas for constructing good heuristic:
 - relax problem definition
 - store precomputed solution costs for subproblems in a pattern database
 - define landmarks
 - learn from experience with the problem class

“Hill Climbing” (Greedy Local) Search and Romanian Roadtrip Example

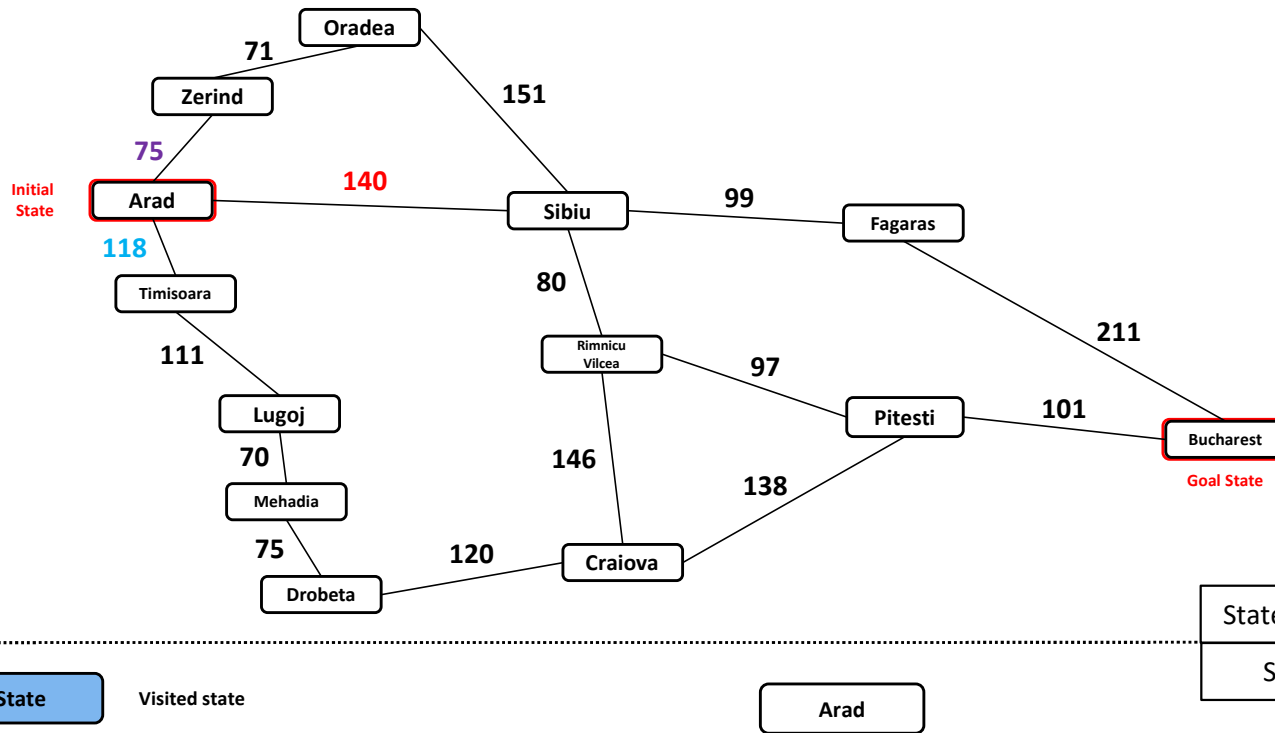
Greedy Local: Evaluation function

Calculate / obtain:

$$f(n) = \text{ACTION-COST}(\text{State}_a, \text{toState}_n, \text{State}_n)$$

A state n with minimum (or maximum) $f(n)$
should be chosen for expansion

Romanian Roadtrip: Greedy Local




Assumption:

We don't "go" to a repeated state

Alternatively:

- We could go to a repeated state and
 - get stuck there
 - or
 - Infinite loop

Romanian Roadtrip: Greedy Local




Assumption:
We don't "go" to a repeated state

State Space Graph

Search Tree

State

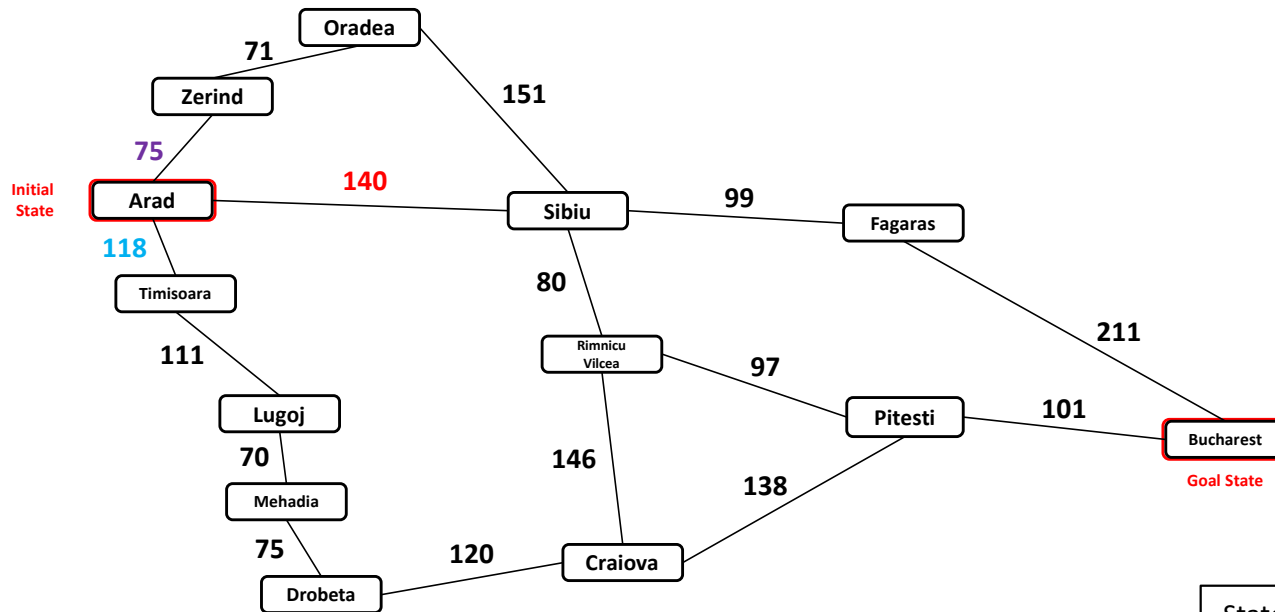
Visited state



Alternatively:

- We could go to a repeated state and
 - get stuck there
- or
- Infinite loop

Romanian Roadtrip: Greedy Local




Assumption:
We don't "go" to a repeated state

State

Visited state

State Space Graph

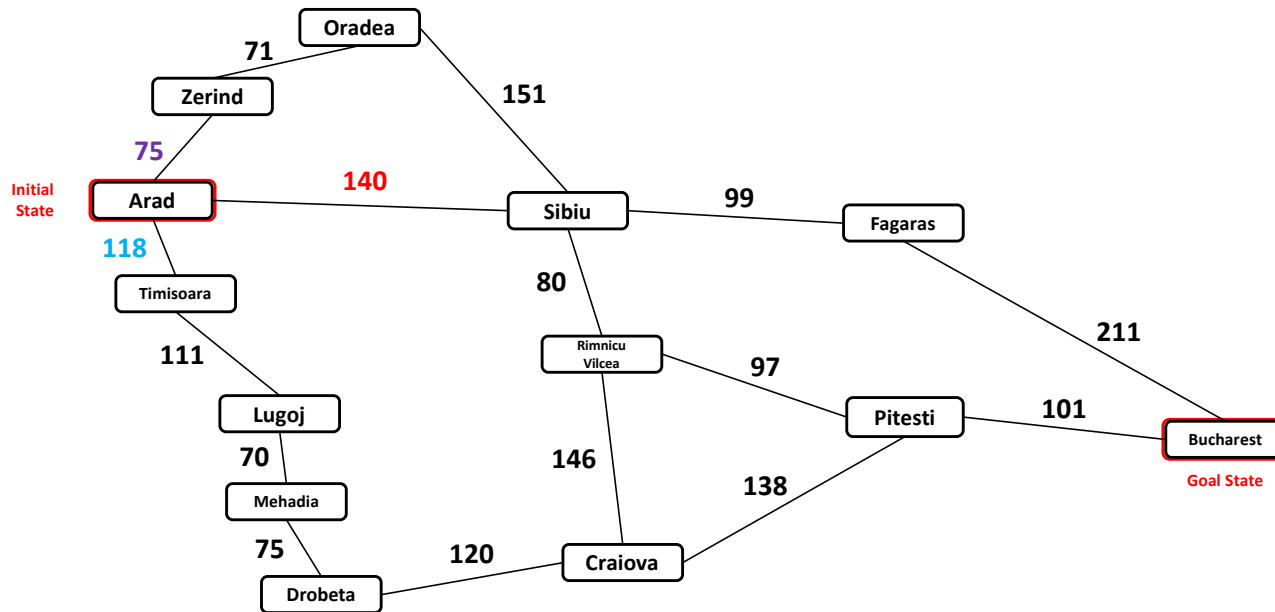
Search Tree



Alternatively:


- We could go to a repeated state and
 - get stuck there
- or
- Infinite loop

Romanian Roadtrip: Greedy Local



Assumption:


We don't "go" to a repeated state




Alternatively:

- We could go to a repeated state and
 - get stuck there
- or
- Infinite loop

Romanian Roadtrip: Greedy Local



Assumption:
We don't "go" to a repeated state



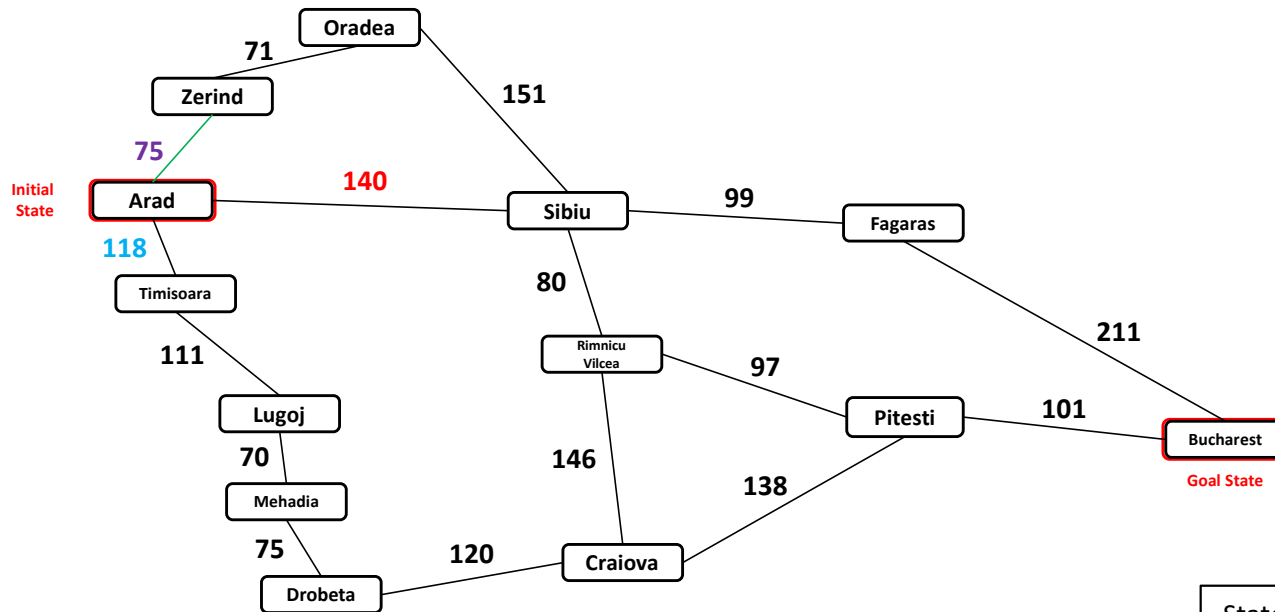
State Space Graph

Search Tree

Alternatively:

- We could go to a repeated state end
 - get stuck there
- or
- Infinite loop

Romanian Roadtrip: Greedy Local




Assumption:
We don't "go" to a repeated state

State Space Graph

Search Tree

State

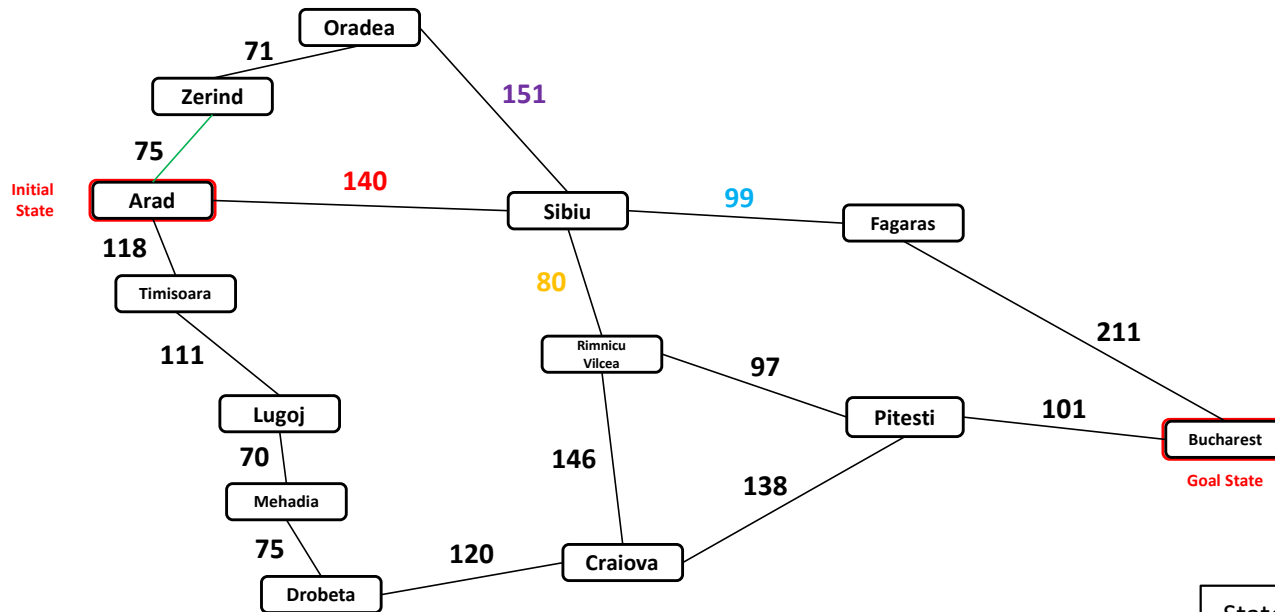
Visited state



Alternatively:

- We could go to a repeated state and
 - get stuck there
- or
- Infinite loop

Romanian Roadtrip: Greedy Local




Assumption:

We don't "go" to a repeated state

State Space Graph

Search Tree

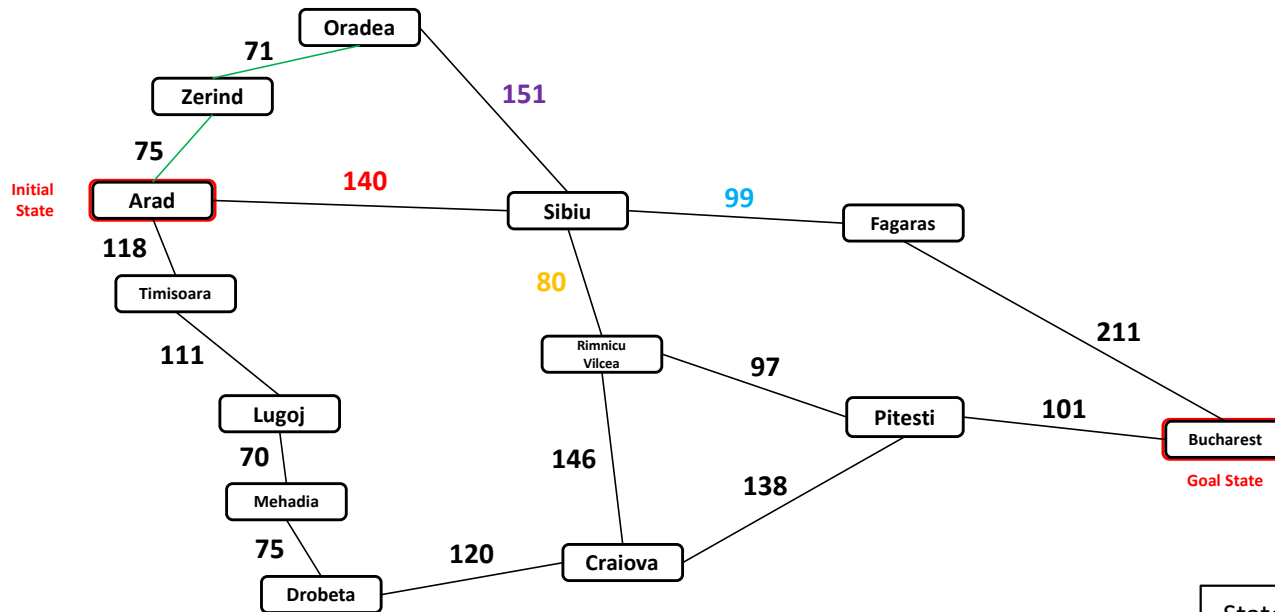
State Visited state



Alternatively:

- We could go to a repeated state end
 - get stuck there
- or
- Infinite loop

Romanian Roadtrip: Greedy Local




Assumption:

We don't "go" to a repeated state

State Space **Graph**

Search **Tree**


State Visited state



Alternatively:

- We could go to a repeated state end
 - get stuck there
 - or
 - Infinite loop

Romanian Roadtrip: Greedy Local

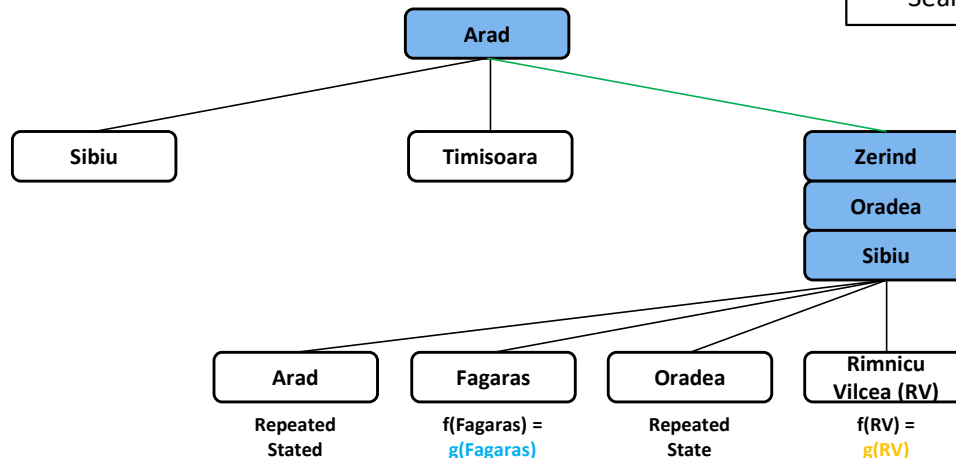


Assumption:
We don't "go" to a repeated state

State Space Graph

Search Tree

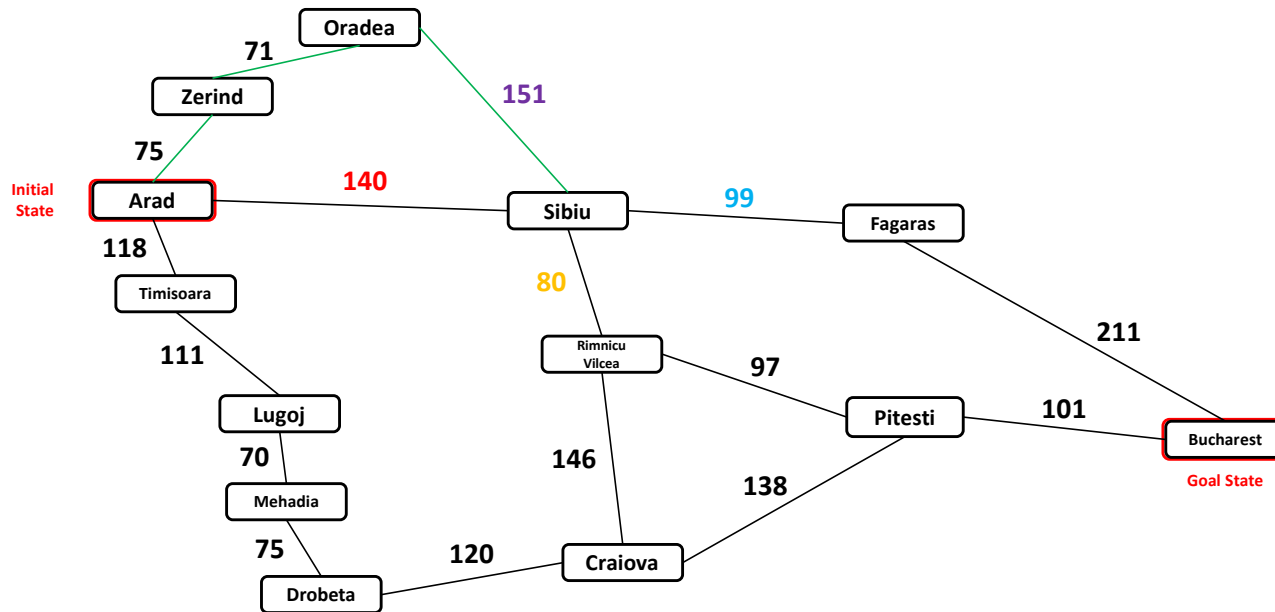
State Visited state



Alternatively:

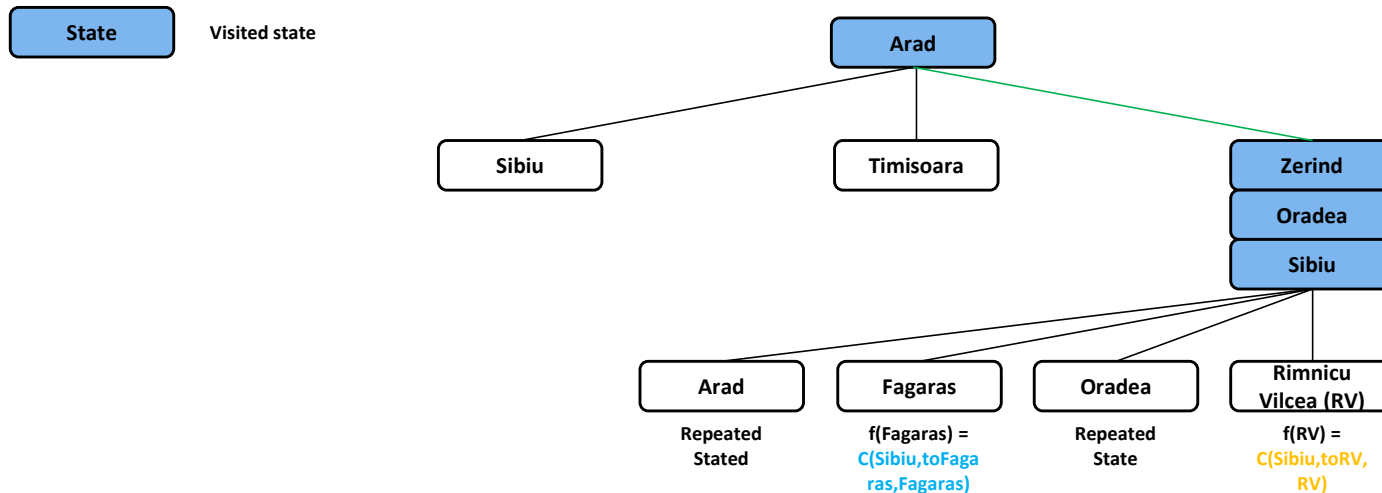
- We could go to a repeated state end
 - get stuck there
- or
- Infinite loop

Romanian Roadtrip: Greedy Local



Assumption:


We don't "go" to a repeated state



Alternatively:

- We could go to a repeated state end
 - get stuck there
- or
- Infinite loop

Romanian Roadtrip: Greedy Local



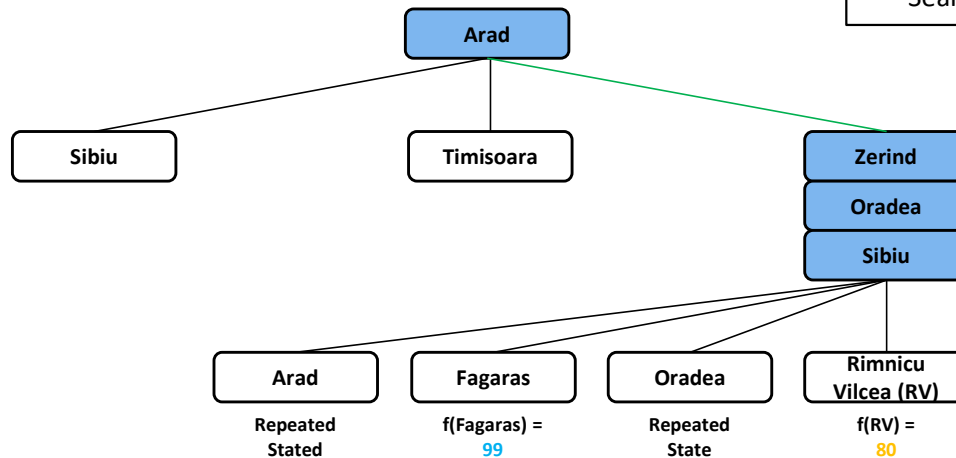
Assumption:
We don't "go" to a repeated state

State Space Graph

Search Tree

State

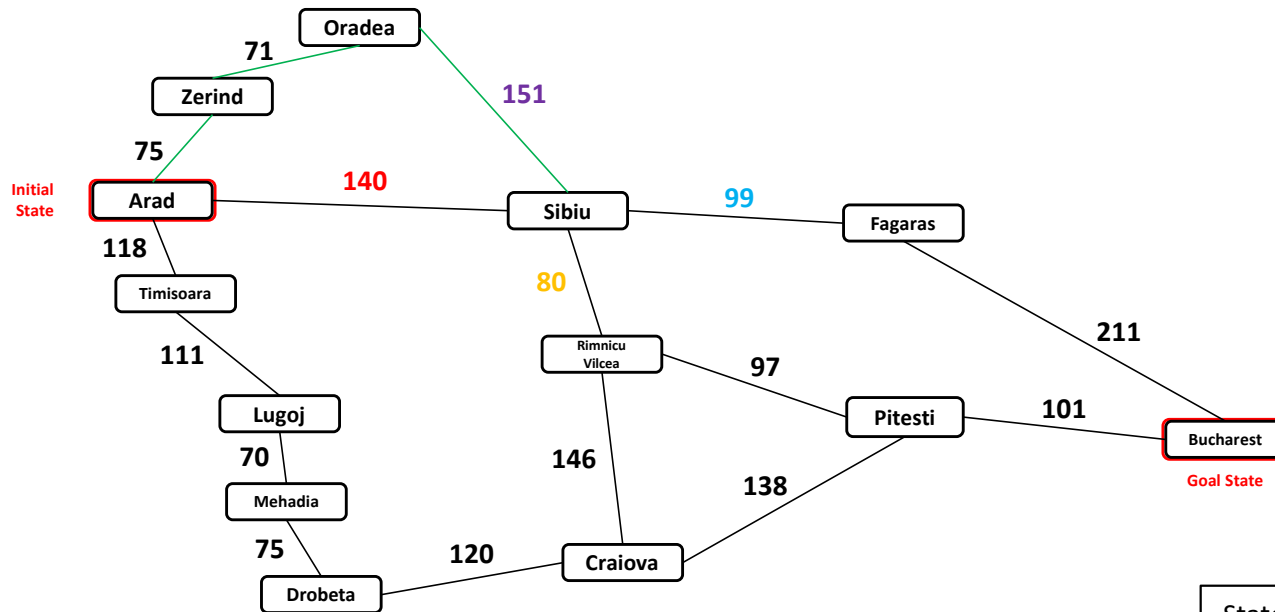
Visited state



Alternatively:

- We could go to a repeated state end
 - get stuck there
- or
- Infinite loop

Romanian Roadtrip: Greedy Local



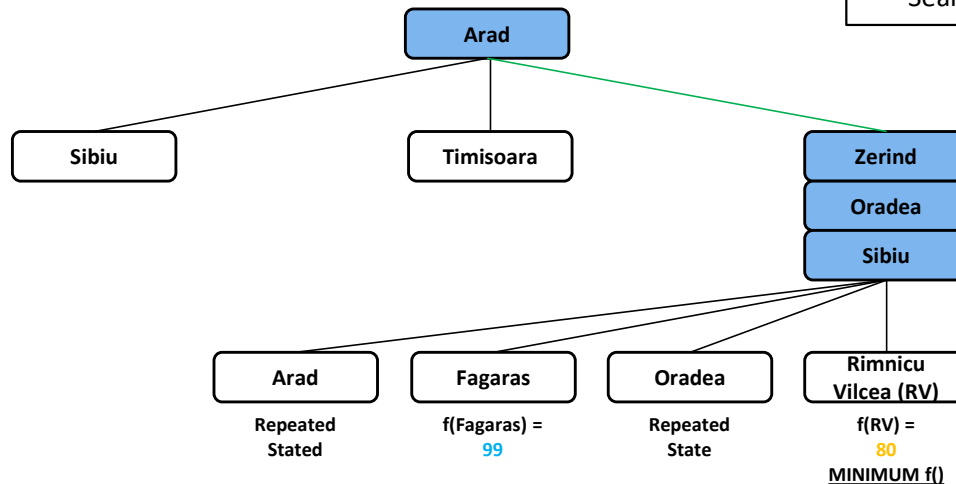
Assumption:
We don't "go" to a repeated state

State Space Graph

Search Tree

State


Visited state



Alternatively:

- We could go to a repeated state end
 - get stuck there
- or
- Infinite loop

Romanian Roadtrip: Greedy Local



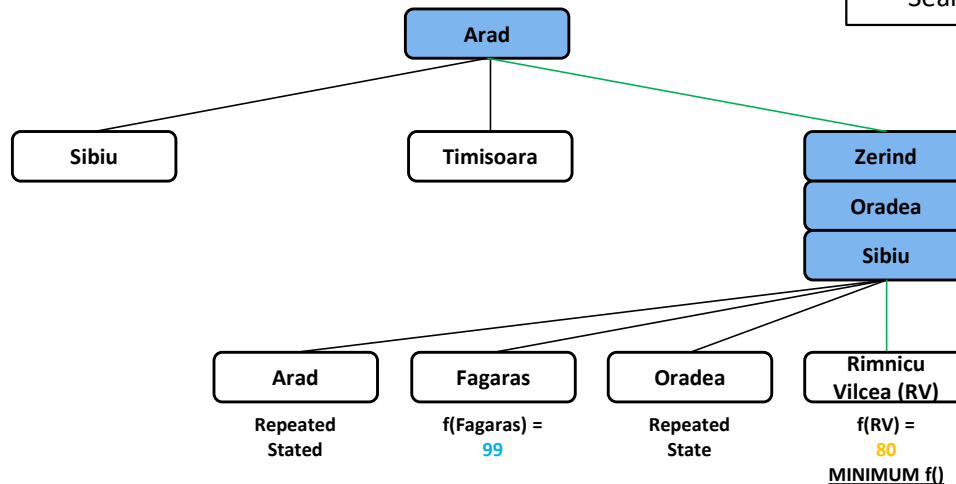
Assumption:
We don't "go" to a repeated state

State Space Graph

Search Tree

State


Visited state



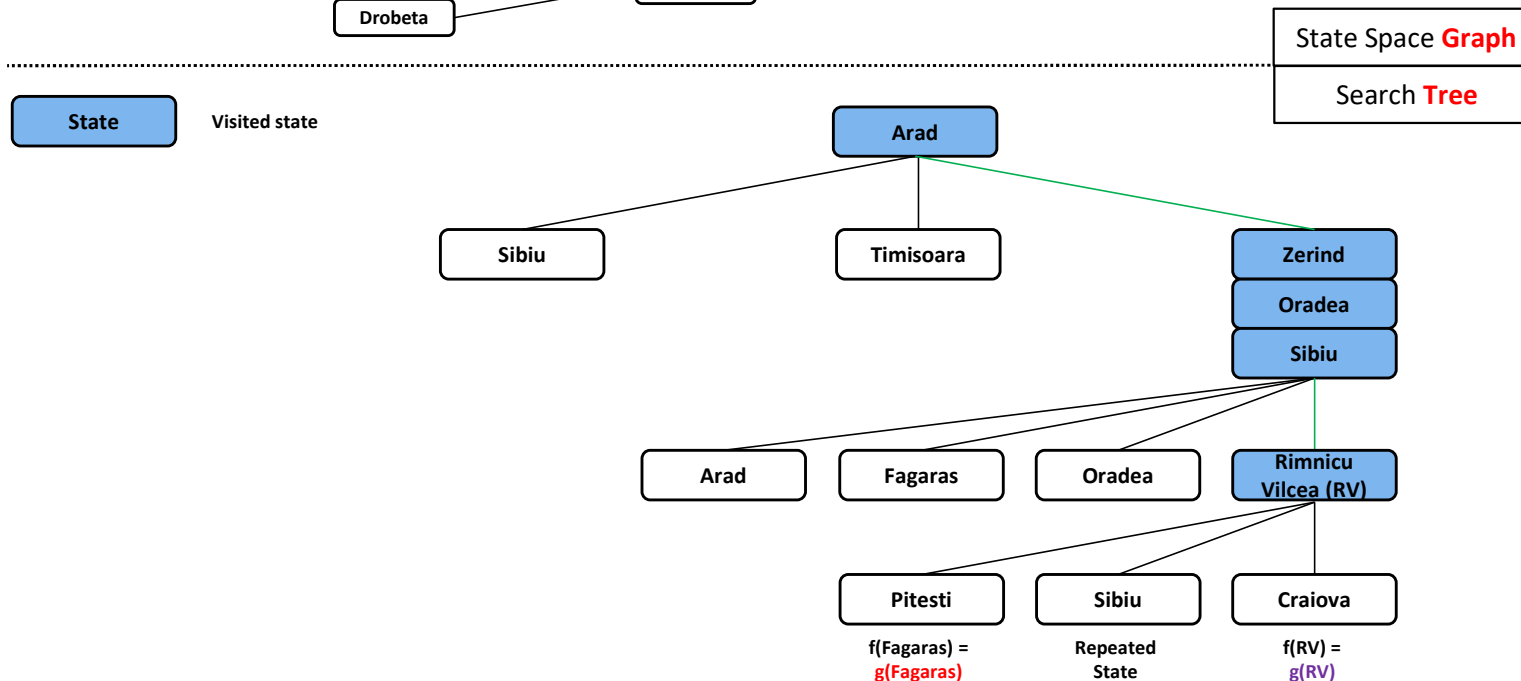
Alternatively:

- We could go to a repeated state end
 - get stuck there
- or
- Infinite loop

Romanian Roadtrip: Greedy Local




Assumption:
We don't "go" to a repeated state



Alternatively:

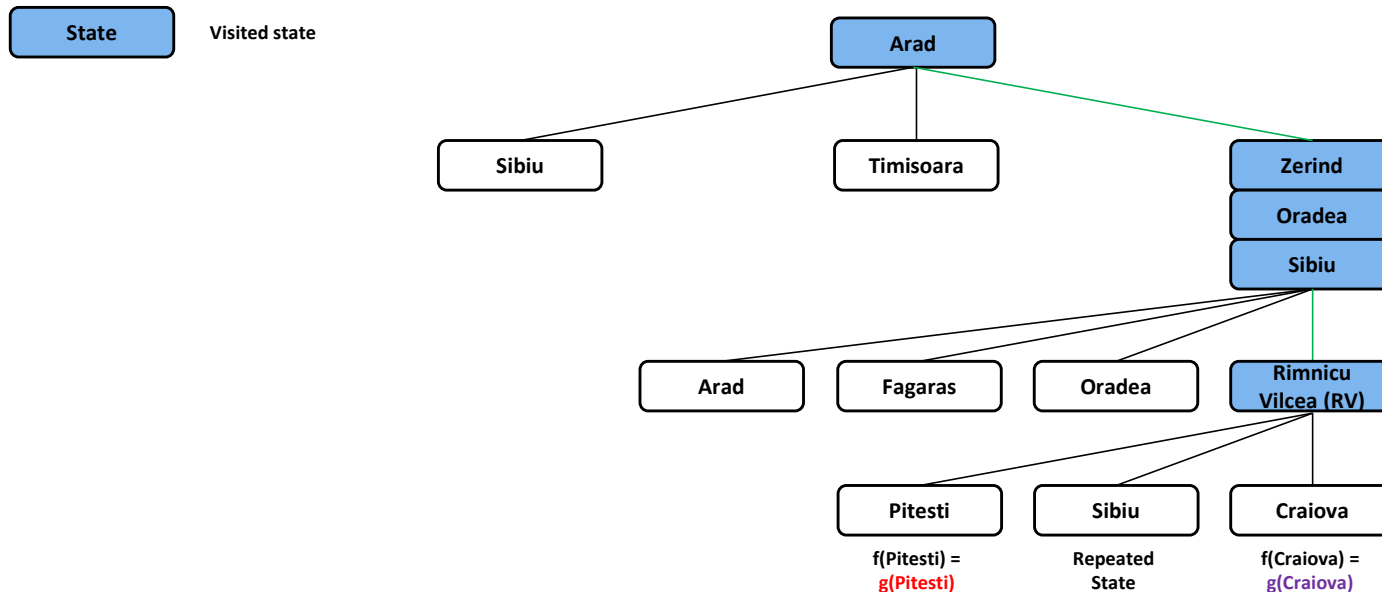
- We could go to a repeated state end
 - get stuck there
- or
- Infinite loop

Romanian Roadtrip: Greedy Local



Assumption:


We don't "go" to a repeated state



Alternatively:

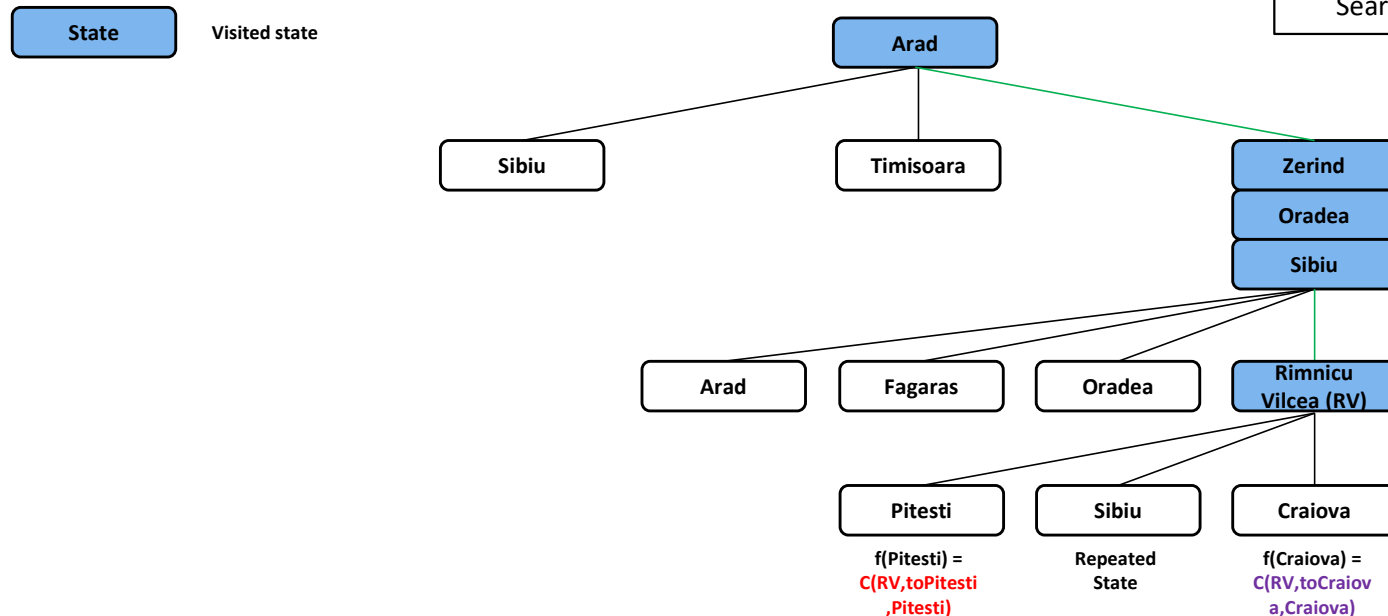
- We could go to a repeated state end
 - get stuck there
- or
- Infinite loop

Romanian Roadtrip: Greedy Local



Assumption:
We don't "go" to a repeated state


State Space Graph
Search Tree



Alternatively:

- We could go to a repeated state end
 - get stuck there
- or
- Infinite loop

Romanian Roadtrip: Greedy Local

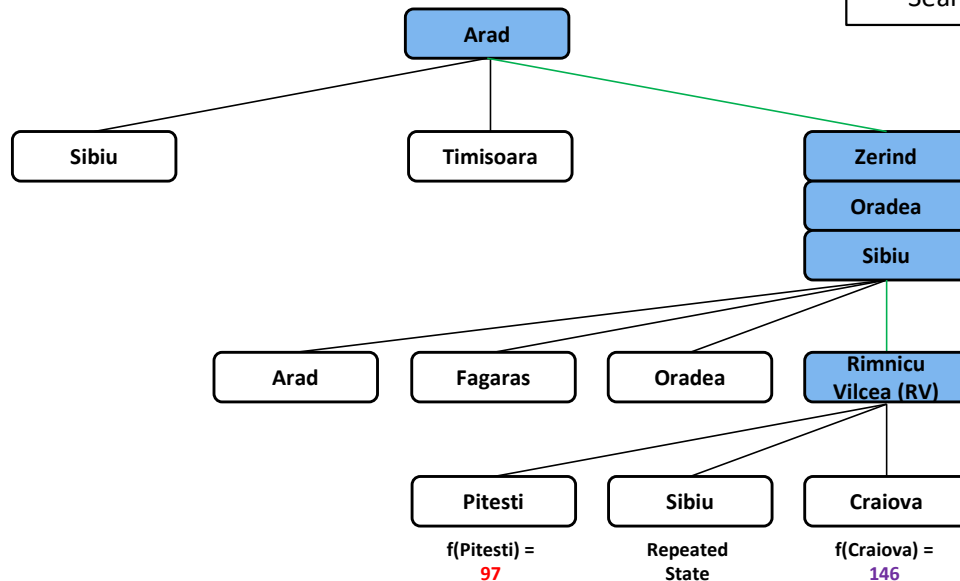


Assumption:
We don't "go" to a repeated state

State Space Graph

Search Tree


State Visited state



Alternatively:

- We could go to a repeated state end
 - get stuck there
- or
- Infinite loop

Romanian Roadtrip: Greedy Local



Assumption:

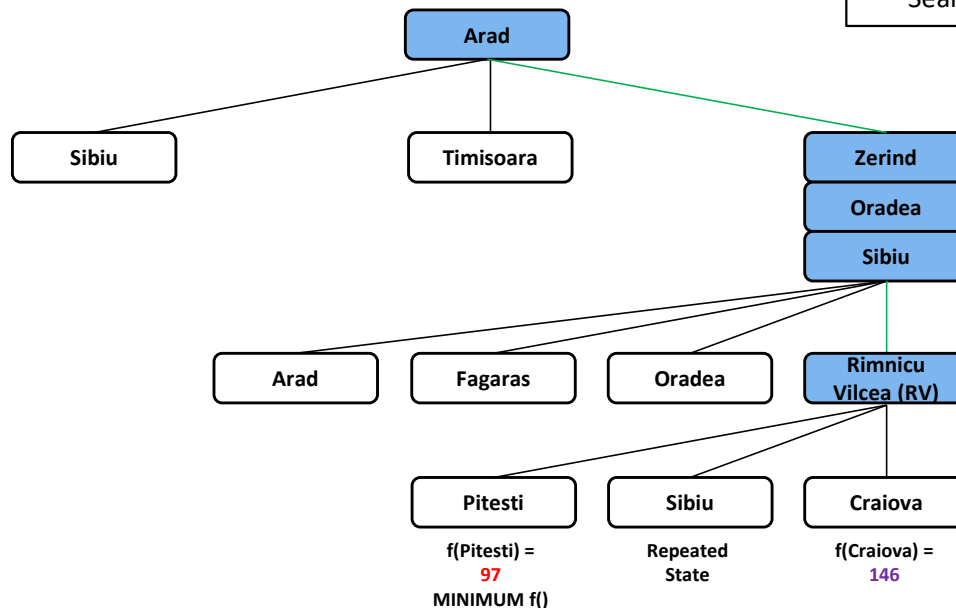
We don't "go" to a repeated state

State Space Graph

Search Tree

State

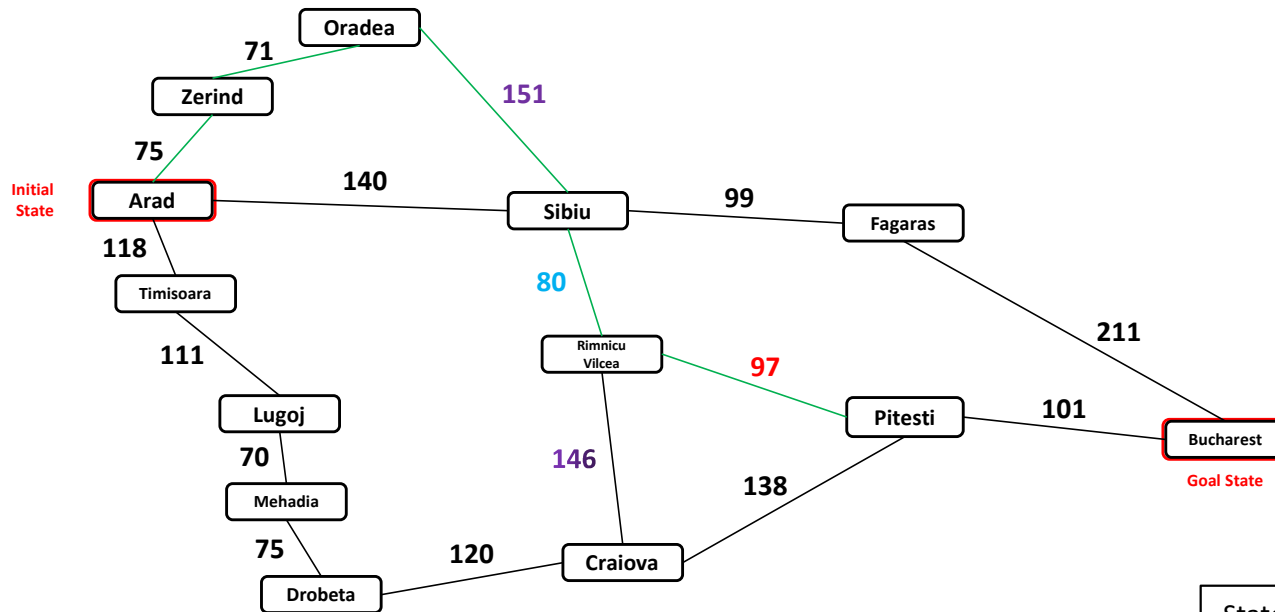
Visited state



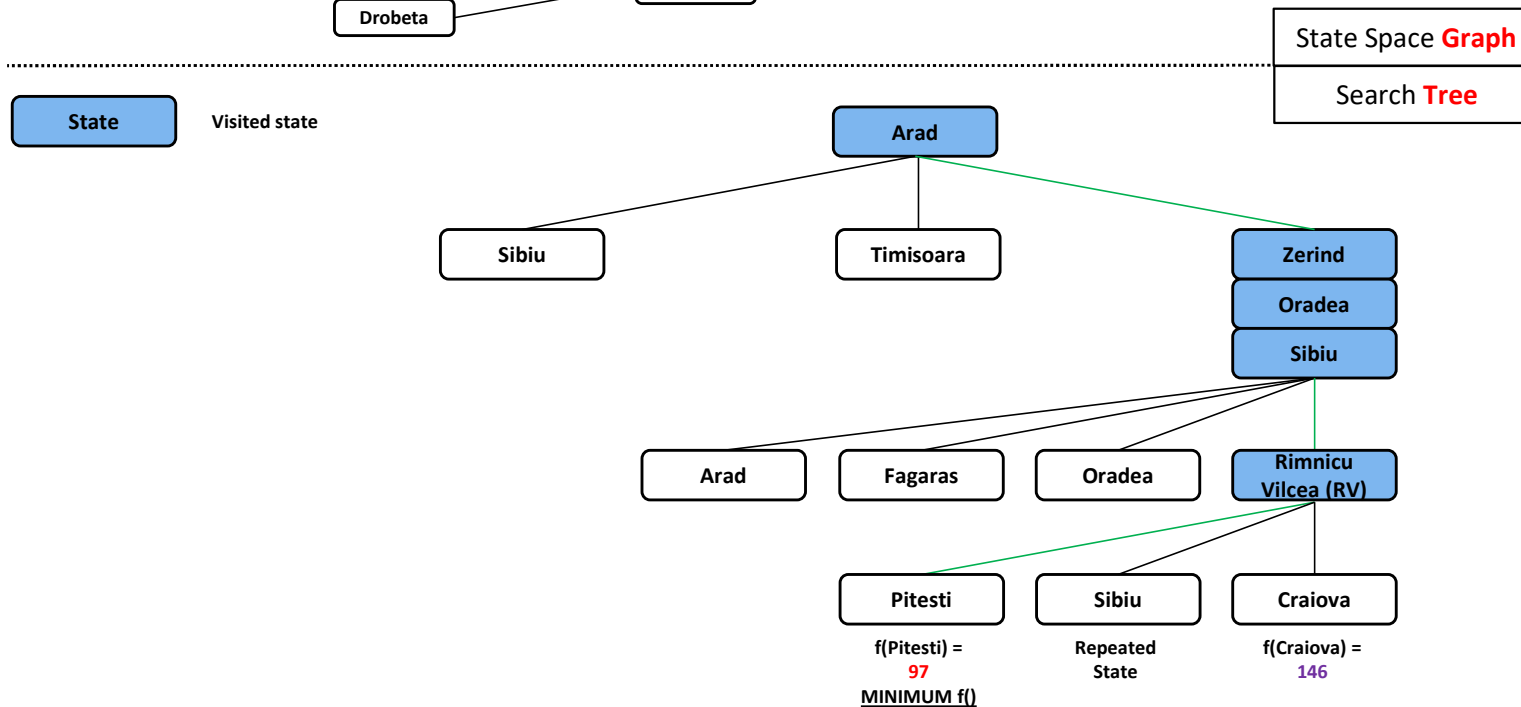
Alternatively:

- We could go to a repeated state end
 - get stuck there
- or
- Infinite loop

Romanian Roadtrip: Greedy Local



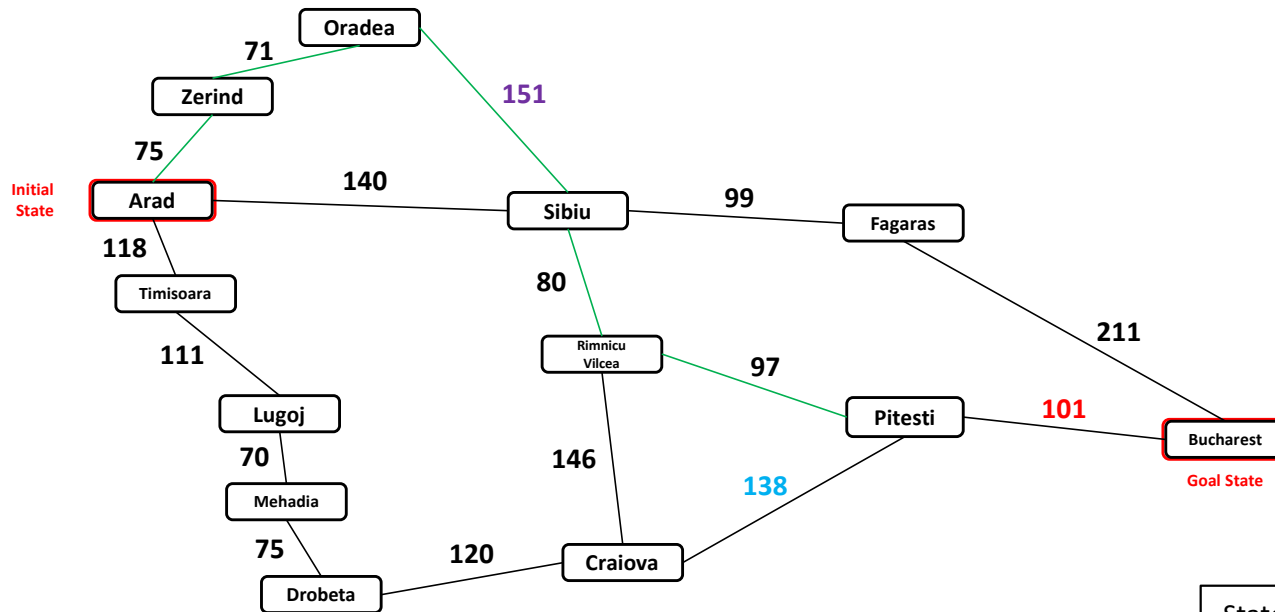
Assumption:
We don't "go" to a repeated state



Alternatively:

- We could go to a repeated state end
 - get stuck there
- or
- Infinite loop

Romanian Roadtrip: Greedy Local



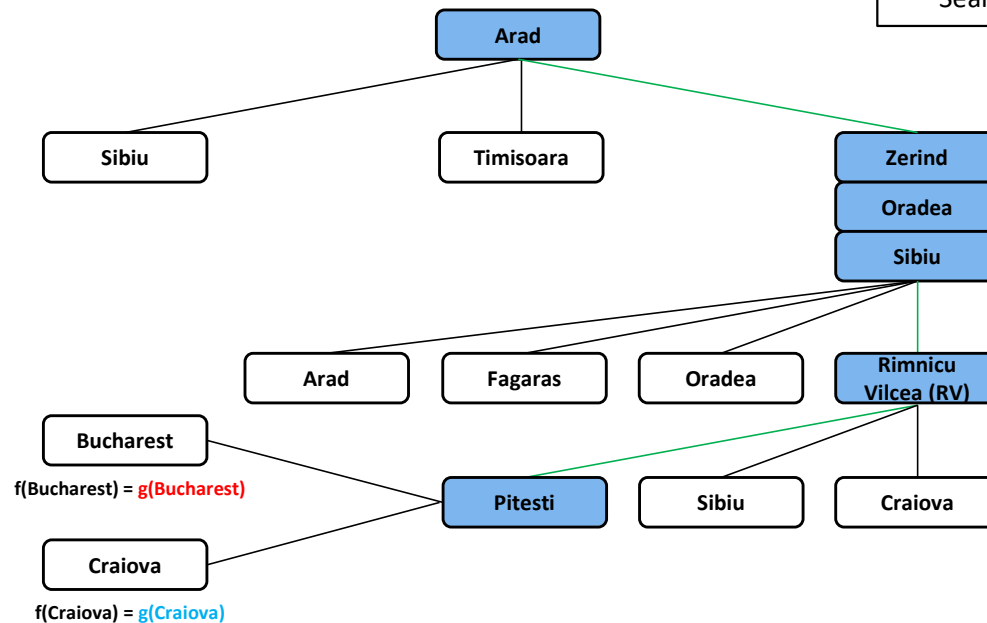
Assumption:
We don't "go" to a repeated state

State Space **Graph**

Search **Tree**

State


Visited state



Alternatively:

- We could go to a repeated state end
 - get stuck there
 - or
 - Infinite loop

Romanian Roadtrip: Greedy Local



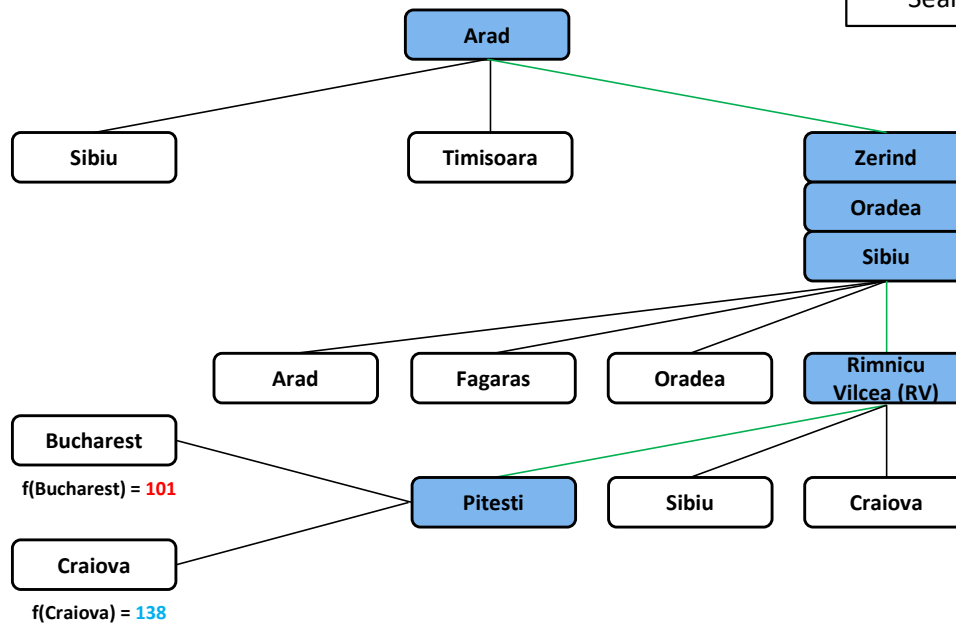
Assumption:
We don't "go" to a repeated state

State Space Graph

Search Tree

State


Visited state



Alternatively:

- We could go to a repeated state end
 - get stuck there
 - or
 - Infinite loop

Romanian Roadtrip: Greedy Local



Assumption:

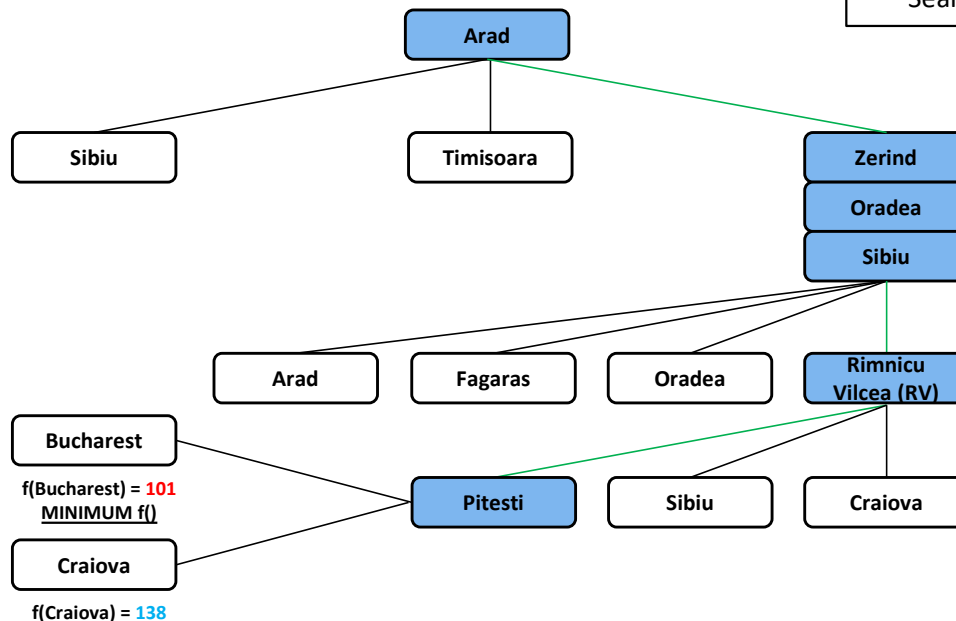
We don't "go" to a repeated state

State Space Graph

Search Tree

State


Visited state



Alternatively:

- We could go to a repeated state end
 - get stuck there
- or
- Infinite loop

Romanian Roadtrip: Greedy Local




Assumption:
We don't "go" to a repeated state

State Space Graph

Search Tree

State Visited state



Goal state: Problem solved! →

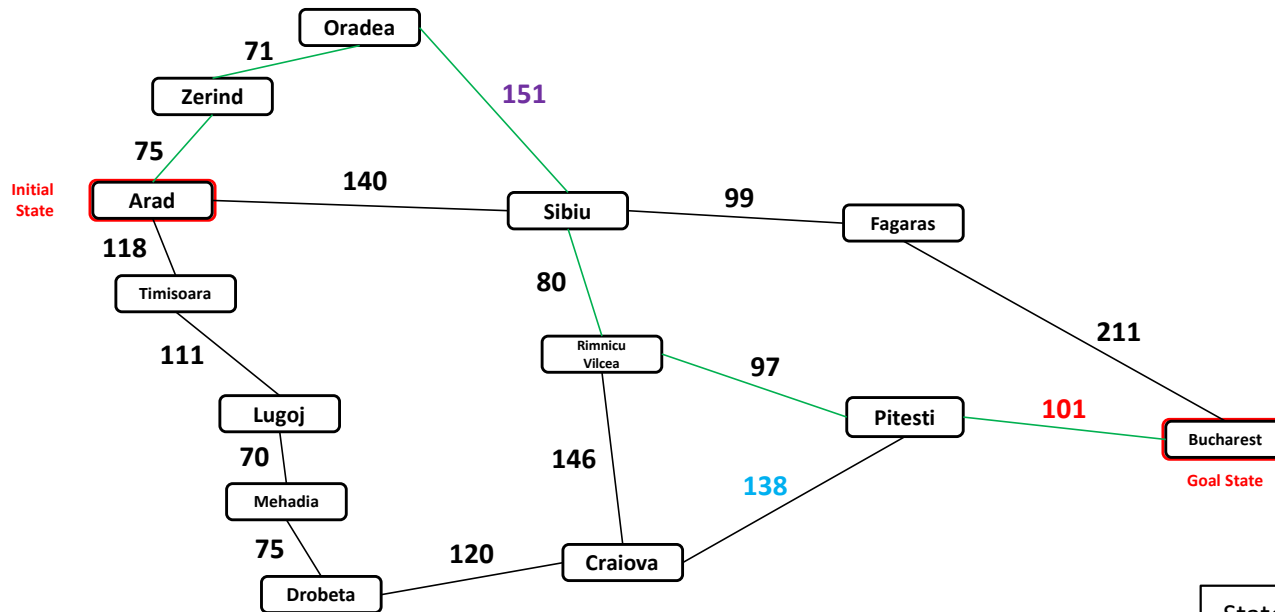
$f(\text{Bucharest}) = 101$
MINIMUM $f()$

$f(\text{Craiova}) = 138$

Alternatively:

- We could go to a repeated state end
 - get stuck there
- or
- Infinite loop

Romanian Roadtrip: Greedy Local

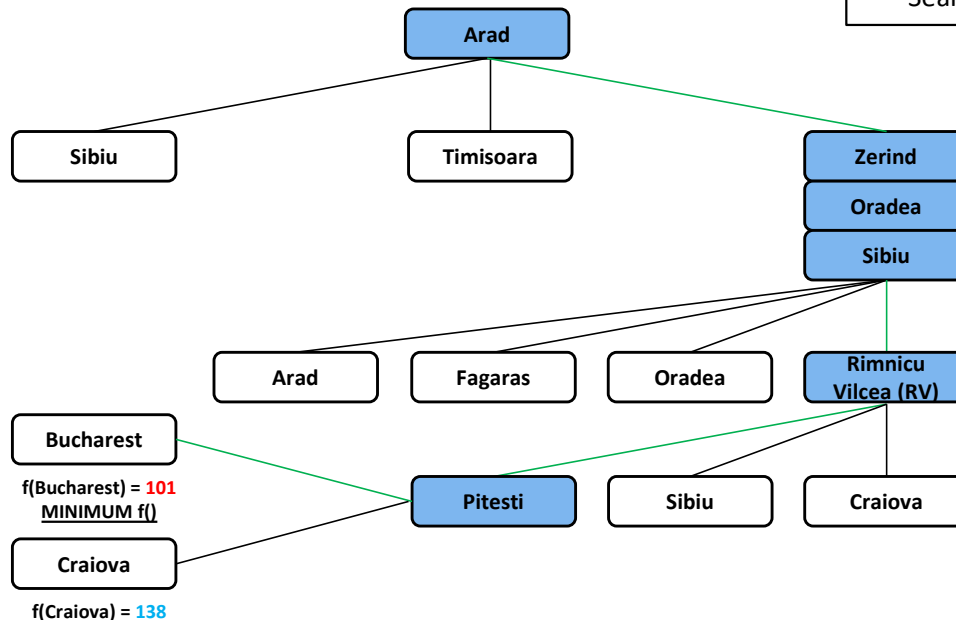


Assumption:
We don't "go" to a repeated state

State Space Graph

Search Tree

State Visited state




Alternatively:

- We could go to a repeated state end
 - get stuck there
- or
- Infinite loop

Search In Complex Environments

Difficult Environment / State Space





Reality: Environment Assumptions

Reality is not simple. What to relax?

- Fully observable?
- Single agent?
- Deterministic?
- Static?
- Episodic or sequential?
- Discrete?
- Known to the agent?

Discrete vs. Continuous Spaces



Local Search Algorithms:

**Can we not care about the path to
the goal?**

Local Search Algorithms

If the **path to the goal does not matter**, we might consider a different class of algorithms.

Local Search Algorithms


- **do not worry about paths** at all.
- Local search algorithms operate:
 - using a **single current state** (rather than multiple paths) and generally **move only to neighbors** of that state.
 - typically, the **paths followed by the search are not retained**

Local Search Algorithms

Although local search algorithms are **not systematic**, they have two key advantages:



- they **use very little memory**—usually a constant amount; and
- they can often **find reasonable solutions in large or infinite (continuous) state spaces** for which systematic algorithms are unsuitable.

1D State Space Landscape





Local search algorithms are useful for **solving pure optimization problems**, in which **the aim is to find the best state** according to an objective function.



Start “Somewhere”



Traverse/Explore Space With “Actions”



Traverse/Explore Space With “Actions”



Hill Climbing Search: Pseudocode

function HILL-CLIMBING(*problem*) **returns** a state that is a local maximum

inputs: *problem*, a problem

local variables: *current*, a node

neighbor, a node

current \leftarrow MAKE-NODE(INITIAL-STATE[*problem*])

loop do

neighbor \leftarrow a highest-valued successor of *current*

if VALUE[*neighbor*] \leq VALUE[*current*] **then return** STATE[*current*]

current \leftarrow *neighbor*

Hill Climbing and Difficult State Spaces / Environments


Hill Climbing (Greedy Local) Search

- The most primitive informed search approach
 - a naive greedy algorithm
 - evaluation function: **value of next state**
 - does not care about the “bigger picture” (for example: total search path cost)
- Practicalities:
 - does not keep track of search history


“Getting Stuck”

- **Local maxima**: a local maximum is a peak that is higher than each of its neighboring states, but lower than the global maximum.
 - Hill-climbing algorithms that reach the vicinity of a local maximum will be drawn upwards towards the peak, but will then be stuck with nowhere else to go
- **Ridge**: ridges result in a sequence of local maxima that is
 - very difficult for greedy algorithms to navigate.
- **Plateau**: a plateau is an area of the state space landscape where the evaluation function is “flat”. It can be a flat local maximum, from which no uphill exit exists, or a shoulder, from which it is possible to make progress.
 - A hill-climbing search might be unable to find its way off the plateau


Hill Climbing Problems: Local Maxima




Hill Climbing Problems: Local Maxima




Hill Climbing Problems: Local Maxima




Hill Climbing Problems: Local Maxima



Hill Climbing Problems: Plateaus



Random Walk



In mathematics, **a random walk**, sometimes known as **a drunkard's walk**, is a random process that describes a path that consists of a succession of random steps on some mathematical space.



Source: https://en.wikipedia.org/wiki/Random_walk

Measuring Searching Performance

Search algorithms can be evaluated in four ways:

- **Completeness**: Is the algorithm guaranteed to find a solution when there is one, and to correctly report failure when there is not?
- **Cost optimality**: Does it find a solution with the lowest path cost of all solutions?
- **Time complexity**: How long does it take to find a solution? (in seconds, actions, states, etc.)
- **Space complexity**: How much memory is needed to perform the search?

Random Walk vs. Hill Climbing



Simulated Annealing: Article

Optimization by Simulated Annealing

- S. Kirkpatrick, C. D. Gelatt Jr., M. P. Vecchi
- *Science* 13 May 1983:
- Vol. 220, Issue 4598, pp. 671-680
- <https://science.sciencemag.org/content/220/4598/671>