

Chapter 2: Intelligent Agents

(2.1): Agents and Environments

(2.2): The Concept of Rationality

(2.3): The Nature of Environments Be comfortable with the PEAS description and environment properties.

(2.4): The Structure of Agents You may be asked to pick the best agent type for some problem and justify your answer.

(2.5): Summary Go through the chapter summary.

Chapter 3: Solving Problems by Search

(3.1): Problem-Solving Agents Be comfortable defining a search problem.

(3.2): Example Problems

Table 1: Uniformed Search Algorithms

Criterion	Breadth-First	Uniform Cost	Depth-First	Depth-Limited	Iterative Deepening	Bidirectional (if applicable)
Complete?	Yes	Yes	No	No	Yes	Yes
Optimal cost?	Yes	Yes	No	No	Yes	Yes
Time	$O(b^d)$	$O(b^{1+C\lceil C^*/\epsilon \rceil})$				

(3.3): Search Algorithms & Uniform Search Strategies Ignore sections 3.4.4 and 3.4.5 for the exam.

(3.5): Informed (Heuristic) Search Strategies Informed search relies on domain-specific knowledge / hints that help locate the goal state. $h(n) = h(\text{State } n) = \text{relevant information about State } n$
You may be asked to solve a search problem by hand.

(3.6): Heuristic Functions

(3.7): Summary Go through the chapter summary. FOCUS ON A* algorithm

Chapter 4: Search in Complex Environments

(4.1): Local Search and Optimization Problems Local search doesn't care about the path to the goal, just getting to the goal. They're useful for pure optimization problems (finding the best state according to an objective function.) Generally use a single current state and generally move to neighbors of that state. Two key advantages are: little memory usage (usually a constant amount) and can find reasonable solutions in large of infinite (continuous) states spaces. The performance can be measured using completeness (guaranteed to find a solution when there is one and report when there isn't), cost-optimality (does it find a solution with the lowest path cost of all solutions), or time or space complexity.

Hill-climbing search

The most primitive informed search approach; it is a naive greedy algorithm and the objective function is the value of the next state. The agent can get stuck on peaks (local maxima), ridges (sequences of peaks), and plateaus (areas where the evaluation function has the same value).

Algorithm 0.1 Hill-climbing search

```
1: function HILL-CLIMBING(problem) returns a state that is a local maximum
2:   current ← problem.INITIAL
3:   while true do
4:     neighbor ← a highest-valued successor state of current
5:     if VALUE(neighbor) ≤ VALUE(current) then return current
6:     end if
7:     current ← neighbor
8:   end while
9: end function
```

Simulated Annealing

Accepts a move if it improves the objective value, and accepts some "bad" moves given some probability depending on the current objective value. Converges to a global optimum; in practice, it can give excellent results.

Algorithm 0.2 Simulated Annealing

```
1: function SIMULATED-ANNEALING(problem, Schedule) returns a solution state
2:   current ← problem.INITIAL
3:   for t = 1 to ∞ do
4:     T ← SCHEDULE(t)
5:     if T == 0 then return current
6:     end if
7:     next ← a randomly selected successor of current
8:     ΔE ← VALUE(current) − VALUE(next)
9:     if ΔE > 0 then current ← next only with probability  $e^{\Delta E/T}$ 
10:    end if
11:  end for
12: end function
```

Evolutionary algorithms

Algorithm 0.3 Genetic Algorithm Pseudocode

```
1: function GENETRIC-ALGORITHM(population, fitness) returns an individual
2:   repeat
3:     weights ← WEIGHTED-BY(population, fitness)
4:     population2 ← empty list
5:     for i = 1 to SIZE(population) do
6:       parent1, parent2 ← WEIGHTED-RANDOM-CHOICES(population, weights)
7:       child ← REPRODUCE(parent1, parent2)
8:       if small random probability then child ← MUTATE(child)
9:       end if
10:      add child to population2
11:    end for
12:    population ← population2
13:  until some individual is fit enough, or enough time has elapsed
14:  return the best individual in population, according to fitness
15: end function
16: function REPRODUCE(parent1, parent2) returns an individual
17:   n ← LENGTH(parent1)
18:   c ← random number from 1 to n
19:   return Append(SUBSTRING(parent1, 1, c), SUBSTRING(parent2, c + 1, n))
20: end function
```

...and everything related to Evolutionary algorithms that I covered in class (especially: EVERYTHING about GENETIC ALGORITHM)
IGNORE TABU SEARCH

Chapter 5: Adversarial Search and Games

(5.1): Game Theory

(5.2): Optimal Decision in Games You may be asked to solve an adversarial problem by hand using Min-Max and alpha-beta pruning. Ignore section 5.2.2.

(5.3): Summary Go through the chapter summary.

Chapter 6: Constraint Satisfaction Problems

(6.1): Defining CSPs You may be asked to formally define a constraint satisfaction problem.

(6.2): Constraint Propagation: Inference in CSPs Ignore sections 6.2.4 and 6.2.5.

(6.3): Backtracking Search for CSPs Ignore sections 6.3.3 and 6.3.4.

(6.4): Summary Go through the chapter summary.

Chapter 7: Ant Colony Optimization